

Problem 3

Application layer protocols: DNS and HTTP

Transport layer protocols: UDP for DNS; TCP for HTTP

Problem 8

a) $RTT_1 + \dots + RTT_n + 2RTT_0 + 8 \times 2RTT_0 = 18RTT_0 + RTT_1 + \dots + RTT_n$

b) $RTT_1 + \dots + RTT_n + 2RTT_0 + 2 \times 2RTT_0 = 6RTT_0 + RTT_1 + \dots + RTT_n$

c) Persistent connection with pipelining. This is the default mode of HTTP.

$$RTT_1 + \dots + RTT_n + 2RTT_0 + RTT_0 = 3RTT_0 + RTT_1 + \dots + RTT_n$$

Persistent connection without pipelining, without parallel connections.

$$RTT_1 + \dots + RTT_n + 2RTT_0 + 8RTT_0 = 10RTT_0 + RTT_1 + \dots + RTT_n$$

Problem 9

a) The time to transmit an object of size L over a link of rate R is L/R. The average time is the average size of the object divided by R:

$$\Delta = (900,000 \text{ bits}) / (15,000,000 \text{ bits/sec}) = 0.06 \text{ sec}$$

The traffic intensity on the link is given by $\beta\Delta = (16 \text{ requests/sec})(0.06 \text{ sec/request}) = 0.96$. Thus, the average access delay is $(0.06 \text{ sec}) / (1 - 0.96) \approx 1.5 \text{ seconds}$. The total average response time is therefore $1.5 \text{ sec} + 3 \text{ sec} = 4.5 \text{ sec}$.

b) The traffic intensity on the access link is reduced by 60% since the 60% of the requests are satisfied within the institutional network. Thus the average access delay is $(0.06 \text{ sec}) / [1 - (0.4)(0.96)] = 0.097 \text{ seconds}$. The response time is approximately zero if the request is satisfied by the cache (which happens with probability 0.6); the average response time is $0.097 \text{ sec} + 3 \text{ sec} = 3.097 \text{ sec}$ for cache misses (which happens 40% of the time). So the average response time is $(0.6)(0 \text{ sec}) + (0.4)(3.097 \text{ sec}) = 1.239 \text{ seconds}$. Thus the average response time is reduced from 4.5 sec to 1.24 sec.

Problem 16

SMTP uses a line containing only a period to mark the end of a message body.

HTTP uses "Content-Length header field" to indicate the length of a message body.

No, HTTP cannot use the method used by SMTP, because HTTP message could be binary data, whereas in SMTP, the message body must be in 7-bit ASCII format.

Problem 22

For calculating the minimum distribution time for client-server distribution, we use the following formula:

$$D_{cs} = \max \{NF/u_s, F/d_{\min}\}$$

Similarly, for calculating the minimum distribution time for P2P distribution, we use the following formula:

$$D_{P2P} = \max \{ F/u_s, F/d_{\min}, NF / \left(u_s + \sum_{i=1}^N u_i \right) \}$$

Where, $F = 20 \text{ Gbits} = 20 * 1000 \text{ Mbits}$

(although it should be $20 * 1024 \text{ Mbits}$),

$u_s = 30 \text{ Mbps}$

$d_{\min} = d_i = 2 \text{ Mbps}$

Client Server

		N		
		10	100	1000
U	300Kbps	10000	66667	666667
	700Kbps	10000	66667	666667
	2MBps	10000	66667	666667

Peer to Peer

		N		
		10	100	1000
U	300Kbps	10000	33333	60606
	700Kbps	10000	20000	27397
	2MBps	10000	10000	10000

Problem 23

a) Consider a distribution scheme in which the server sends the file to each client, in parallel, at a rate of u_s/N . Note that this rate is less than each of the client's download rate, since by assumption $u_s/N \leq d_{\min}$. Thus each client can also receive at rate u_s/N . Since each client receives at rate u_s/N , the time for each client to receive the entire file is $F/(u_s/N) = NF/u_s$. Since all the clients receive the file in NF/u_s , the overall distribution time is also NF/u_s .

b) Consider a distribution scheme in which the server sends the file to each client, in parallel, at a rate of d_{\min} . Note that the aggregate rate, $N d_{\min}$, is less than the server's link rate u_s , since by assumption $u_s/N \geq d_{\min}$. Since each client receives at rate d_{\min} , the time for each client to receive the entire file is F/d_{\min} . Since all the clients receive the file in this time, the overall distribution time is also F/d_{\min} .

Problem 26

Yes. His first claim is possible, as long as there are enough peers staying in the swarm for a long enough time. Bob can always receive data through optimistic unchoking by other peers.

His second claim is also true. He can run a client on each host, let each client "free-ride," and combine the collected chunks from the different hosts into a single file. He can even write a small scheduling program to make the different hosts ask for different chunks of the file. This is actually a kind of Sybil attack in P2P networks.