


CUDA

Scott Cheng

Recall: Reduction

Given a vector $[v_0, v_1, \dots, v_n]$ it returns a scalar $v_0 + v_1 + \dots + v_n$



Performance for 4M element reduction



	Time (2 ²² ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
Kernel 4: first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x
Kernel 5: unroll last warp	0.536 ms	31.289 GB/s	1.8x	15.01x
Kernel 6: completely unrolled	0.381 ms	43.996 GB/s	1.41x	21.16x
Kernel 7: multiple elements per thread	0.268 ms	62.671 GB/s	1.42x	30.04x

Kernel 7 on 32M elements: 73 GB/s!

Warp Shuffle Functions

T `__shfl_sync(unsigned mask, T var, int srcLane, int width=warpSize);`

Direct copy from indexed lane

T `__shfl_up_sync(unsigned mask, T var, unsigned int delta, int width=warpSize);`

Copy from a lane with lower ID relative to caller

T `__shfl_down_sync(unsigned mask, T var, unsigned int delta, int width=warpSize);`

Copy from a lane with higher ID relative to caller

T `__shfl_xor_sync(unsigned mask, T var, int laneMask, int width=warpSize);`

Copy from a lane based on bitwise XOR of own lane ID


Example: Broadcast of a single value across a warp

```
#include <stdio.h>
```

```
__global__ void bcast(int arg) {  
    int laneId = threadIdx.x & 0x1f;  
    int value; // Note unused variable for all threads except lane 0  
    if (laneId == 0)  
        value = func(arg);  
    // Synchronize all threads in warp, and get "value" from lane 0  
    value = __shfl_sync(0xffffffff, value, 0);  
    // every thread in the warp get the same "value" now  
}
```


Reduction: __shfl_xor_sync

```
for (int i=1; i<32; i*=2)
    value += __shfl_xor_sync(-1, value, i);
```



Reduction: __shfl_down_sync

```
for (int i=16; i>0; i=i/2)
    value += __shfl_down_sync(-1, value, i);
```



Warp Reduce Functions

Supported by devices of compute capability 8.x or higher.

// add/min/max

```
unsigned __reduce_add_sync(unsigned mask, unsigned value);  
unsigned __reduce_min_sync(unsigned mask, unsigned value);  
unsigned __reduce_max_sync(unsigned mask, unsigned value);  
int __reduce_add_sync(unsigned mask, int value);  
int __reduce_min_sync(unsigned mask, int value);  
int __reduce_max_sync(unsigned mask, int value);
```

// and/or/xor

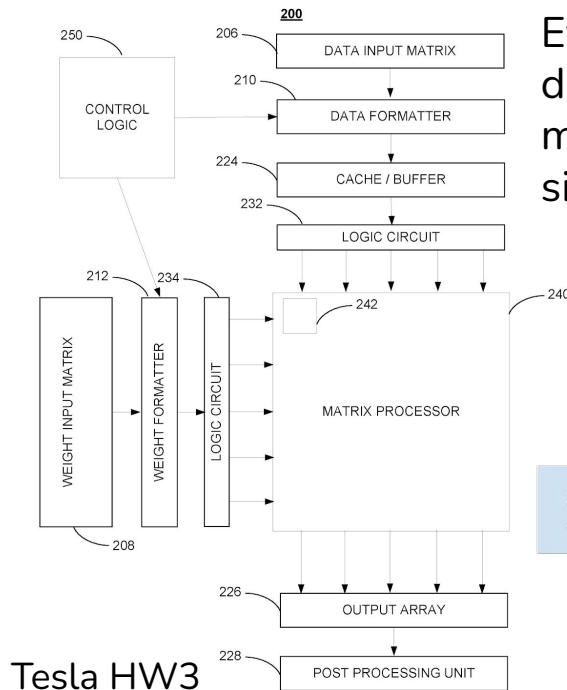
```
unsigned __reduce_and_sync(unsigned mask, unsigned value);  
unsigned __reduce_or_sync(unsigned mask, unsigned value);  
unsigned __reduce_xor_sync(unsigned mask, unsigned value);
```


Tensor Core

Proliferation of TCUs

Many companies have now developed chips with TCUs to accelerate DNN computation:


- NVIDIA TensorCores
- Google's TPU
- Tesla HW3
- Apple's A11
- Intel AMX/DLBoost
- Many many startups



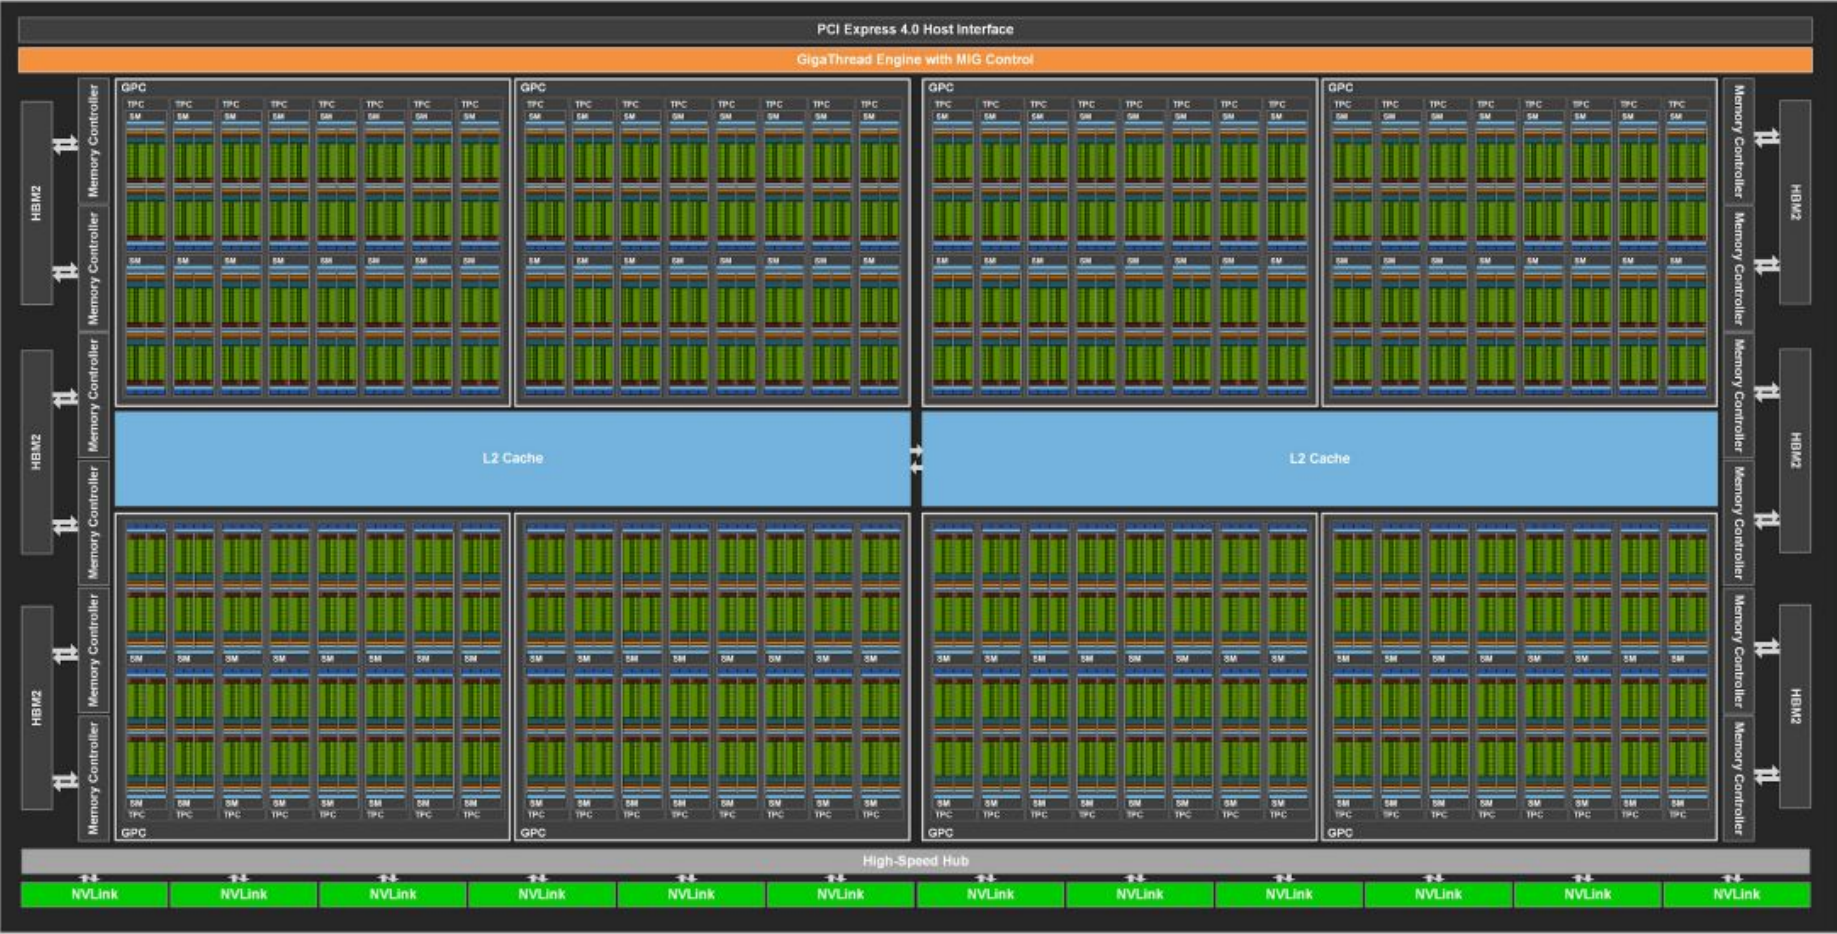
Tesla HW3

Even though they are named differently, they are all designed as matrix multiplication units for fixed size matrices.

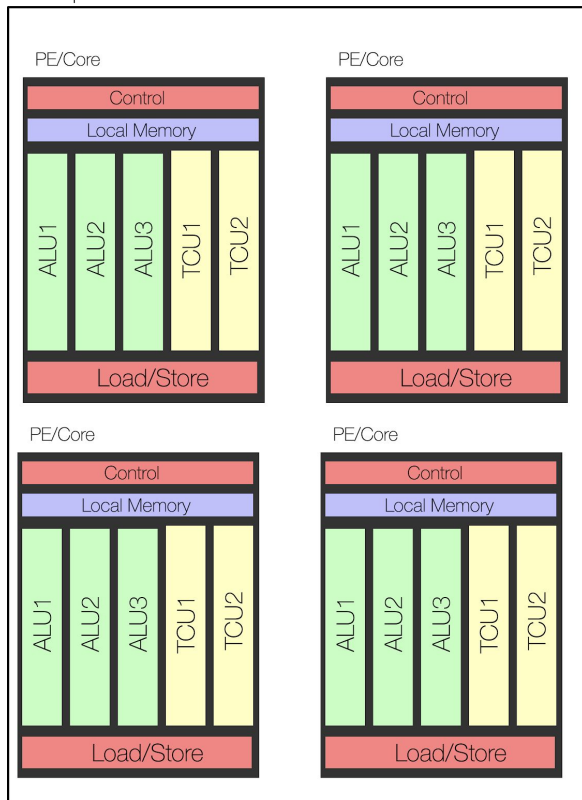
Google TPU



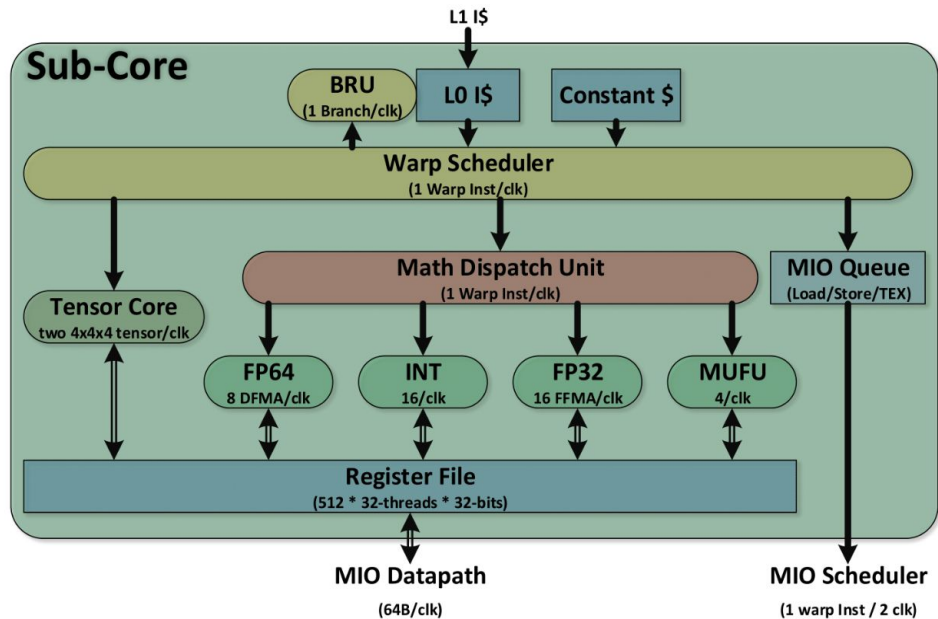
A100 block diagram: 108 SMs



Chip





NVIDIA V100 SubCore



TensorCore Operations

- Fixed size matrix multiplication: usually 4x4, 16x16, or 64x64
- NVIDIA supports non-square matrix dimensions



Warp Matrix Functions

```
template<typename Use, int m, int n, int k, typename T, typename Layout=void> class fragment;
```

```
void load_matrix_sync(fragment<...> &a, const T* mptr, unsigned ldm);
```

```
void load_matrix_sync(fragment<...> &a, const T* mptr, unsigned ldm, layout_t layout);
```

```
void store_matrix_sync(T* mptr, const fragment<...> &a, unsigned ldm, layout_t layout);
```

```
void fill_fragment(fragment<...> &a, const T& v);
```

```
void mma_sync(fragment<...> &d, const fragment<...> &a, const fragment<...> &b, const  
fragment<...> &c, bool satf=false);
```

wmma::fragment

Element type of fragment. Float is only supported for accumulator fragments (mixed precision)

Layout of fragment: column or row major order.

```
wmma::fragment<wmma::matrix_a | wmma::matrix_b | wmma::accumulator, M, N, K, half | float, col|row> a_frag;
```

Kind of fragment:
An A matrix, B matrix, or a C (accumulator) matrix

Dimensions of fragment. For our purposes these are 16, 16, 16, but some other values are possible

A memory fragments (internally a set of registers).

Result = A * B + C

wmma:fill_fragment



```
wmma::fill_fragment(acc_frag, half(C));
```

Fills the fragment with some constant value **C**

wmma::load_matrix_sync



```
wmma::load_matrix_sync(a_frag, a + aRow + aCol * lda, lda)
```

Loads data from global or shared memory with specified stride into the fragment. Here we load a tile of a matrix (at `&a[aRow+aCol*lda]`) into `a_frag` with `lda` as the stride

wmma:store_matrix_sync



```
wmma::store_matrix_sync(c + cRow + cCol * ldc, c_frag, ldc, wmma::mem_col_major);
```

Stores the fragment into global or shared memory with specified stride and layout.

Example

```
1  #include <mma.h>
2  using namespace nvcuda::wmma;
3  __global__ void dot_wmma_16x16(half *a, half *b, half *c) {
4      fragment<matrix_a, 16, 16, 16, half, col_major> a_frag;
5      fragment<matrix_b, 16, 16, 16, half, row_major> b_frag;
6      fragment<accumulator, 16, 16, 16, half> c_frag;
7      load_matrix_sync(a_frag, a, /* leading dim */ 16);
8      load_matrix_sync(b_frag, b, /* leading dim */ 16);
9      fill_fragment(c_frag, 0.0f);
10     mma_sync(c_frag, a_frag, b_frag, c_frag);
11     store_matrix_sync(c, c_frag, 16, row_major);
12 }
```

Listing 1: A simple CUDA kernel performing $\langle 16, 16, 16 \rangle$ matrix multiplication ($C = A.B + C$) in half precision using the CUDA WMMA API.

Recall

Warp reduction is performed using the shuffle instructions

```
1  __device__ half warp_reduce(half val) {
2      for (int offset=WARP_SIZE/2; offset>0; offset/=2)
3          val += __shfl_down_sync(0xFFFFFFFFFU, val, mask);
4      return val; }
5  __device__ half warp_scan(half val) {
6      for (int offset=1; offset<WARP_SIZE; offset*=2) {
7          auto n = __shfl_up_sync(0xFFFFFFFFFU, val, mask);
8          if (laneid >= offset) val += n; }
9      return val; }
```

Listing 2: NVIDIA's recommended warp-level reduction and scan implementations utilizing shuffle instructions.

We want to do the same thing, but now using TensorCores

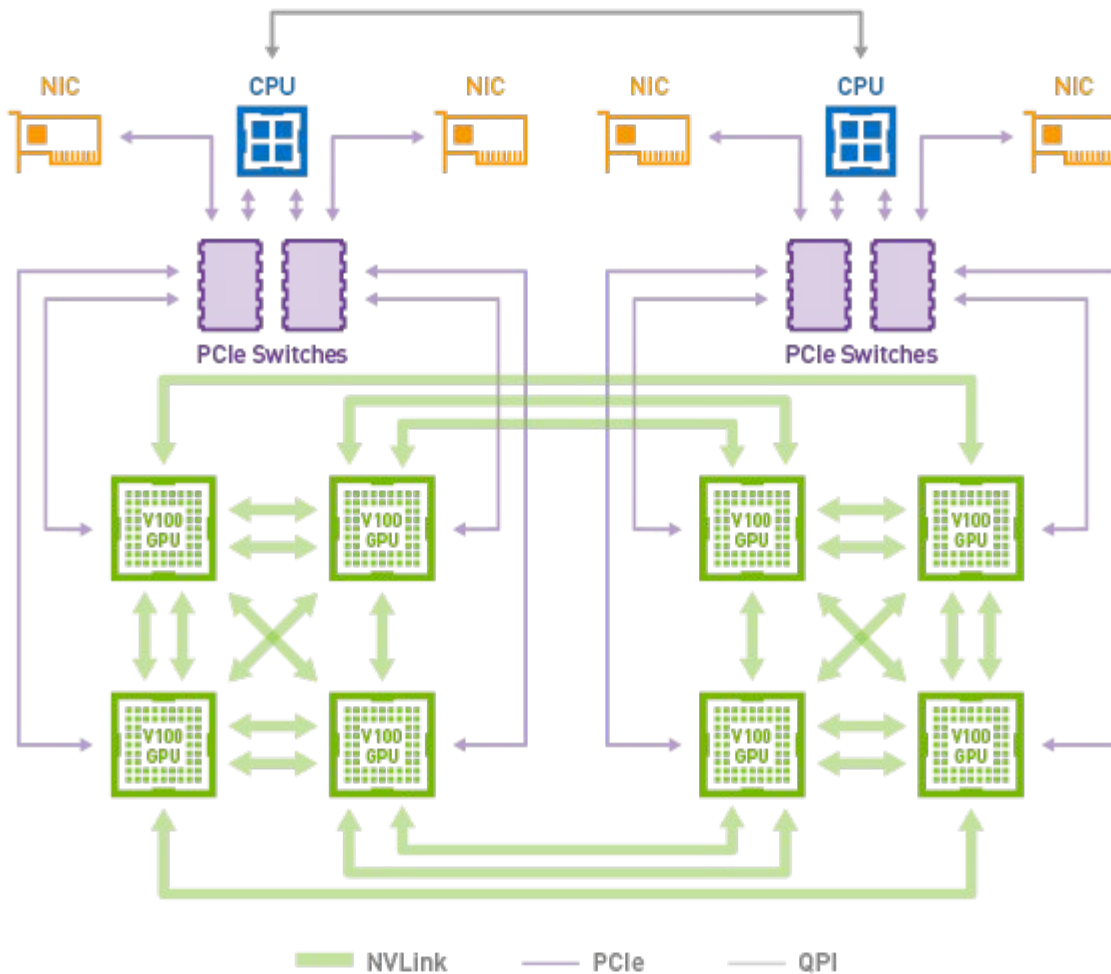
Reduction

Easy to represent it as a matrix multiplication

$$\text{Reduction}[a_1, a_2, \dots, a_n] = \begin{pmatrix} 1 & 1 & \dots & 1 \\ 0 & 0 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 0 \end{pmatrix} \cdot \begin{pmatrix} a_1 & a_2 & \dots & a_n \\ 0 & 0 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 0 \end{pmatrix}^T = \begin{pmatrix} \sum_{i=1}^n a_i & 0 & \dots & 0 \\ 0 & 0 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 0 \end{pmatrix}$$

There is a lot of compute waste here

NVLink



PCIe Gen3: 32GB/sec

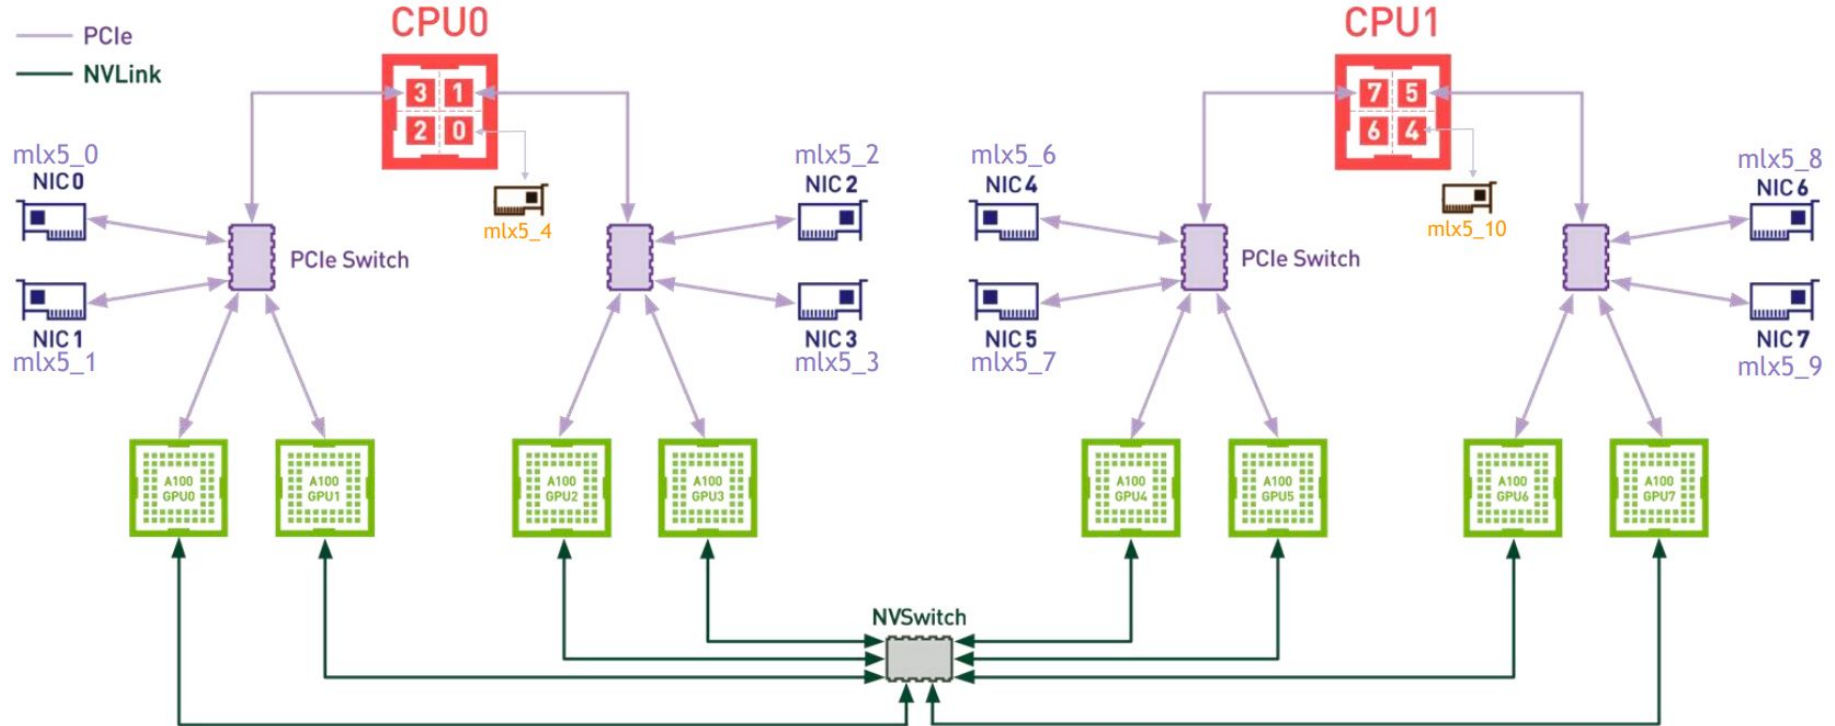
HBM2 BW: 900GB/sec

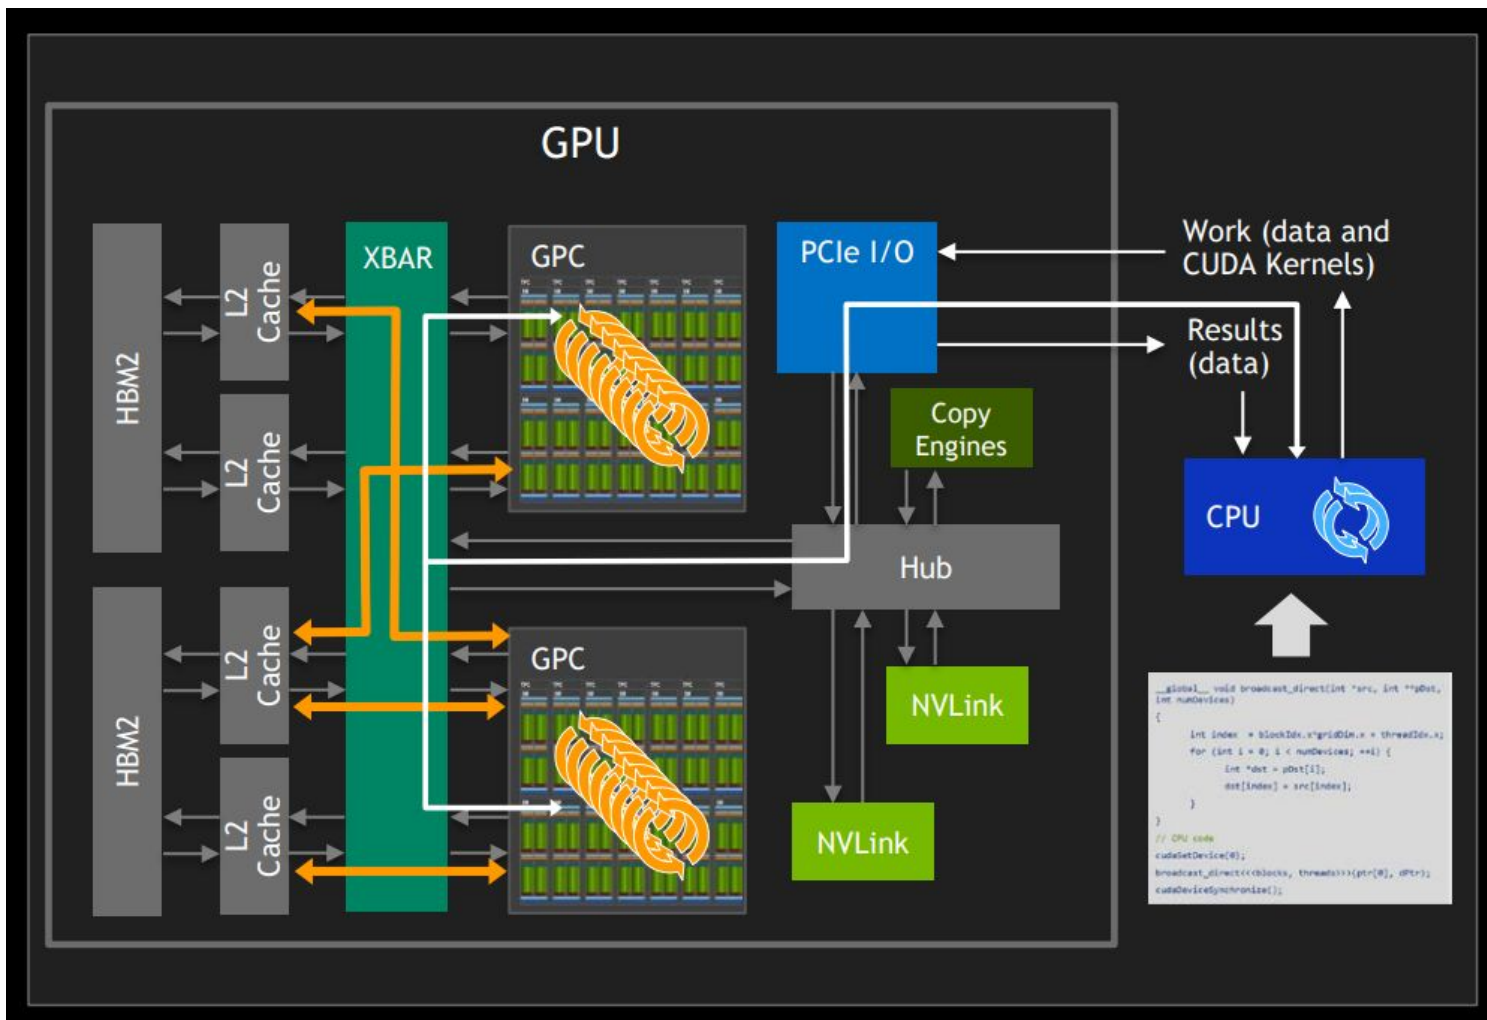
NVLINK 2.0: 300GB/sec

DGX A100

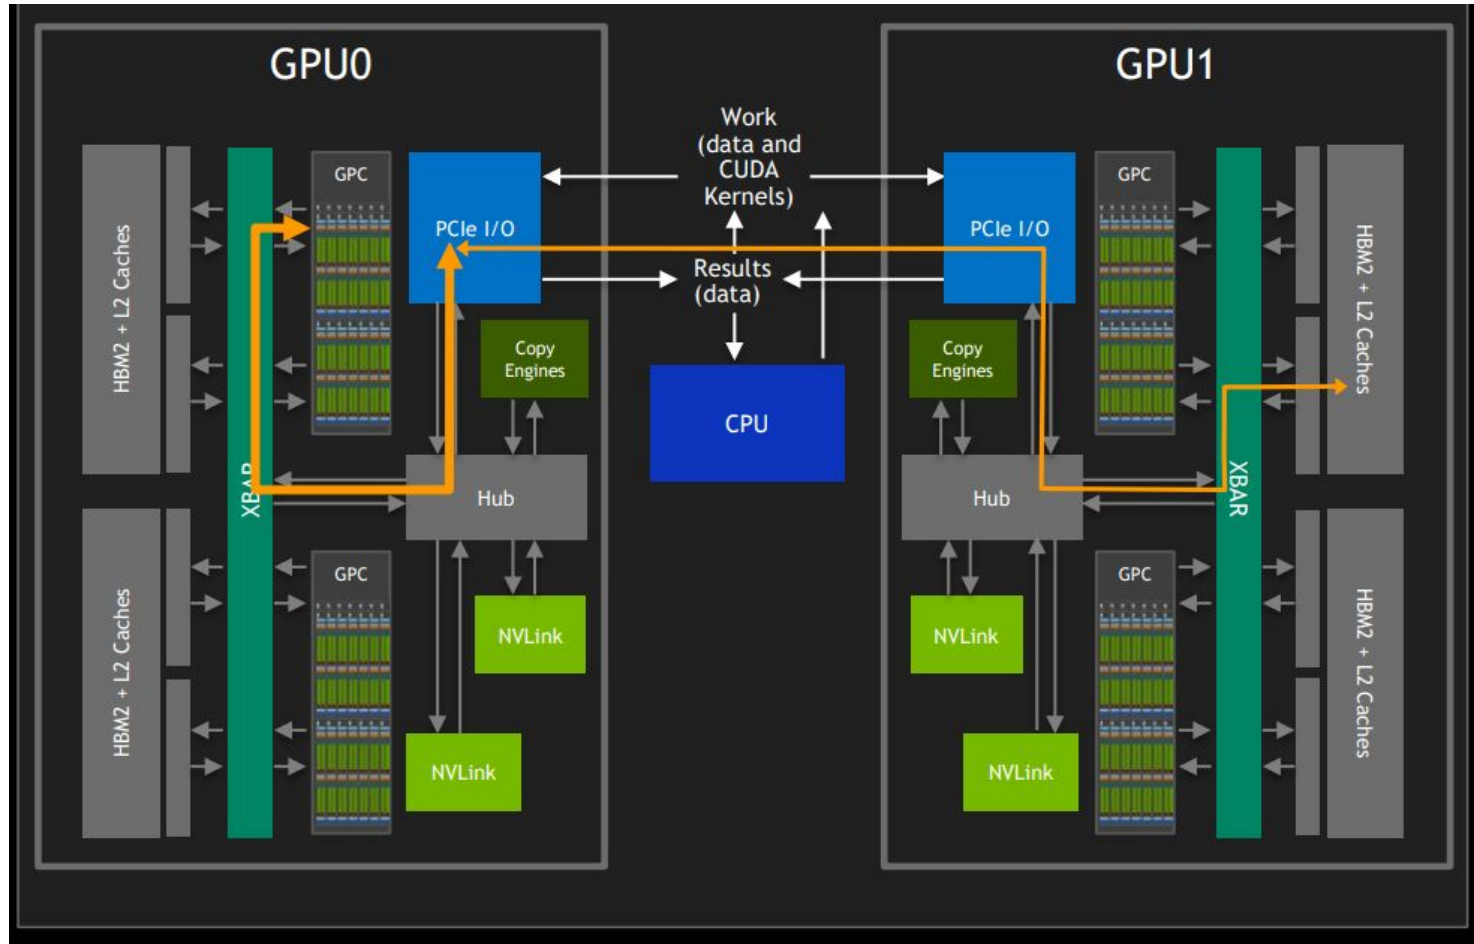
High-level Topology Overview (with options)

Data plane (can be used as eth or IB)
Compute plane (IB)

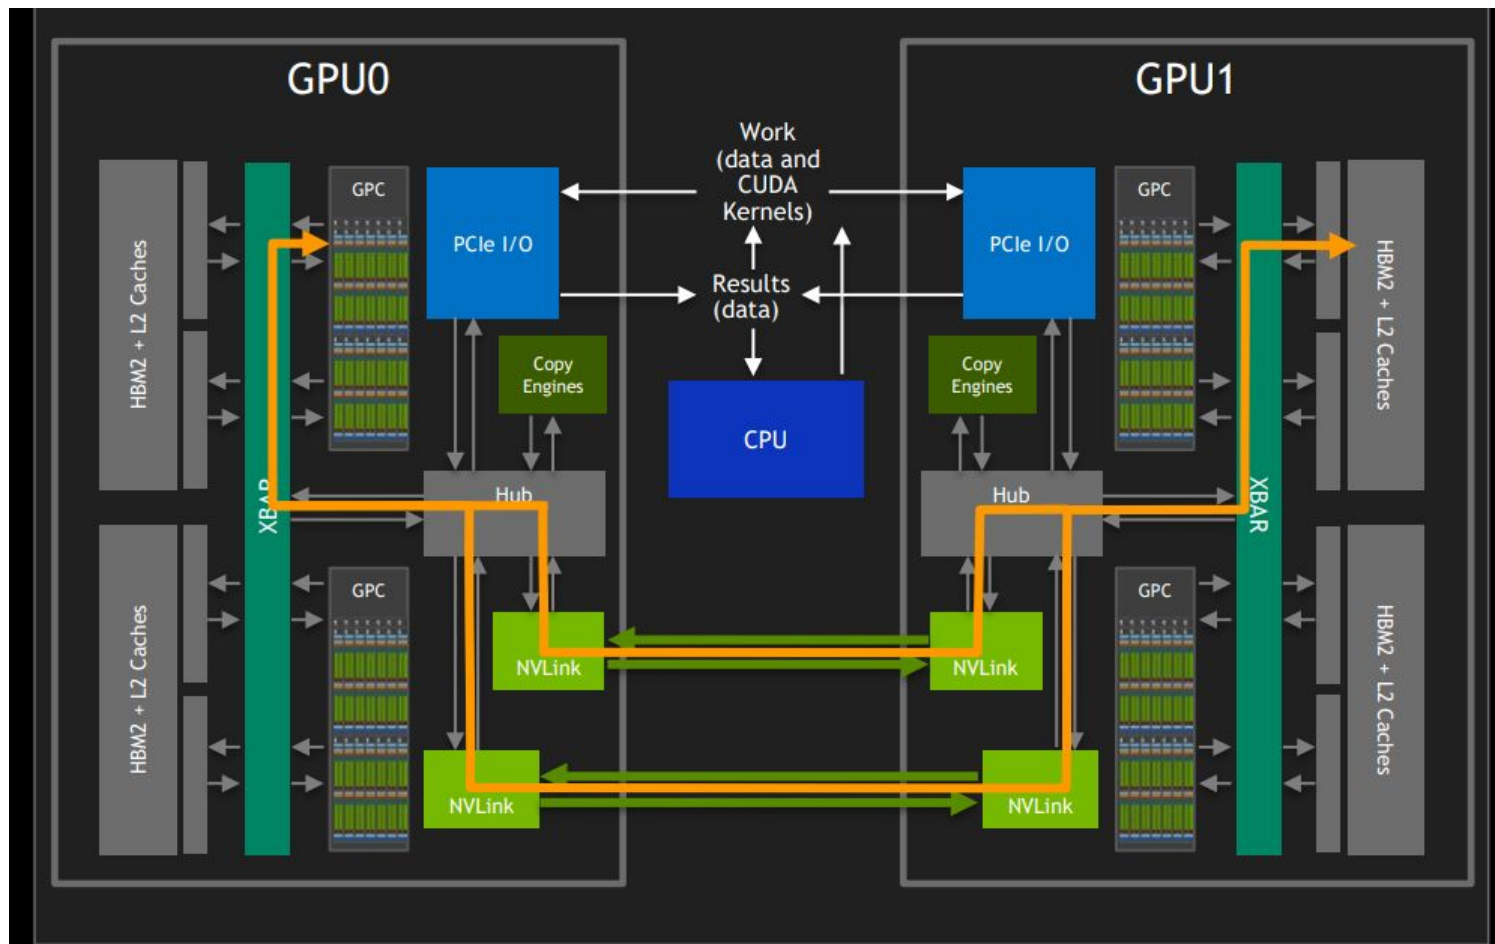




Interactions with CPU compete with GPU-to-GPU

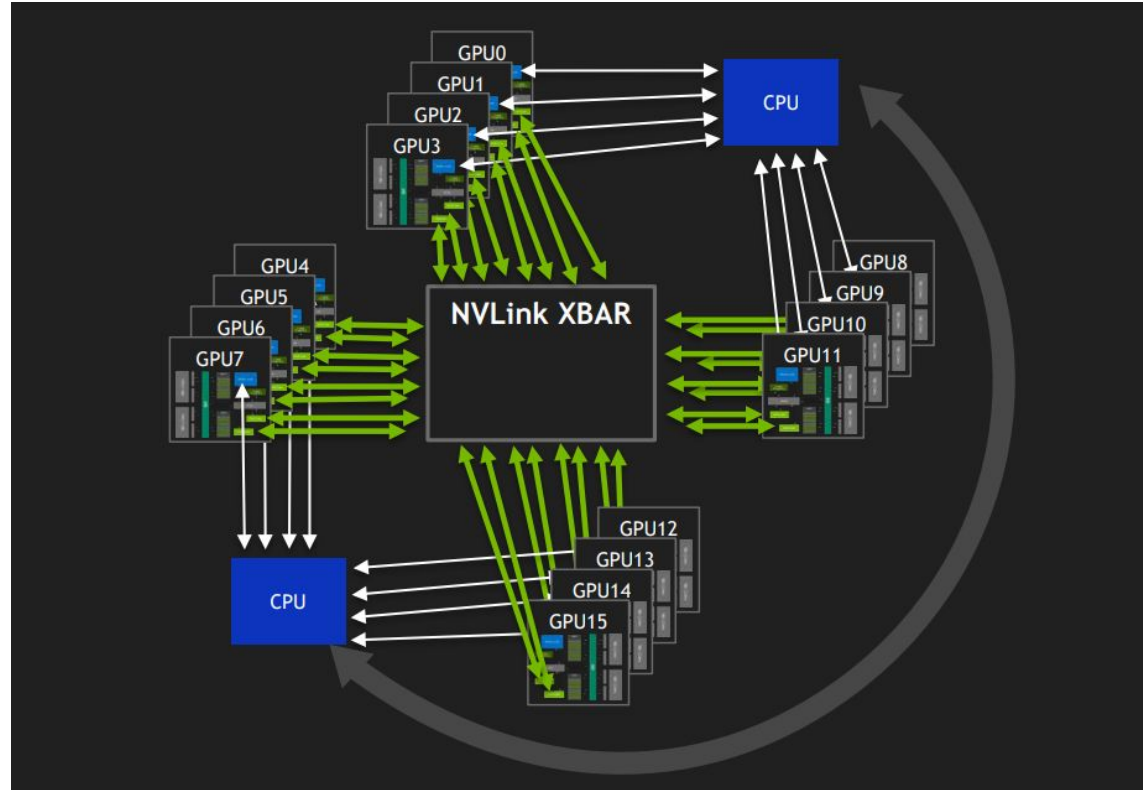


- NVLinks are effectively a “bridge” between XBARs
- No collisions with PCIe traffic




THE “ONE GIGANTIC GPU” IDEAL



- Highest number of GPUs possible
- Single GPU Driver process controls all work across all GPUs
- From perspective of GPCs, all HBM2s can be accessed without intervention by other processes (LD/ST instructions, Copy Engine RDMA, everything “just works”)
- Access to all HBM2s is independent of PCIe
- Bandwidth across bridged XBARs is as high as possible (some NUMA is unavoidable)




NVSwitch

- No NVSwitch
 - Connect GPU directly
 - Aggregate NVLinks into gangs for higher bandwidth
 - Interleaved over the links to prevent camping
 - Max bandwidth between two GPUs limited to the bandwidth of the gang
- With NVSwitch
 - Interleave traffic across all the links and to support full bandwidth between any pair of GPUs
 - Traffic to a single GPU is nonblocking, so long as aggregate bandwidth of six NVLinks is not

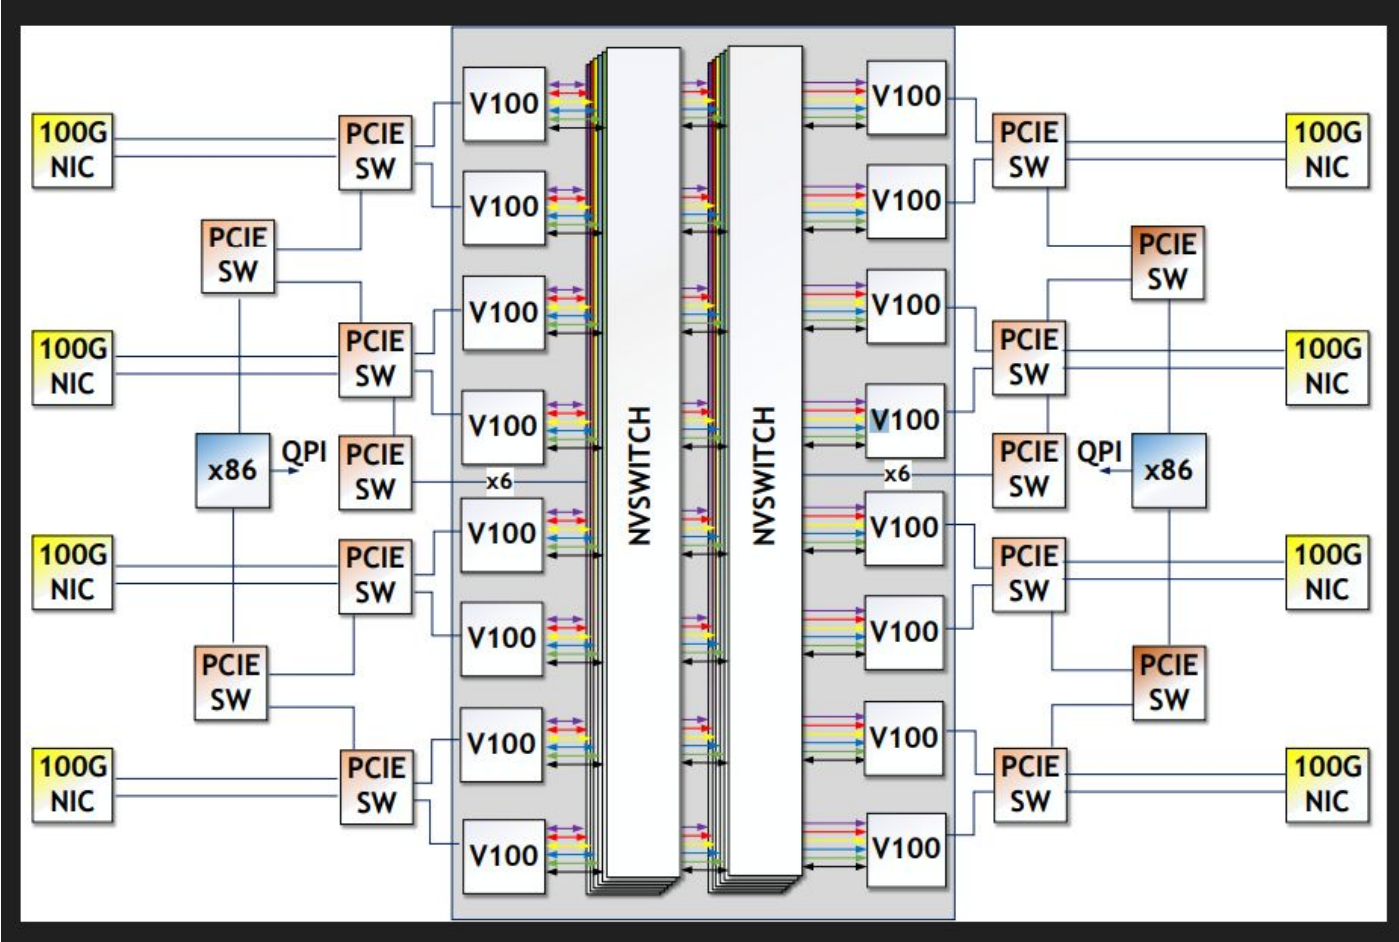




DGX-2




DGX-2 PCIe Network



NVLINK-ENABLED SERVER GENERATIONS


Any-to-Any Connectivity with NVSwitch



2016

DGX-1 (P100)


140GB/s Bisection BW
40GB/s AllReduce BW



2018

DGX-2 (V100)

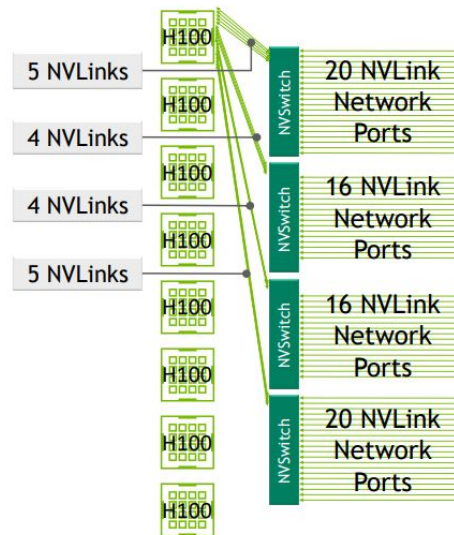
2.4TB/s Bisection BW
75GB/s AllReduce BW



2020

DGX A100

2.4TB/s Bisection BW
150GB/s AllReduce BW



2022

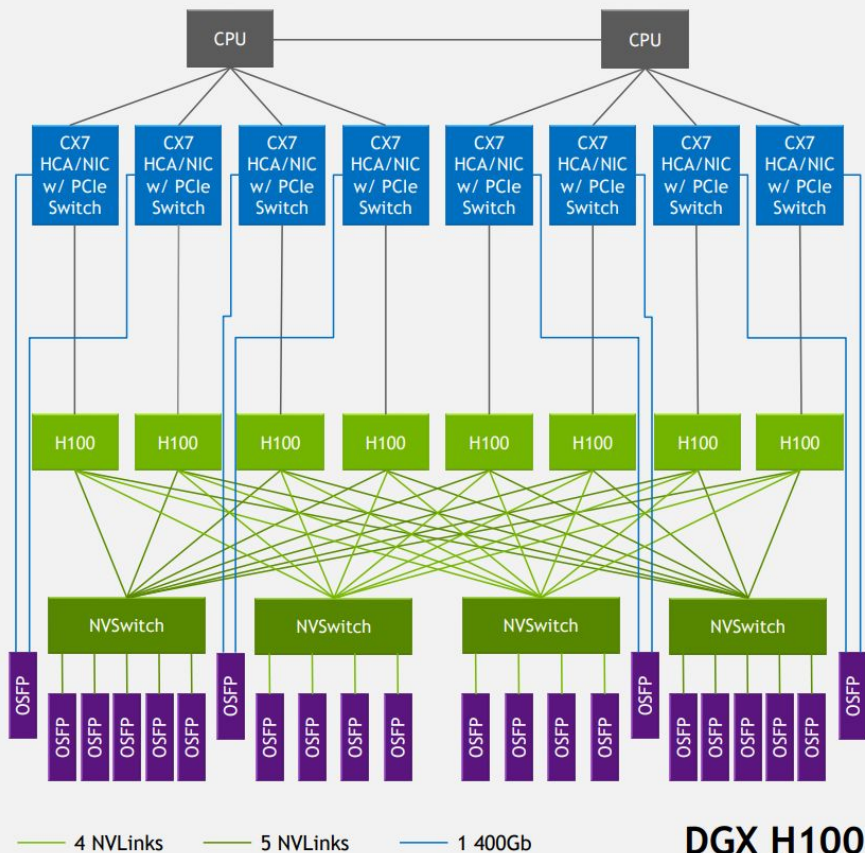
DGX H100

3.6TB/s Bisection BW
450GB/s AllReduce BW

MAPPING TO TRADITIONAL NETWORKING

NVLink Network is Tightly Integrated with GPU

Concept	Traditional Example	NVLink Network
Physical Layer	400G electrical/optical media	Custom-FW OSFP
Data Link Layer	Ethernet	NVLink custom on-chip HW and FW
Network Layer	IP	New NVLink Network Addressing and Management Protocols
Transport Layer	TCP	NVLink custom on-chip HW and FW
Session Layer	Sockets	SHARP groups CUDA export of Network addresses of data-structures
Presentation Layer	TSL/SSL	Library abstractions (e.g., NCCL, NVSHMEM)
Application Layer	HTTP/FTP	AI Frameworks or User Apps
NIC	PCIe NIC (card or chip)	Functions embedded in GPU and NVSwitch
RDMA Off-Load	NIC Off-Load Engine	GPU-internal Copy Engine
Collectives Off-Load	NIC/Switch Off-Load Engine	NVSwitch-internal SHARP Engines
Security Off-Load	NIC Security Features	GPU-internal Encryption and “TLB” Firewalls
Media Control	NIC Cable Adaptation	NVSwitch-internal OSFP-cable controllers



DGX H100

DGX H100: DATA-NETWORK CONFIGURATION

Full-BW Intra-Server NVLink

- All 8 GPUs can simultaneously saturate 18 NVLinks to other GPUs within server
- Limited only by over-subscription from multiple other GPUs

Half-BW NVLink Network

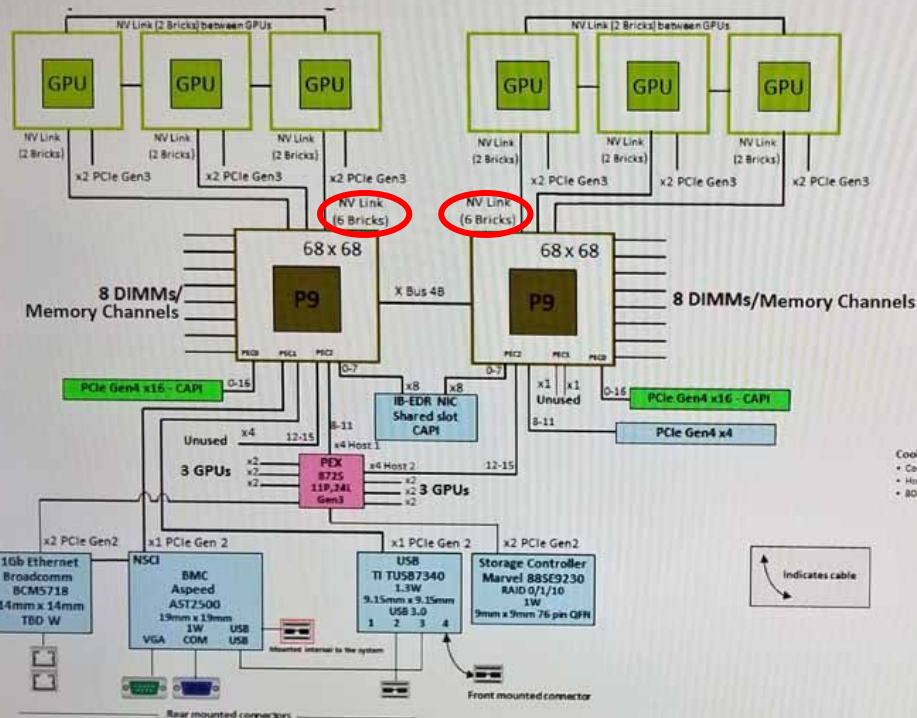
- All 8 GPUs can half-subscribe 18 NVLinks to GPUs in other servers
- 4 GPUs can saturate 18 NVLinks to GPUs in other servers
- Equivalent of full-BW on AllReduce with SHARP
- Reduction in All2All BW is a balance with server complexity and costs

Multi-Rail InfiniBand/Ethernet

- All 8 GPUs can independently RDMA data over its own dedicated 400 Gb/s HCA/NIC
- 800 GBps of aggregate full-duplex to non-NVLink Network devices

IBM HPC System Utilizing PCIe Gen4

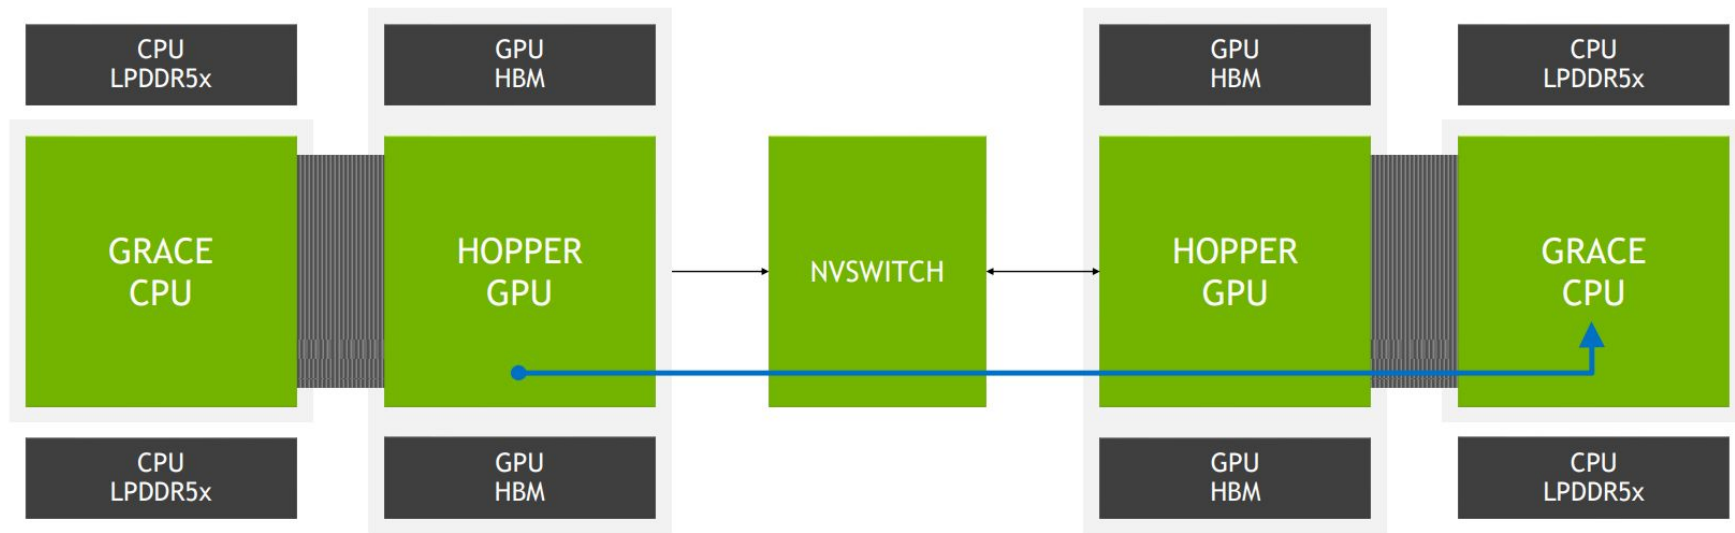
IBM Power9



Rank	System	Cores	Rmax (PFlop/s)	Rpeak (PFlop/s)	Power (kW)
1	Frontier - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE DOE/SC/Oak Ridge National Laboratory United States	8,730,112	1,102.00	1,685.65	21,100
2	Supercomputer Fugaku - Supercomputer Fugaku, A64FX 48C 2.26GHz, Tofu interconnect D, Fujitsu RIKEN Center for Computational Science Japan	7,630,848	442.01	537.21	29,899
3	LUMI - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE EuroHPC/CSC Finland	2,220,288	309.10	428.70	6,016
4	Leonardo - BullSequana XH2000, Xeon Platinum 8358 32C 2.6GHz, NVIDIA A100 SXM4 64 GB, Quad-rail NVIDIA HDR100 Infiniband, Atos EuroHPC/CINECA Italy	1,463,616	174.70	255.75	5,610
5	Summit - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM DOE/SC/Oak Ridge National Laboratory United States	2,414,592	148.60	200.79	10,096


NVLINK-C2C

Superchip Scaling | CPU/GPU | Extended GPU Memory



Enables remote NVLINK connected GPUs, to access Grace's memory at native NVLINK speeds

TOP500 Supercomputers by Processor Family



Reference

- Dakkak, Abdul et al. “Accelerating reduction and scan using tensor core units.” *Proceedings of the ACM International Conference on Supercomputing* (2018): n. Pag.
- (HC2018) NVSWITCH AND DGX-2 NVLINK-SWITCHING CHIP AND SCALE-UP COMPUTE SERVER
- (HC2022) THE NVLINK-NETWORK SWITCH: NVIDIA’S SWITCH CHIP FOR HIGH COMMUNICATION-BANDWIDTH SUPERPODS