

1. The MPI program sums  $n$  numbers, and thus  $T_S = n$ .

$$T_p = (\text{sum } n/p \text{ numbers in each process}) + (\text{sum} + \text{sendrecv}) * (\text{stages of reduction})$$

$$T_p = \frac{n}{p} + 2 \log P$$

$$T_o = PT_p - T_S = (n + 2P \log P) - n = 2P \log P$$

$$\text{Isoefficiency function: } W = K(2P \log P)$$

2.

a)

First, each processor performs a local prefix sum of its  $n/p$  numbers. In the second step, the  $p$  processors compute prefix sums of  $p$  numbers by using the last prefix sum resulting from the local computation on each processor. This step takes  $(1 + t_s + t_w) \log p$  time. Finally, the result of the parallel prefix sums operation of the second step is added to all the  $n/p$  prefix sums of the first step at each processor. Therefore,  $T_p = 2(n/p) + (1 + t_s + t_w) \log p$ .

b)

```
double pdf[N], cdf[N];
MPI_Init(&argc, &argv);
int rank, procs;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &procs);
int L=rank*N/procs, R=(rank+1)*N/procs;
for (int i=L; i<R; i++) {
    pdf[i] = rand();
}
cdf[L] = pdf[L];
for (int i=L+1; i<R; i++) {
    cdf[i] = cdf [i-1] + pdf[i];
}
// MPI_Barrier(MPI_COMM_WORLD); // optional here
double base;
MPI_Scan(&cdf[R-1], &base, 1, MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);
base -= cdf[R-1];
for (int i=L; i<R; i++) {
    cdf[i] += base;
}
for (int i=L; i<R; i++) {
    printf("process %d: cumulative_sum[%d] = %lf\n", rank, i, cdf[i]);
}
MPI_Finalize();
```

3.

- b) AllReduce
- c) AlltoAll (Personalized)
- d) Bcast

4. The loop is performing a series of reductions on the same data, with the result being placed onto processor 0, 1, ..., N, respectively. An MPI\_AllReduce would have the same effect and execute in  $\log P$  time.