

Programming Shared Address Space Platforms

CSE 531

Spring 2023

Mahmut Taylan Kandemir

Topic Overview

- Thread Basics
- The POSIX Thread API
- Synchronization Primitives in Pthreads
- Controlling Thread and Synchronization Attributes
- Composite Synchronization Constructs
- OpenMP: a Standard for Directive Based Parallel Programming

Overview of Programming Models

- Programming models provide support for expressing **concurrency** and **synchronization**.
- **Process-based models** assume that all data associated with a process is *private*, by default, unless otherwise specified.
- **Lightweight processes** and **threads** assume that all memory is *global*.
- **Directive-based programming models** extend the threaded model by facilitating creation and synchronization of threads.

Overview of Programming Models

- A **thread** is a single stream of control in the flow of a program. A program like:

```
for (row = 0; row < n; row++)  
    for (column = 0; column < n; column++)  
        c[row][column] =  
            dot_product( get_row(a, row),  
                        get_col(b, col));
```

can be transformed to:

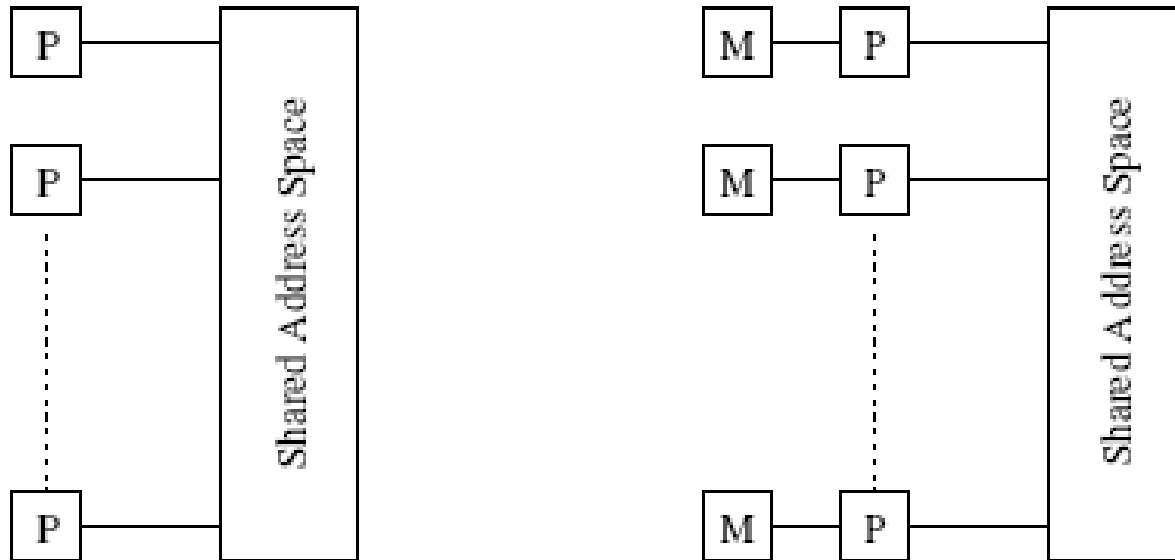
```
for (row = 0; row < n; row++)  
    for (column = 0; column < n; column++)  
        c[row][column] =  
            create_thread( dot_product(get_row(a, row),  
                                      get_col(b, col)));
```

In this case, one may think of the thread as an instance of a function that returns before the function has finished executing.

Thread Basics

- All memory in the logical machine model of a thread is globally accessible to every thread.
- The stack corresponding to the function call is generally treated as being local to the thread for liveness reasons.
 - Remember static memory, stack and heap?
- This implies a logical machine model with both global memory (default) and local memory (stacks).
- It is important to note that such a flat model may result in very poor performance since memory is physically distributed in typical machines.
 - Productivity vs Performance tradeoff

Thread Basics



- The logical machine model of a thread-based programming paradigm.
- In a distributed shared address space machines, the cost to access a physically local memory may be an order of magnitude less than that of accessing a remote memory.
- Caches also skew memory access times.

Thread Basics

- Threads provide software portability.
- Inherent support for latency hiding.
- Scheduling and load balancing.
- Ease of programming and widespread use.

The POSIX Thread API

- Commonly referred to as Pthreads, POSIX has emerged as the standard threads API, supported by most vendors.
- The concepts discussed here are largely independent of the API and can be used for programming with other thread APIs (NT threads, Solaris threads, Java threads, etc.) as well.

Thread Basics: Creation and Termination

- Pthreads provides two basic functions for specifying concurrency in a program:

```
#include <pthread.h>

int pthread_create (
    pthread_t *thread_handle, const pthread_attr_t
    *attribute,
    void * (*thread_function) (void *),
    void *arg);

int pthread_join (
    pthread_t thread,
    void **ptr);
```

- The function `pthread_create` invokes function `thread_function` as a thread
- A call to `pthread_join` waits for the termination of the thread whose id is given by `thread`

Thread Basics: Creation and Termination (Example)

```
#include <pthread.h>
#include <stdlib.h>
#define MAX_THREADS 512
void *compute_pi (void *);
....
main() {
    ...
    pthread_t p_threads[MAX_THREADS];
    pthread_attr_t attr;
    pthread_attr_init (&attr);
    for (i=0; i< num_threads; i++) {
        hits[i] = i;
        pthread_create(&p_threads[i], &attr, compute_pi,
            (void *) &hits[i]);
    }
    for (i=0; i< num_threads; i++) {
        pthread_join(p_threads[i], NULL);
        total_hits += hits[i];
    }
    ...
}
```

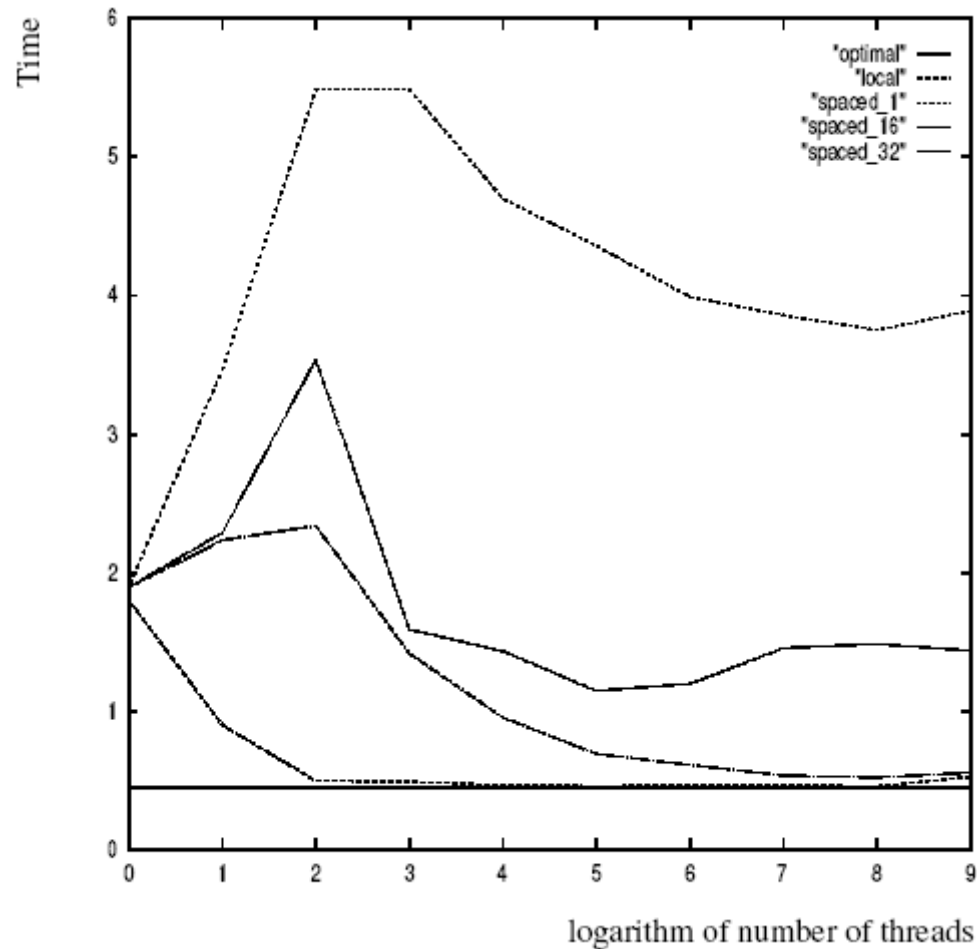
Thread Basics: Creation and Termination (Example)

```
void *compute_pi (void *s) {
    int seed, i, *hit_pointer;
    double rand_no_x, rand_no_y;
    int local_hits;
    hit_pointer = (int *) s;
    seed = *hit_pointer;
    local_hits = 0;
    for (i = 0; i < sample_points_per_thread; i++) {
        rand_no_x =(double)(rand_r(&seed))/(double)((2<<14)-1);
        rand_no_y =(double)(rand_r(&seed))/(double)((2<<14)-1);
        if (((rand_no_x - 0.5) * (rand_no_x - 0.5) +
            (rand_no_y - 0.5) * (rand_no_y - 0.5)) < 0.25)
            local_hits ++;
        seed *= i;
    }
    *hit_pointer = local_hits;
    pthread_exit(0);
}
```

Programming and Performance Notes

- Note the use of the function `rand_r` (instead of superior random number generators such as `drand48`).
- Executing this on a 4-processor machine, we observe a 3.91 fold speedup at 32 threads. This corresponds to a parallel efficiency of 0.98!
- We can also modify the program slightly to observe the effect of false-sharing.
- The program can also be used to assess the secondary cache line size.

Programming and Performance Notes



- Execution time of the `compute_pi` program.

Synchronization Primitives in Pthreads

- When multiple threads attempt to manipulate the same data item, the results can often be incoherent if proper care is not taken to synchronize them.
- Consider:

```
/* each thread tries to update variable best_cost as follows */  
if (my_cost < best_cost)  
    best_cost = my_cost;
```
- Assume that there are two threads, the initial value of `best_cost` is 100, and the values of `my_cost` are 50 and 75 at threads `t1` and `t2`.
- Depending on the schedule of the threads, the value of `best_cost` could be 50 or 75!
- The value 75 does not correspond to any serialization of the threads.

Mutual Exclusion

- The code in the previous example corresponds to a **critical segment**; i.e., a segment that must be executed by only *one thread at any time*.
- Critical segments in Pthreads are implemented using mutex locks.
- Mutex-locks have two states: *locked* and *unlocked*. At any point of time, only one thread can lock a mutex lock. A lock is an **atomic operation**.
- A thread entering a critical segment first tries to get a lock. It goes ahead when the lock is granted.

Mutual Exclusion

- The Pthreads API provides the following functions for handling mutex-locks:

```
int pthread_mutex_lock (  
    pthread_mutex_t *mutex_lock);  
int pthread_mutex_unlock (  
    pthread_mutex_t *mutex_lock);  
int pthread_mutex_init (  
    pthread_mutex_t *mutex_lock,  
    const pthread_mutexattr_t *lock_attr);
```


Mutual Exclusion

- We can now write our previously incorrect code segment as:

```
pthread_mutex_t minimum_value_lock;
...
main() {
    ....
    pthread_mutex_init(&minimum_value_lock, NULL);
    ....
}
void *find_min(void *list_ptr) {
    ....
    pthread_mutex_lock(&minimum_value_lock);
    if (my_min < minimum_value)
        minimum_value = my_min;
    /* and unlock the mutex */
    pthread_mutex_unlock(&minimum_value_lock);
}
```

Producer-Consumer Computation Using Locks

- A common use of mutex-locks is in establishing a **producer-consumer** relationship among threads.
- The producer-consumer scenario imposes the following constraints:
 - The producer thread must not overwrite the shared buffer when the previous task has not been picked up by a consumer thread.
 - The consumer threads must not pick up tasks until there is something present in the shared data structure.
 - Individual consumer threads should pick up tasks one at a time.

Producer-Consumer Using Locks

```
pthread_mutex_t task_queue_lock;
int task_available;
...
main() {
    ....
    task_available = 0;
    pthread_mutex_init(&task_queue_lock, NULL);
    ....
}
void *producer(void *producer_thread_data) {
    ....
    while (!done()) {
        inserted = 0;
        create_task(&my_task);
        while (inserted == 0) {
            pthread_mutex_lock(&task_queue_lock);
            if (task_available == 0) {
                insert_into_queue(my_task);
                task_available = 1;
                inserted = 1;
            }
            pthread_mutex_unlock(&task_queue_lock);
        }
    }
}
```

We use a variable called `task_available`: If this variable is 0, consumer threads must wait, but the producer thread can insert tasks into the shared data structure `task_queue`. If `task_available` is equal to 1, the producer thread must wait to insert the task into the shared data structure, but one of the consumer threads can pick up the task available.

Producer-Consumer Using Locks

```
void *consumer(void *consumer_thread_data) {
    int extracted;
    struct task my_task;
    /* local data structure declarations */
    while (!done()) {
        extracted = 0;
        while (extracted == 0) {
            pthread_mutex_lock(&task_queue_lock);
            if (task_available == 1) {
                extract_from_queue(&my_task);
                task_available = 0;
                extracted = 1;
            }
            pthread_mutex_unlock(&task_queue_lock);
        }
        process_task(my_task);
    }
}
```

Types of Mutexes

- Pthreads supports three *types* of mutexes – **normal**, **recursive**, and **error-check**.
- A **normal mutex** deadlocks if a thread that already has a lock tries a second lock on it.
- A **recursive mutex** allows a single thread to lock a mutex as many times as it wants. It simply increments a count on the number of locks. A lock is relinquished by a thread when the count becomes zero.
- An **error check mutex** reports an error when a thread with a lock tries to lock it again (as opposed to deadlocking in the first case, or granting the lock, as in the second case).
- The type of the mutex can be set in the **attributes object** before it is passed at time of initialization.

Overheads of Locking

- Locks represent **serialization points** since critical sections must be executed by threads *one after the other*.
- Encapsulating large segments of the program within locks can lead to significant *performance degradation*.
- It is often possible to reduce the idling overhead associated with locks using an alternate function, `pthread_mutex_trylock`.

```
int pthread_mutex_trylock (  
    pthread_mutex_t *mutex_lock);
```
- `pthread_mutex_trylock` is typically much faster than `pthread_mutex_lock` on typical systems since it does not have to deal with queues associated with locks for multiple threads waiting on the lock.

Alleviating Locking Overhead (Example)

```
/* Finding k matches in a list */
void *find_entries(void *start_pointer) {
    /* This is the thread function */
    struct database_record *next_record;
    int count;
    current_pointer = start_pointer;
    do {
        next_record = find_next_entry(current_pointer);
        count = output_record(next_record);
    } while (count < requested_number_of_records);
}

int output_record(struct database_record *record_ptr) {
    int count;
    pthread_mutex_lock(&output_count_lock);
    output_count++;
    count = output_count;
    pthread_mutex_unlock(&output_count_lock);
    if (count <= requested_number_of_records)
        print_record(record_ptr);
    return (count);
}
```

Alleviating Locking Overhead (Example)

```
/* rewritten output_record function */
int output_record(struct database_record
    *record_ptr) {
    int count;
    int lock_status;
    lock_status=pthread_mutex_trylock(&output_count_lock);
    if (lock_status == EBUSY) {
        insert_into_local_list(record_ptr);
        return(0);
    }
    else {
        count = output_count;
        output_count += number_on_local_list + 1;
        pthread_mutex_unlock(&output_count_lock);
        print_records(record_ptr, local_list,
            requested_number_of_records - count);
        return(count + number_on_local_list + 1);
    }
}
```


Condition Variables for Synchronization

- Indiscriminate use of locks can result in *idling overhead* from blocked threads.
- A **condition variable** allows a thread to block itself until specified data reaches a predefined state.
- A condition variable is associated with this predicate. When the predicate becomes true, the condition variable is used to signal one or more threads waiting on the condition.
- A single condition variable may be associated with more than one predicate.
- A condition variable always has a *mutex* associated with it. A thread locks this mutex and tests the predicate defined on the shared variable.
- If the predicate is not true, the thread waits on the condition variable associated with the predicate using the function `pthread_cond_wait`.

Condition Variables for Synchronization

- Pthreads provides the following functions for condition variables:

```
int pthread_cond_wait(pthread_cond_t *cond,  
    pthread_mutex_t *mutex);  
int pthread_cond_signal(pthread_cond_t *cond);  
int pthread_cond_broadcast(pthread_cond_t *cond);  
int pthread_cond_init(pthread_cond_t *cond,  
    const pthread_condattr_t *attr);  
int pthread_cond_destroy(pthread_cond_t *cond);
```

Producer-Consumer Using Condition Variables

```
pthread_cond_t cond_queue_empty, cond_queue_full;
pthread_mutex_t task_queue_cond_lock;
int task_available;
/* other data structures here */
main() {
    /* declarations and initializations */
    task_available = 0;
    pthread_init();
    pthread_cond_init(&cond_queue_empty, NULL);
    pthread_cond_init(&cond_queue_full, NULL);
    pthread_mutex_init(&task_queue_cond_lock, NULL);
    /* create and join producer and consumer threads */
}
```

Producer-Consumer Using Condition Variables

```
void *producer(void *producer_thread_data) {
    int inserted;
    while (!done()) {
        create_task();
        pthread_mutex_lock(&task_queue_cond_lock);
        while (task_available == 1)
            pthread_cond_wait(&cond_queue_empty,
                             &task_queue_cond_lock);
        insert_into_queue();
        task_available = 1;
        pthread_cond_signal(&cond_queue_full);
        pthread_mutex_unlock(&task_queue_cond_lock);
    }
}
```

Producer-Consumer Using Condition Variables

```
void *consumer(void *consumer_thread_data) {
    while (!done()) {
        pthread_mutex_lock(&task_queue_cond_lock);
        while (task_available == 0)
            pthread_cond_wait(&cond_queue_full,
                             &task_queue_cond_lock);
        my_task = extract_from_queue();
        task_available = 0;
        pthread_cond_signal(&cond_queue_empty);
        pthread_mutex_unlock(&task_queue_cond_lock);
        process_task(my_task);
    }
}
```

Controlling Thread and Synchronization Attributes

- The Pthreads API allows a programmer to change the default attributes of entities using **attributes objects**.
- An attributes object is a data-structure that describes entity (thread, mutex, condition variable) properties.
- Once these properties are set, the attributes object can be passed to the method initializing the entity.
- Enhances modularity, readability, and ease of modification.

Attributes Objects for Threads

- Use `pthread_attr_init` to create an attributes object.
- Individual properties associated with the attributes object can be changed using the following functions:

```
pthread_attr_setdetachstate,  
pthread_attr_setguardsize_np,  
pthread_attr_setstacksize,  
pthread_attr_setinheritsched,  
pthread_attr_setschedpolicy, and  
pthread_attr_setschedparam
```

Attributes Objects for Mutexes

- Initialize the attributes object using function:
`pthread_mutexattr_init.`
- The function `pthread_mutexattr_settype_np` can be used for setting the type of mutex specified by the mutex attributes object.

```
pthread_mutexattr_settype_np (  
pthread_mutexattr_t *attr,  
int type);
```
- Here, `type` specifies the type of the mutex and can take one of:
 - `PTHREAD_MUTEX_NORMAL_NP`
 - `PTHREAD_MUTEX_RECURSIVE_NP`
 - `PTHREAD_MUTEX_ERRORCHECK_NP`

Composite Synchronization Constructs

- By design, Pthreads provide support for a basic set of operations.
- Higher level constructs can be built using basic synchronization constructs.
- We discuss two such constructs:
 - Read-Write Locks and Barriers.

Read-Write Locks

- In many applications, a data structure is read frequently but written infrequently. For such applications, we should use **read-write locks**.
- A **read lock** is granted when there are other threads that may already have read locks.
- If there is a **write lock** on the data (or if there are queued write locks), the thread performs a **condition wait**.
- If there are multiple threads requesting a write lock, they must perform a condition wait.
- With this description, we can design functions for read locks `mylib_rwlock_rlock`, write locks `mylib_rwlock_wlock`, and unlocking `mylib_rwlock_unlock`.

Read-Write Locks

- The lock data type `mylib_rwlock_t` holds the following:
 - a count of the number of readers,
 - the writer (a 0/1 integer specifying whether a writer is present),
 - a condition variable `readers_proceed` that is signaled when readers can proceed,
 - a condition variable `writer_proceed` that is signaled when one of the writers can proceed,
 - a count `pending_writers` of pending writers, and
 - a mutex `read_write_lock` associated with the shared data structure

Read-Write Locks

```
typedef struct {
    int readers;
    int writer;
    pthread_cond_t readers_proceed;
    pthread_cond_t writer_proceed;
    int pending_writers;
    pthread_mutex_t read_write_lock;
} mylib_rwlock_t;

void mylib_rwlock_init (mylib_rwlock_t *l) {
    l -> readers = l -> writer = l -> pending_writers
        = 0;
    pthread_mutex_init(&(l -> read_write_lock),
        NULL);
    pthread_cond_init(&(l -> readers_proceed), NULL);
    pthread_cond_init(&(l -> writer_proceed), NULL);
}
```

Read-Write Locks

```
void mylib_rwlock_rlock(mylib_rwlock_t *l) {
    /* if there is a write lock or pending writers, perform
       condition wait.. else increment count of readers and grant
       read lock */
    pthread_mutex_lock(&(l -> read_write_lock));
    while ((l -> pending_writers > 0) || (l -> writer > 0))
        pthread_cond_wait(&(l -> readers_proceed),
            &(l -> read_write_lock));
    l -> readers ++;
    pthread_mutex_unlock(&(l -> read_write_lock));
}
```

Read-Write Locks

```
void mylib_rwlock_wlock(mylib_rwlock_t *l) {
    /* if there are readers or writers, increment pending
       writers count and wait. On being woken, decrement
       pending writers count and increment writer count */

    pthread_mutex_lock(&(l -> read_write_lock));
    while ((l -> writer > 0) || (l -> readers > 0)) {
        l -> pending_writers ++;
        pthread_cond_wait(&(l -> writer_proceed),
            &(l -> read_write_lock));
    }
    l -> pending_writers --;
    l -> writer ++;
    pthread_mutex_unlock(&(l -> read_write_lock));
}
```

Read-Write Locks

```
void mylib_rwlock_unlock(mylib_rwlock_t *l) {
/* if there is a write lock then unlock, else if there are
   read locks, decrement count of read locks. If the count
   is 0 and there is a pending writer, let it through, else
   if there are pending readers, let them all go through */
pthread_mutex_lock(&(l -> read_write_lock));
if (l -> writer > 0)
    l -> writer = 0;
else if (l -> readers > 0)
    l -> readers --;
pthread_mutex_unlock(&(l -> read_write_lock));
if ((l -> readers == 0) && (l -> pending_writers > 0))
    pthread_cond_signal(&(l -> writer_proceed));
else if (l -> readers > 0)
    pthread_cond_broadcast(&(l -> readers_proceed));
}
```

Barriers

- As in MPI (will be discussed later), a **barrier** holds a thread until all threads participating in the barrier have reached it.
- Barriers can be implemented using a counter, a mutex and a condition variable.
- A single integer is used to keep track of the number of threads that have reached the barrier.
- If the count is less than the total number of threads, the threads execute a condition wait.
- The *last thread* entering (and setting the count to the number of threads) *wakes up* all the threads using a condition broadcast.

Barriers

```
typedef struct {  
    pthread_mutex_t count_lock;  
    pthread_cond_t ok_to_proceed;  
    int count;  
} mylib_barrier_t;  
void mylib_init_barrier(mylib_barrier_t *b) {  
    b -> count = 0;  
    pthread_mutex_init(&(b -> count_lock), NULL);  
    pthread_cond_init(&(b -> ok_to_proceed), NULL);  
}
```

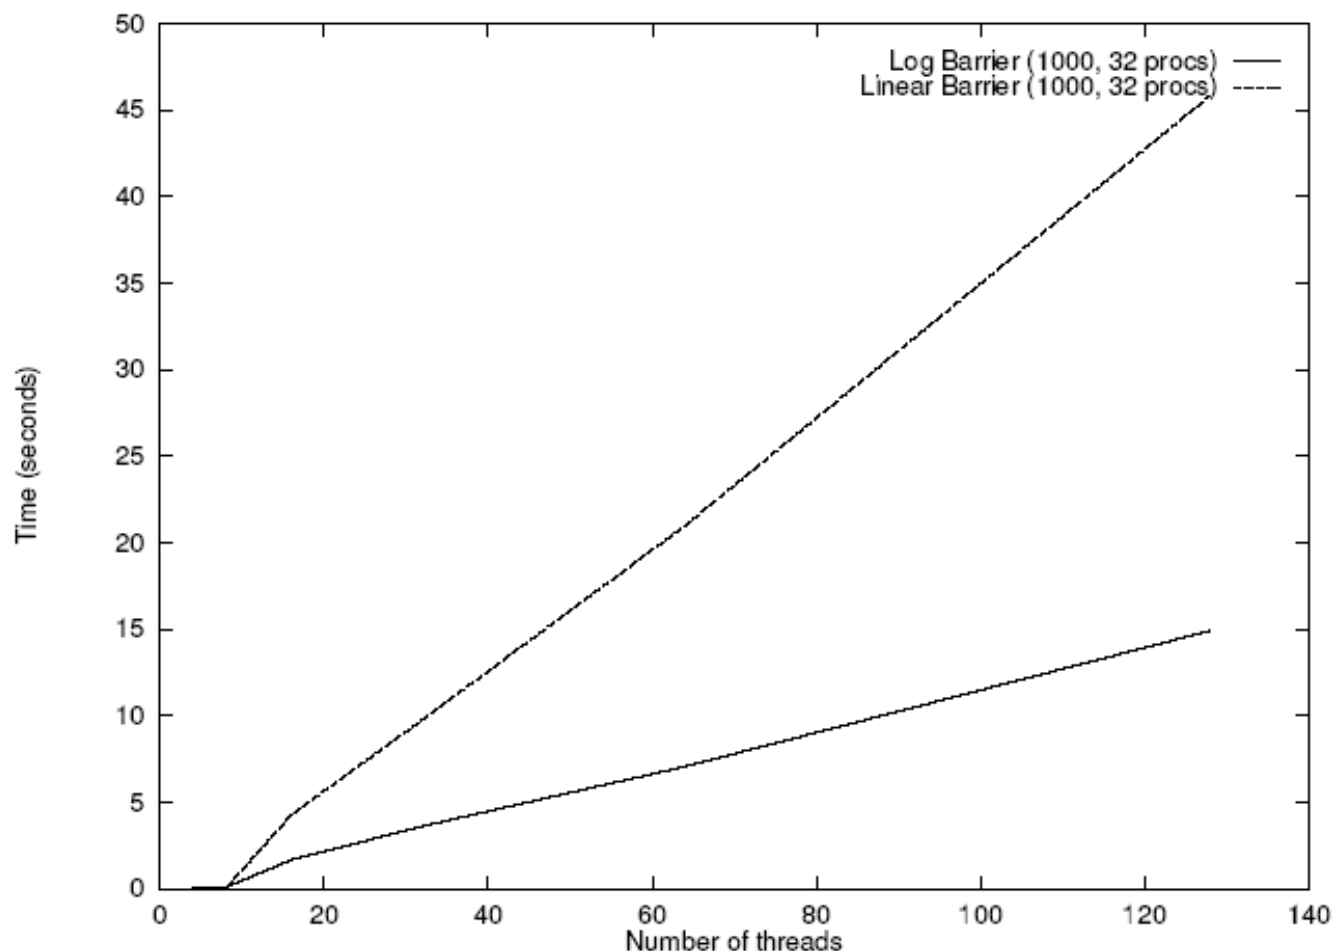
Barriers

```
void mylib_barrier (mylib_barrier_t *b, int
    num_threads) {
    pthread_mutex_lock(&(b -> count_lock));
    b -> count ++;
    if (b -> count == num_threads) {
        b -> count = 0;
        pthread_cond_broadcast(&(b -> ok_to_proceed));
    }
    else
        while (pthread_cond_wait(&(b -> ok_to_proceed),
            &(b -> count_lock)) != 0);
    pthread_mutex_unlock(&(b -> count_lock));
}
```

Barriers

- The barrier described above is called a **linear barrier**.
- The trivial lower bound on execution time of this function is therefore $O(n)$ for n threads.
- This implementation of a barrier can be sped up using multiple barrier variables organized in a *tree* – called **Combining Tree Barrier**
- In **k-tree barrier**, all threads are equally divided into subgroups of k threads and a first-round synchronizations are done within these subgroups.
- Once all subgroups have done their synchronizations, the first thread in each subgroup enters the second level for further synchronization.
- In the second level, like in the first level, the threads form new subgroups of k threads and synchronize within groups, sending out one thread in each subgroup to next level, and so on.
- Eventually, in the final level there is only one subgroup to be synchronized. After the final-level synchronization, the releasing signal is transmitted to upper levels and all threads get past the barrier.
- This is also called a **log barrier** and its runtime grows as $O(\log p)$.

Barrier



- Execution time of 1000 sequential and logarithmic barriers as a function of number of threads on a 32 processor SGI Origin 2000.

Tips for Designing Asynchronous Programs

- Set up all the requirements for a thread before actually creating it. This includes initializing the data, setting thread attributes, thread priorities, mutex attributes, etc
- Never rely on scheduling assumptions when exchanging data.
 - Threads are usually switched at semi-deterministic ways
- Do not rely on scheduling as a means of synchronization.
- Where possible, define and use group synchronizations and data replication. This can significantly improve program performance.

OpenMP Introduction

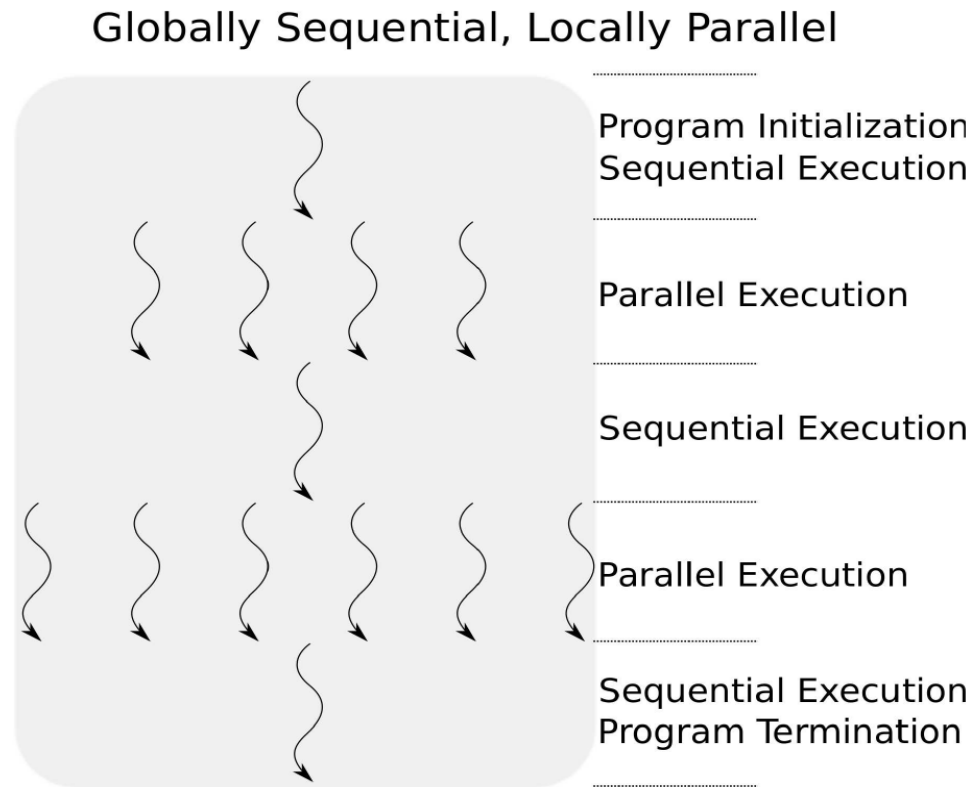
- The decomposition of a sequential program into components that can execute in parallel is a tedious enterprise.
- **OpenMP** has been designed to alleviate much of the effort involved, by accommodating the incremental conversion of sequential programs into parallel ones, with the assistance of the compiler.
- OpenMP relies on **compiler directives** for decorating portions of the code that the compiler will attempt to parallelize.

OpenMP History

- OpenMP : Open Multi-Processing is an API for shared-memory programming.
- OpenMP was specifically designed for parallelizing existing sequential programs.
- Uses compiler directives and a library of functions to support its operation.
- OpenMP v.1 was published in 1998.
- OpenMP v.5.2 was published in November 2021. Now OpenMP supports GPU accelerators.
- Standard controlled by the OpenMP Architecture Review Board (ARB).
- GNU Compiler support:
 - GCC 9 supports OpenMP 5.0
 - GCC 12 supports OpenMP 5.1

OpenMP Paradigm

- OpenMP programs are Globally Sequential, Locally Parallel.
- Programs follow the fork-join paradigm:



OpenMP Paradigm

- OpenMP is a directive-based API that can be used with FORTRAN, C, and C++ for programming *shared address space* machines.
- OpenMP directives provide support for concurrency, synchronization, and data handling while obviating the need for explicitly setting up mutexes, condition variables, data scope, and initialization.

OpenMP Programming Model

- OpenMP directives in C and C++ are based on the `#pragma` compiler directives.
- A directive consists of a directive name followed by clauses.

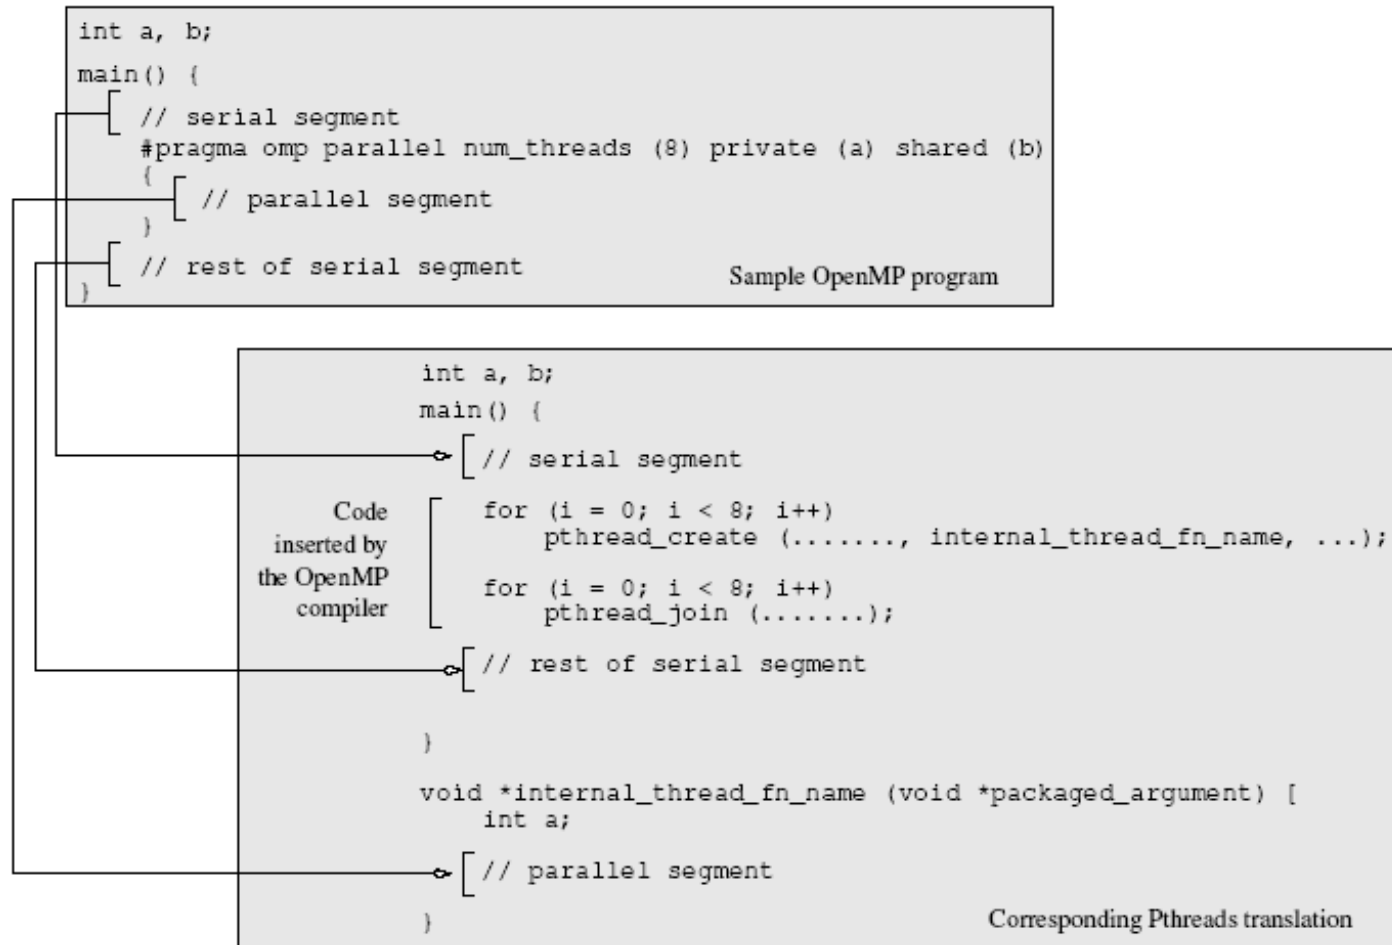
```
#pragma omp directive [clause list]
```
- OpenMP programs execute *serially* until they encounter the `parallel` directive, which creates a group of threads.

```
#pragma omp parallel [clause list]  
/* structured block */
```
- The main thread that encounters the `parallel` directive becomes the *master* of this group of threads and is assigned the thread id 0 within the group.

OpenMP Programming Model

- The clause list is used to specify conditional parallelization, number of threads, and data handling.
 - **Conditional Parallelization:** The clause `if (scalar expression)` determines whether the parallel construct results in creation of threads.
 - **Degree of Concurrency:** The clause `num_threads(integer expression)` specifies the number of threads that are created.
 - **Data Handling:** The clause `private (variable list)` indicates that variables local to each thread. The clause `firstprivate (variable list)` is similar to the `private`, except values of variables are initialized to corresponding values before the parallel directive. The clause `shared (variable list)` indicates that variables are shared across all the threads.

OpenMP Programming Model



- A sample OpenMP program along with its **Pthreads translation** that might be performed by an OpenMP compiler.

OpenMP Programming Model

```
#pragma omp parallel if (is_parallel== 1) num_threads(8) \  
    private (a) shared (b) firstprivate(c) {  
    /* structured block */  
}
```

- A **parallel region** is a block of code executed by multiple threads simultaneously

```
#pragma omp parallel [clause[[,] clause] ...]  
{  
    "this will be executed in parallel"  
} (implied barrier)
```

- If the value of the variable `is_parallel` equals one, eight threads are created.
- Each of these threads gets private copies of variables `a` and `c`, and shares a single value of variable `b`.
- The value of each copy of `c` is initialized to the value of `c` before the parallel directive.
- The default state of a variable is specified by the clause `default` (`shared`) or `default` (`none`).

Reduction Clause in OpenMP

- The **reduction clause** specifies how multiple local copies of a variable at different threads are combined into a single copy at the master when threads exit.
- The usage of the `reduction clause` is `reduction (operator: variable list)`.
- The variables in the list are implicitly specified as being private to threads.
- The operator can be one of `+`, `*`, `-`, `&`, `|`, `^`, `&&`, and `||`.

```
#pragma omp parallel reduction(+: sum) num_threads(8) {  
/* compute local sums here */  
}  
/*sum here contains sum of all local instances of sums */
```

OpenMP Programming: Example

```
main ()
{
    int i, n, chunk;
    float x[100], y[100], result;
    /* Some initializations */
    n = 100; chunk = 10; result = 0.0;
    for (i=0; i < n; i++)
        { x[i] = ... ; y[i] = ...; }
    #pragma omp parallel for \
default(shared) private(i) \
schedule(static,chunk) \
reduction(+:result)
    for (i=0; i < n; i++)
        result = result + (x[i] * y[i]);
    printf("Final result= %f\n",result);
}
```

OpenMP Programming: Example

```
/* *****  
An OpenMP version of a threaded program to compute PI.  
***** */  
#pragma omp parallel default(private) shared (npoints) \  
  reduction(+: sum) num_threads(8)  
{  
    num_threads = omp_get_num_threads();  
    sample_points_per_thread = npoints / num_threads;  
    sum = 0;  
    for (i = 0; i < sample_points_per_thread; i++) {  
        rand_no_x = (double) (rand_r(&seed)) / (double) ((2<<14)-1);  
        rand_no_y = (double) (rand_r(&seed)) / (double) ((2<<14)-1);  
        if (((rand_no_x - 0.5) * (rand_no_x - 0.5) +  
            (rand_no_y - 0.5) * (rand_no_y - 0.5)) < 0.25)  
            sum ++;  
    }  
}
```


Specifying Concurrent Tasks in OpenMP

- The `parallel` directive can be used in conjunction with other directives to specify *concurrency across iterations and tasks*.
- OpenMP provides two directives – `for` and `sections` – to specify *concurrent* iterations and tasks.
- The `for` directive is used to *split* parallel iteration spaces across threads. The general form of a `for` directive is as follows:

```
#pragma omp for [clause list]
/* for loop */
```

- The clauses that can be used in this context are: `private`, `firstprivate`, `lastprivate`, `reduction`, `schedule`, `nowait`, and `ordered`.

Specifying Concurrent Tasks in OpenMP: Example

```
#pragma omp parallel default(private) shared (npoints) \  
    reduction(+: sum) num_threads(8)  
{  
    sum = 0;  
    #pragma omp for  
    for (i = 0; i < npoints; i++) {  
        rand_no_x =(double) (rand_r(&seed)) / (double) ((2<<14)-1);  
        rand_no_y =(double) (rand_r(&seed)) / (double) ((2<<14)-1);  
        if (((rand_no_x - 0.5) * (rand_no_x - 0.5) +  
            (rand_no_y - 0.5) * (rand_no_y - 0.5)) < 0.25)  
            sum ++;  
    }  
}
```

Assigning Iterations to Threads

- The `schedule` clause of the `for` directive deals with the assignment of iterations to threads.
- The general form of the `schedule` directive is `schedule(scheduling_class[, parameter])`.
- OpenMP supports four scheduling classes: `static`, `dynamic`, `guided`, and `runtime`.

omp parallel and omp for

```
#pragma omp parallel
{
    #pragma omp for
    for(int i = 1; i < 100; ++i)
        { ... }
}
```

```
#pragma omp parallel for
{
    for(int i = 1; i < 100; ++i)
        { ... }
}
```

`#pragma omp parallel` spawns a group of threads, while `#pragma omp for` divides loop iterations between the spawned threads. You can do both things at once with the *fused* `#pragma omp parallel for` directive.

Controlling the Schedule

- By the schedule clause `schedule`. Syntax:
`#pragma omp parallel for schedule(`
 `static | dynamic |`
 `guided | auto | runtime`
 `[, chunk_size])`
- The runtime option delegates the scheduling decision for the execution of the program, where a previous setting (e.g., via `OMP_SCHEDULE`) can be inspected for suggestions. This is exclusive to the `schedule` clause only.


How to Select a Schedule Option

- **static**: If iterations are *homogeneous*
- **dynamic**: If execution cost varies
- **guided**: if execution cost varies and the number of iteration groups is too high, or the per-iteration cost increases as the loop progresses.
 - The singularity of guided is that the chunk size starts as large and gradually decreases, in order to progressively improve load balancing. These decreasing chunk sizes will eventually fall to 1 by default, unless a minimum chunk size is passed via the optional argument (except the last chunk to assign, which may have fewer iterations that asked).
 - a.k.a. **tapering**
- **runtime**: decision is made at runtime

Assigning Iterations to Threads: Example

```
/* static scheduling of matrix multiplication loops */  
#pragma omp parallel default(private) shared (a, b, c, dim) \  
  num_threads(4)  
  #pragma omp for schedule(static)  
  for (i = 0; i < dim; i++) {  
    for (j = 0; j < dim; j++) {  
      c(i,j) = 0;  
      for (k = 0; k < dim; k++) {  
        c(i,j) += a(i, k) * b(k, j);  
      }  
    }  
  }  
}
```

Assigning Iterations to Threads: Example



- Three different schedules using the static scheduling class of OpenMP.

Parallel For Loops

- Often, it is desirable to have a sequence of `for`-directives within a parallel construct that do not execute an implicit barrier at the end of each `for` directive.
- OpenMP provides a clause - `nowait`, which can be used with a `for` directive.

Parallel For Loops: Example

```
#pragma omp parallel
{
    #pragma omp for nowait
        for (i = 0; i < nmax; i++)
            if (isEqual(name, current_list[i])
                processCurrentName(name);
    #pragma omp for
        for (i = 0; i < mmax; i++)
            if (isEqual(name, past_list[i])
                processPastName(name);
}
```

Task Parallelism

- The `sections` directive can be used to setup individual work items that will be executed by threads. Their relative order of execution (or by which thread it is done) is *unknown*.

```
#pragma omp parallel
{
    ...
    #pragma omp sections
    ...
}

// OR

#pragma omp parallel sections
{
    ...
}
```

The section/sections Directives

- The individual work items are contained in blocks decorated by `section` directives:

```
#pragma omp parallel sections
{
  #pragma omp section
  {
    // concurrent block 0
  }
  ...
  #pragma omp section
  {
    // concurrent block M-1
  }
}
```

The sections Directive: Example

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        {
            taskA();
        }
        #pragma omp section
        {
            taskB();
        }
        #pragma omp section
        {
            taskC();
        }
    }
}
```

Nesting `parallel` Directives

- Nested parallelism can be enabled using the `OMP_NESTED` environment variable.
- If the `OMP_NESTED` environment variable is set to `TRUE`, nested parallelism is enabled.
- In this case, each parallel directive creates a new team of threads.

Synchronization Constructs in OpenMP

- OpenMP provides a variety of synchronization constructs:

```
#pragma omp barrier
```

```
#pragma omp single [clause list]  
    structured block
```

```
#pragma omp master  
    structured block
```

```
#pragma omp critical [(name)]  
    structured block
```

```
#pragma omp ordered  
    structured block
```

Synchronization Constructs

- **critical** : allows only one thread at a time, to enter the structured block that follows. The syntax involves an optional identifier:

```
#pragma omp critical [ ( identifier ) ]  
{  
    // structured block  
}
```

- The identifier allows the establishment of **named critical sections**. All critical directives without an identifier are assumed to have the same name, and use the same mutex.
- **atomic** : this is a lightweight version of the critical construct. Only a single statement (not a block) can follow.

OpenMP Library Functions

- In addition to directives, OpenMP also supports a number of functions that allow a programmer to control the execution of threaded programs.

```
/* thread and processor count */  
void omp_set_num_threads (int  
    num_threads);  
int omp_get_num_threads ();  
int omp_get_max_threads ();  
int omp_get_thread_num ();  
int omp_get_num_procs ();  
int omp_in_parallel();
```

OpenMP Library Functions

```
/* controlling and monitoring thread creation */  
void omp_set_dynamic (int dynamic_threads);  
int omp_get_dynamic ();  
void omp_set_nested (int nested);  
int omp_get_nested ();  
/* mutual exclusion */  
void omp_init_lock (omp_lock_t *lock);  
void omp_destroy_lock (omp_lock_t *lock);  
void omp_set_lock (omp_lock_t *lock);  
void omp_unset_lock (omp_lock_t *lock);  
int omp_test_lock (omp_lock_t *lock);
```

- In addition, all lock routines also have a nested lock counterpart
- for recursive mutexes.

Environment Variables in OpenMP

- `OMP_NUM_THREADS`: This environment variable specifies the default number of threads created upon entering a parallel region.
- `OMP_SET_DYNAMIC`: Determines if the number of threads can be dynamically changed.
- `OMP_NESTED`: Turns on nested parallelism.
- `OMP_SCHEDULE`: Scheduling of for-loops if the clause specifies runtime

Explicit Threads versus Directive Based Programming

- Directives layered on top of threads facilitate a variety of thread-related tasks.
- A programmer is rid of the tasks of initializing attributes objects, setting up arguments to threads, partitioning iteration spaces, etc.
- There are some drawbacks to using directives as well.
- An artifact of explicit threading is that data exchange is more apparent. This helps in alleviating some of the overheads from data movement, false sharing, and contention.
- Explicit threading also provides a richer API in the form of condition waits, locks of different types, and increased flexibility for building composite synchronization operations.
- Finally, since explicit threading is used more widely than OpenMP, tools and support for Pthreads programs are easier to find (NOTE: This may not be true soon...)