

# Midterm 1 Review

Scott Cheng

Jamboard Link:


<https://jamboard.google.com/d/1fXiHXHQgxWwHwdqLH8SBhlsYSnphPKfS88Ed11Njmzg/edit?usp=sharing>

# Midterm 1

- slide-group-1 ~ 7
- (openmp-supp, mpi-supp)

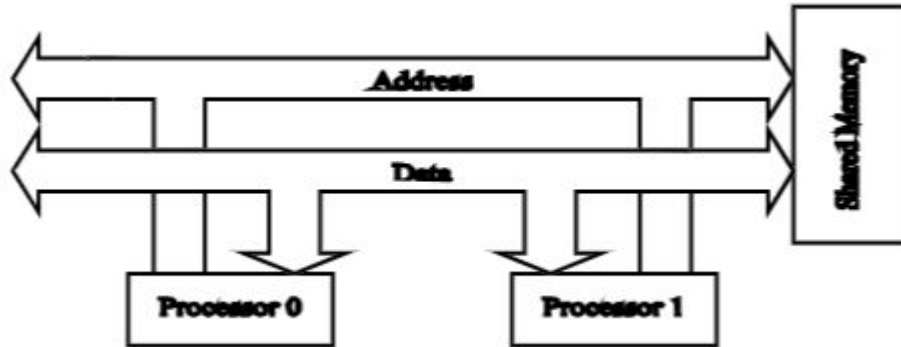
# Terms related to Interconnection Networks

- Distance between two nodes
- Diameter
- Bisection Width
- Arc Connectivity




# Network Topologies: Buses

- What's the distance between any two nodes?
- Bandwidth of the shared bus is a major bottleneck
  - => How to improve?



# Network Topologies: Crossbars

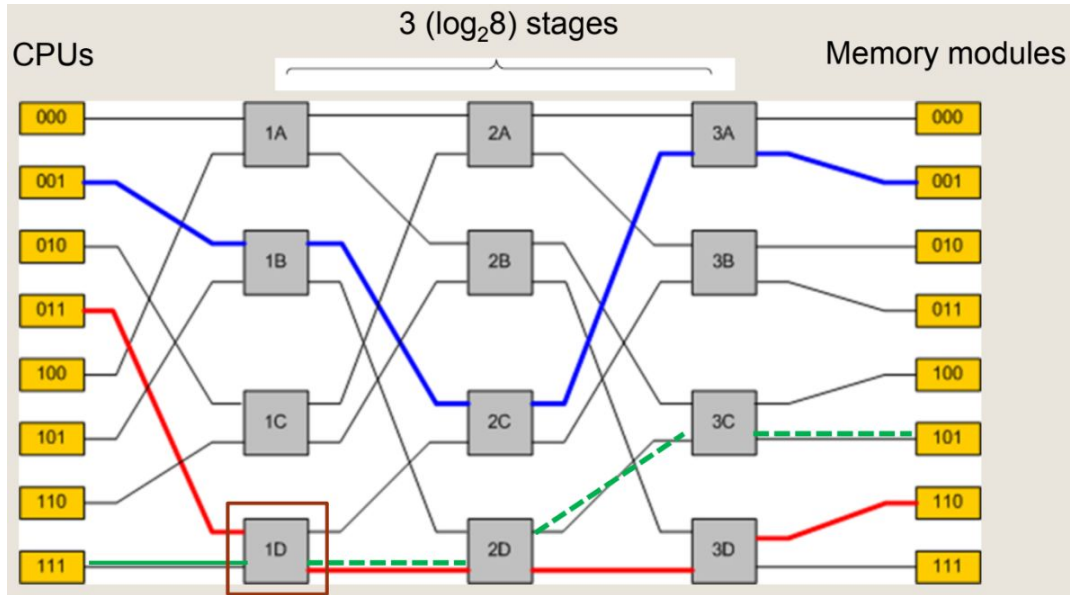
- The cost of a crossbar of  $p$  processors grows as  $\_\_\_O(?)\_\_\_$
- Blocking or non-blocking?



# Network Topologies: Multistage Omega Networks

- The cost of an omega network of  $p$  processors grows as  $\Theta(\log^2 p)$
- Blocking or non-blocking?


Example: an 8x8 Omega Network (with 2x2 crossbars)




- $\log_2(p) = 3$  stages
- each stage has 4 crossbars
- total number of crossbars? 12
- total number of links? 24

# In general, multistage networks are blocking

Can you construct a non-blocking multistage network?



Crossbar is non-blocking



Omega is blocking

Butterfly?

Network	Diameter	Bisection Width	Arc Connectivity	Cost (No. of links)
Completely-connected	1	$p^2/4$	$p - 1$	$p(p - 1)/2$
Star	2	1	1	$p - 1$
Complete binary tree	$2 \log((p + 1)/2)$	1	1	$p - 1$
Linear array	$p - 1$	1	1	$p - 1$
2-D mesh, no wraparound	$2(\sqrt{p} - 1)$	$\sqrt{p}$	2	$2(p - \sqrt{p})$
2-D wraparound mesh	$2\lfloor \sqrt{p}/2 \rfloor$	$2\sqrt{p}$	4	$2p$
Hypercube	$\log p$	$p/2$	$\log p$	$(p \log p)/2$
Crossbar	1	$p$	1	$p^2$
Omega Network	$\log p$	$p/2$	2	$p/2$
Dynamic Tree	$2 \log p$	1	2	$p - 1$



# Performance Metrics for Parallel Systems: Total Parallel Overhead

- Let  $T_{all}$  be the total time collectively spent by all the processing elements.
- $T_s$  is the serial time.

Speedup:

$$T_p = ?$$

$$S = ?$$

The overhead function:

$$T_o = ?$$

Efficiency:

$$E = ?$$

Cost Optimality?

# Isoefficiency function

Problem size  $W$  is defined as the asymptotic number of operations associated with the best serial algorithm to solve the problem.

$$W = KT_o(W, p)$$

# Example: Sorting

Consider a sorting algorithm that uses  $n$  processing elements to sort the list in time  $(\log n)^2$

- Since the serial runtime of a (comparison-based) sort is  $n \log n$ , the speedup and efficiency of this algorithm are given by  $\frac{n \log n}{n (\log n)^2}$  and  $\frac{n \log n}{n (\log n)^2}$  respectively.
- The  $p T_p$  product of this algorithm is  $n (\log n)^2$

Cost Optimal?

## Example: Adding n numbers

Consider the minimum execution time for adding  $n$  numbers.

$$T_P = \frac{n}{p} + 2 \log p.$$

$$t_s + t_w W$$


$$T_0 = ?$$

$$T_P^{\min} = ?$$

Isoefficiency function?

$$T_P^{\text{cost\_optimal}} = ?$$

# Reduction



**Examples:** averaging of Monte Carlo samples; convergence testing; image comparison metrics; matrix operations.

## Example: Adding N numbers

- *Reduction* combines every element in a collection into one element using an *associative* operator.

```
b = 0;  
for (i=0; i<n; ++i) {  
    b += f(B[i]);  
}
```

- Reordering of the operations is often needed to allow for parallelism.
- A tree reordering requires associativity.


# Scan

## Example: Prefix sum

- *Scan* computes all partial reductions of a collection

```
A[0] = B[0] + init;  
for (i=1; i<n; ++i) {  
    A[i] = B[i] + A[i-1];  
}
```

- Operator must be (at least) associative.
- Diagram shows one possible parallel implementation using three-phase strategy
- We'll consider different implementations later



**Examples:** random number generation,  
pack, tabulated integration, time series  
analysis

## Example: FFT



The parallel runtime of a parallel implementation of the FFT algorithm with  $p$  processing elements is given by  $T_p = (n/p) \log n + 10(n/p) \log p$  for an input sequence of length  $n$ .

Cost Optimal?

$$T_p^{\min} = ?$$

# Example: Matrix Multiplication (Rowwise 1D Partitioning)


$$T_P = \frac{n^2}{p} + t_s \log p + t_w n$$






# Example: Matrix Multiplication (2D Partitioning)


$$T_P \approx \frac{n^2}{p} + t_s \log p + t_w \frac{n}{\sqrt{p}} \log p$$




(a) Initial data distribution and communication steps to align the vector along the diagonal



(b) One-to-all broadcast of portions of the vector along process columns



(c) All-to-one reduction of partial results




(d) Final distribution of the result vector

# Communication Patterns

- One-to-All Broadcast and All-to-One Reduction
- All-to-All Broadcast and Reduction
- All-Reduce and Prefix Sum
- Scatter and Gather
- All-to-All Personalized Communication
- Circular Shift


# One-to-All Broadcast and All-to-One Reduction

- Ring: recursive doubling




$$T = (t_s + t_w m) \log p$$

- Mesh/Hypercube



- Balanced binary tree



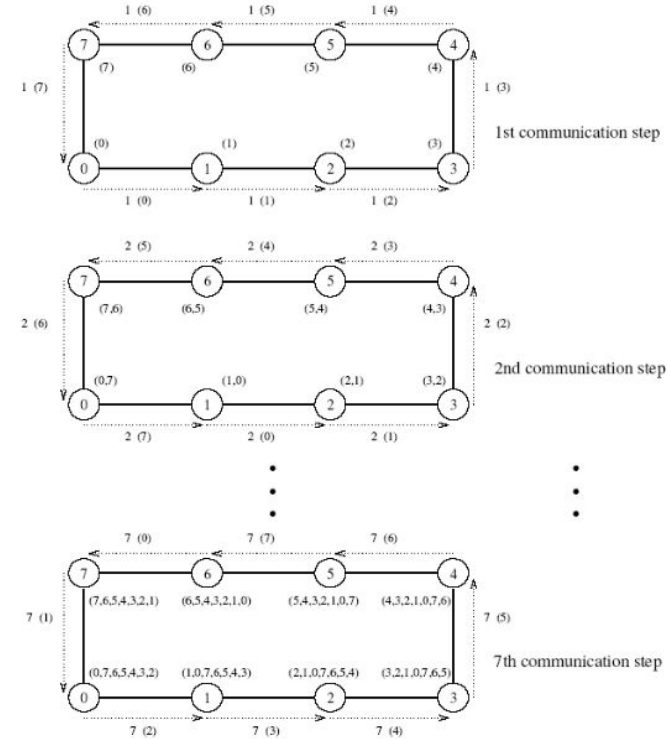
# All-to-All Broadcast and Reduction

- On a ring, the time is given by:  $(t_s + t_w m)(p-1)$ .
- On a mesh, the time is given by:  $2t_s(\sqrt{p} - 1) + t_w m(p-1)$ .
- On a hypercube, we have:

$$T = \sum_{i=1}^{\log p} (t_s + 2^{i-1} t_w m)$$

$$= t_s \log p + t_w m(p - 1).$$

- **Balanced binary tree**  $(t_s + t_w mp/2) \log p$



# Embedding


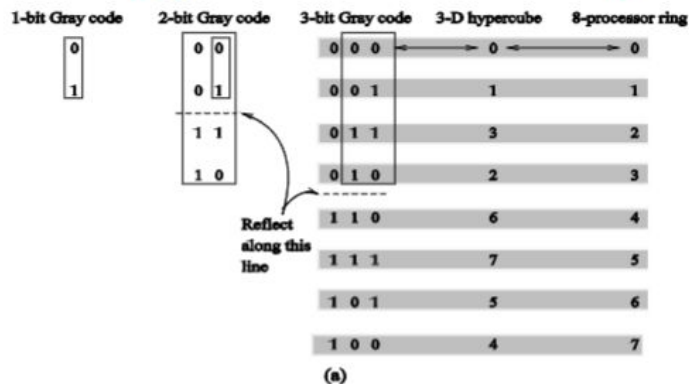
$$G(V, E) \rightarrow G'(V', E')$$

- Congestion?
- Dilation?
- Expansion?

# Embedding

- Embedding a ring into a hypercube
  - binary reflected gray code (RGC)
- Embedding a mesh into a hypercube
  - Concatenated binary RGC
- Embedding a mesh into a ring
- Embedding a hypercube into a mesh?
- Embedding a complete binary tree into a hypercube?
- Embedding a mesh of trees into a hypercube?
- Embedding a ring into a complete binary tree?



# Embedding a Linear Array into a Hypercube: Example



(a) A three-bit reflected Gray code ring; and (b) its embedding into a three-dimensional hypercube.

- Congestion?
- Dilation?

# Embedding a Mesh into a Hypercube




- Congestion?
- Dilation?


(a) A  $4 \times 4$  mesh illustrating the mapping of mesh nodes to the nodes in a four-dimensional hypercube; and (b) a  $2 \times 4$  mesh embedded into a three-dimensional hypercube.



# Embedding a Mesh into a Linear Array: Example



(a) Mapping a linear array into a 2D mesh (congestion 1).



(b) Inverting the mapping - mapping a 2D mesh into a linear array (congestion 5)

- Congestion?
- Dilation?

(a) Embedding a 16 node linear array into a 2-D mesh; and (b) the inverse of the mapping. Solid lines correspond to links in the linear array and normal lines to links in the mesh.