

CSE 541: Database Systems I



Hash Tables

Recap: Indexing

User queries and transactions:

UPDATE(Key, record...)

SELECT(Key...)



Alternatives for index entries:

- Key + actual data record
- Key + RID
- Key + list of RIDs

➔ Use a key to find data entry

Can be implemented using different index structures

- Hash tables
 - Equality search only
- Trees and skip lists
 - Equality + range search
 - "Order preserving"

Hash Tables

Hash table: a collection of buckets, each bucket holds keys (and associated payload) that are hashed into the same bucket using a hash function H

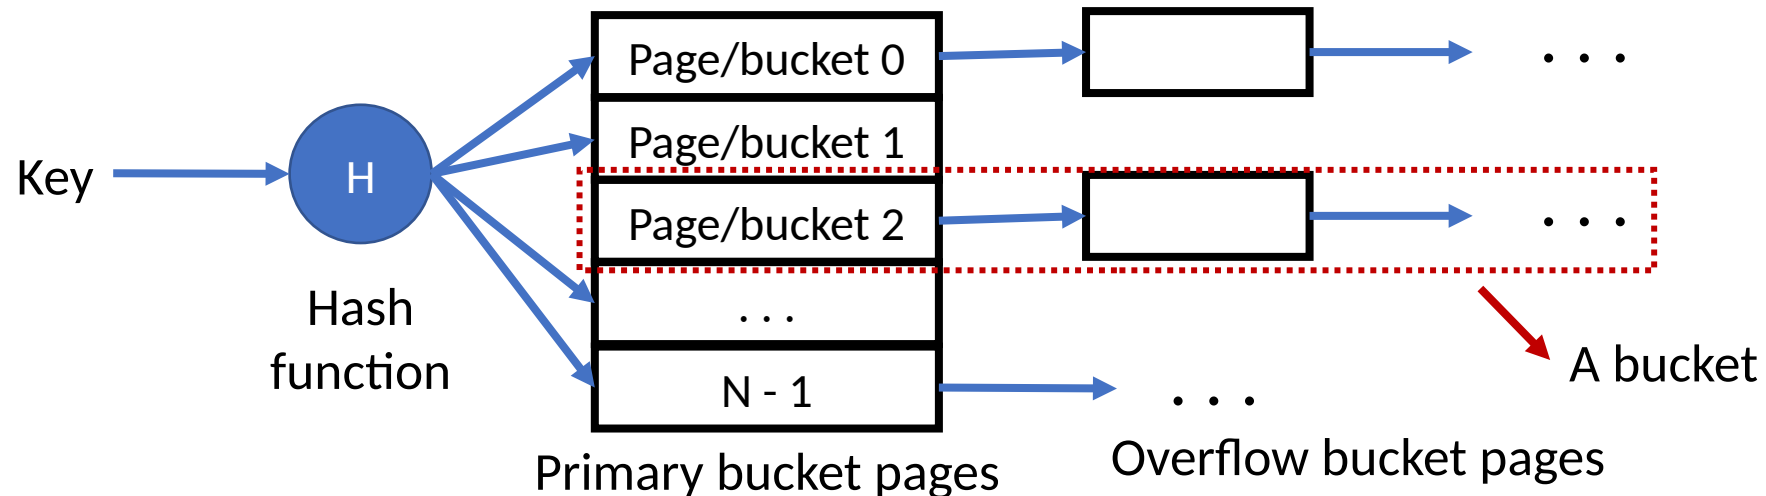
- Hash function maps key to an index into bucket number
 - Collisions: for different keys $K1$ and $K2$, $H(K1) == H(K2)$
 - So each bucket may contain multiple data entries
- Good for point queries
 - Arguably the fastest option
- Typically cannot do range scans
 - Order is usually not preserved by hash functions
 - Possible to have: $K1 < K2$ but $H(K1) > H(K2)$

Techniques



- Static hashing
- Dynamic: extendible hashing and linear hashing

Static Hashing

- Fixed number of buckets
- Each bucket consists of one or more pages
- Each bucket page has fixed capacity
 - No space in single page: create overflow pages chained together per bucket



Static Hashing

- Hash function must evenly distribute keys into different buckets
 - Usually works well: $H(\text{key}) = (A * \text{key} + B)$  A and B are constants that need to be tuned
 - Bucket number = $H(\text{key}) \bmod N$  N: Hash table size

Main problem: the number of buckets is fixed

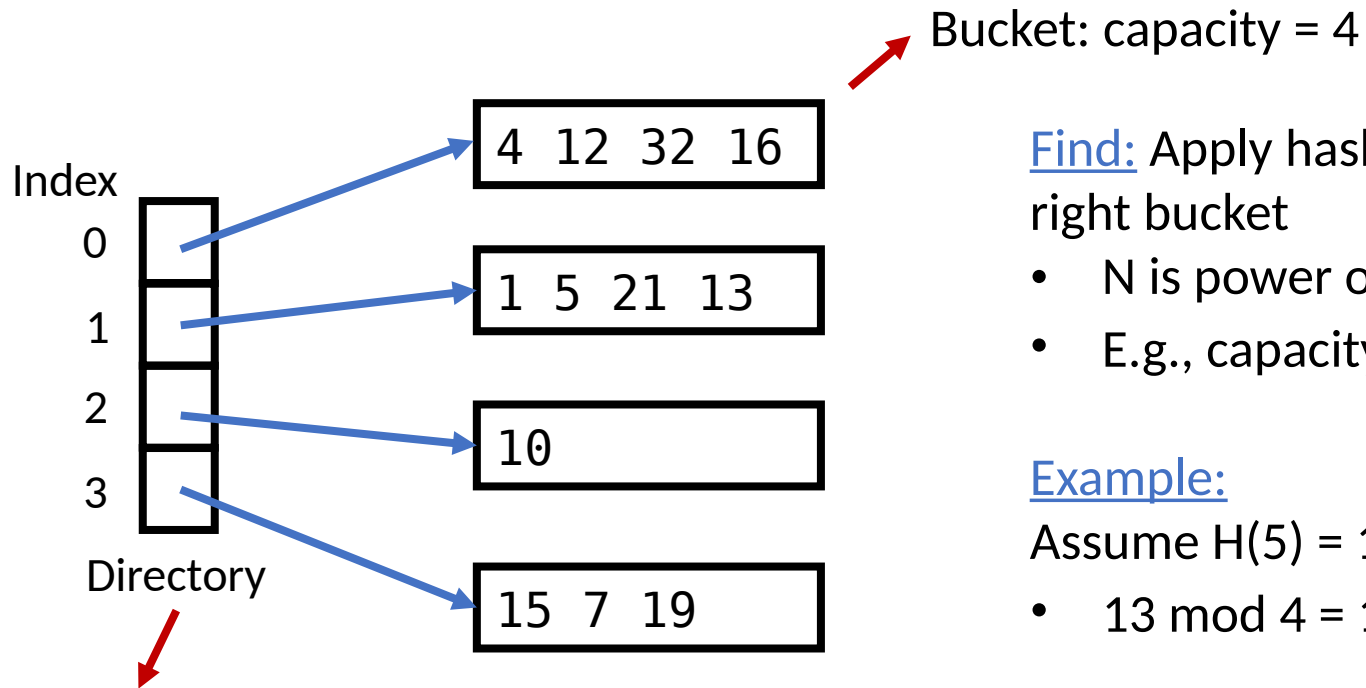
- Chains of overflow pages can become long
 - Affect search performance
- Waste space when data size reduced significantly
 - Empty buckets
- High cost to re-organize/expand by increasing the number of buckets: too much data copying

Extensible Hashing

Use a level of indirection to easily add more buckets

- Add a directory of pointers to buckets
 - Much smaller than actual bucket pages, cheap to resize/expand
- Double the directory when a bucket is full
 - Only split buckets that are full
 - No overflow pages

Extensible Hashing



Array with four elements

- Each element is a pointer to a bucket page (page offset in file)

Find: Apply hash function, modulo directory size (N) to find the right bucket

- N is power of 2, mod N **same as taking the last $\log_2 N$ bits**
- E.g., capacity = 4, take 2 bits

Example:

Assume $H(5) = 13 = 11\mathbf{01}_2$

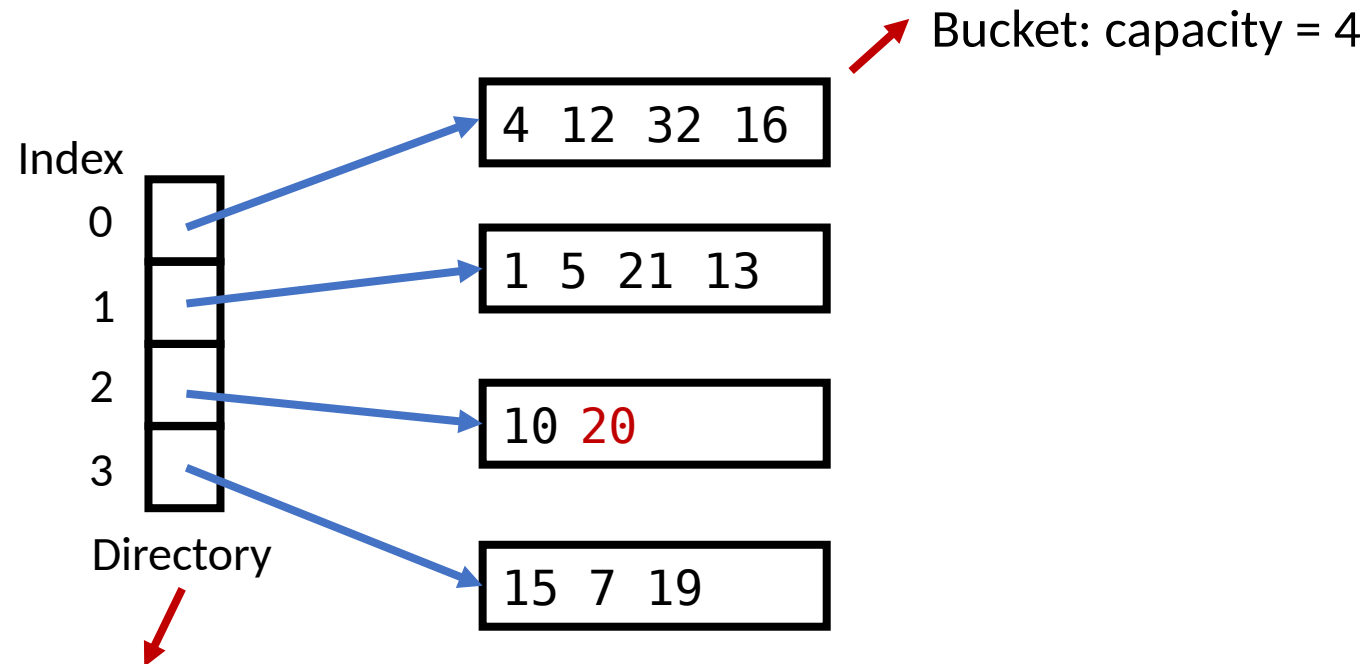
- $13 \bmod 4 = 1$

Insert: Use Find() to arrive at the right bucket first, then add new data entry

- Done if there is enough space
- Split if the bucket is already full

Extensible Hashing

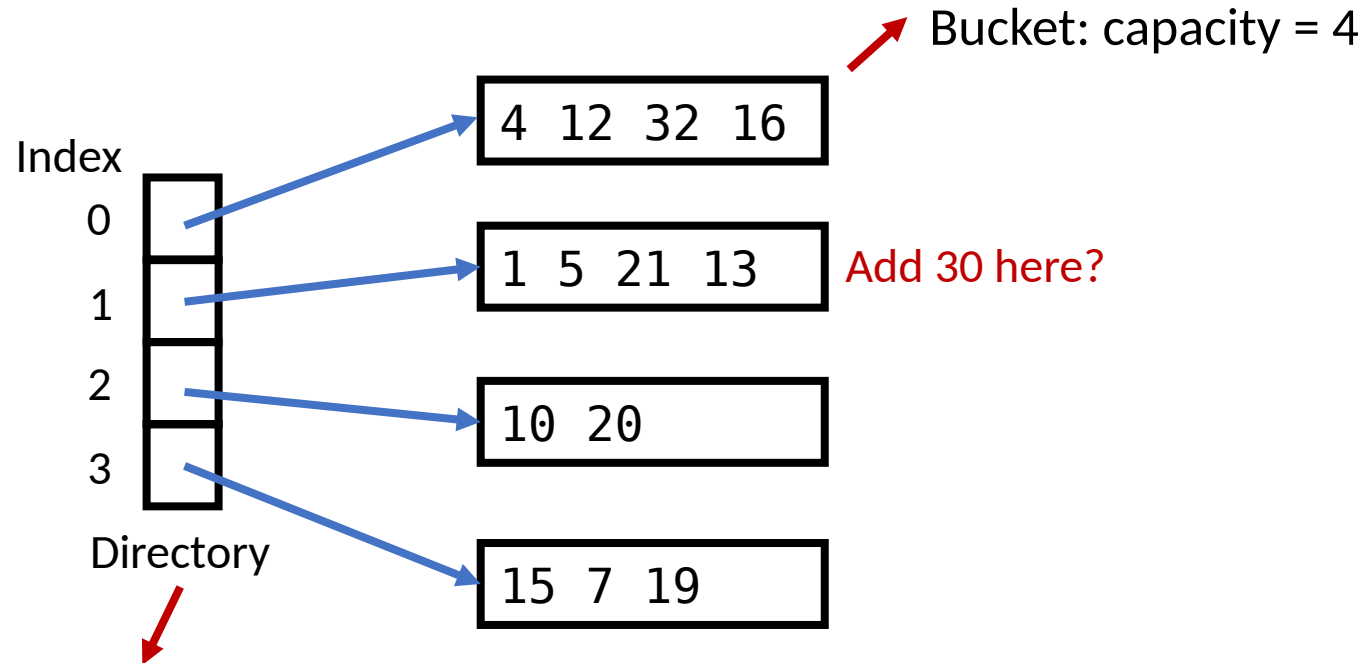
- Insert 20: assume $H(20) \bmod 4 = 2$, bucket 2 has enough space



Array with four elements, each is a pointer to a bucket page

Extensible Hashing

- Insert 30: assume $H(30) \bmod 4 = 1$, bucket 1 is already full



Array with four elements, each is a pointer to a bucket page


Extensible Hashing

Insert 30: assume $H(30) \bmod 4 = 1$, bucket 1 is already full

- Double the directory size to 8

Corresponding elements:

- Buckets 0 and 4
 - Buckets 1 and 5
 - Buckets 2 and 6
 - Buckets 3 and 7
-
- Old vs. new buckets: differ in the number of bits used

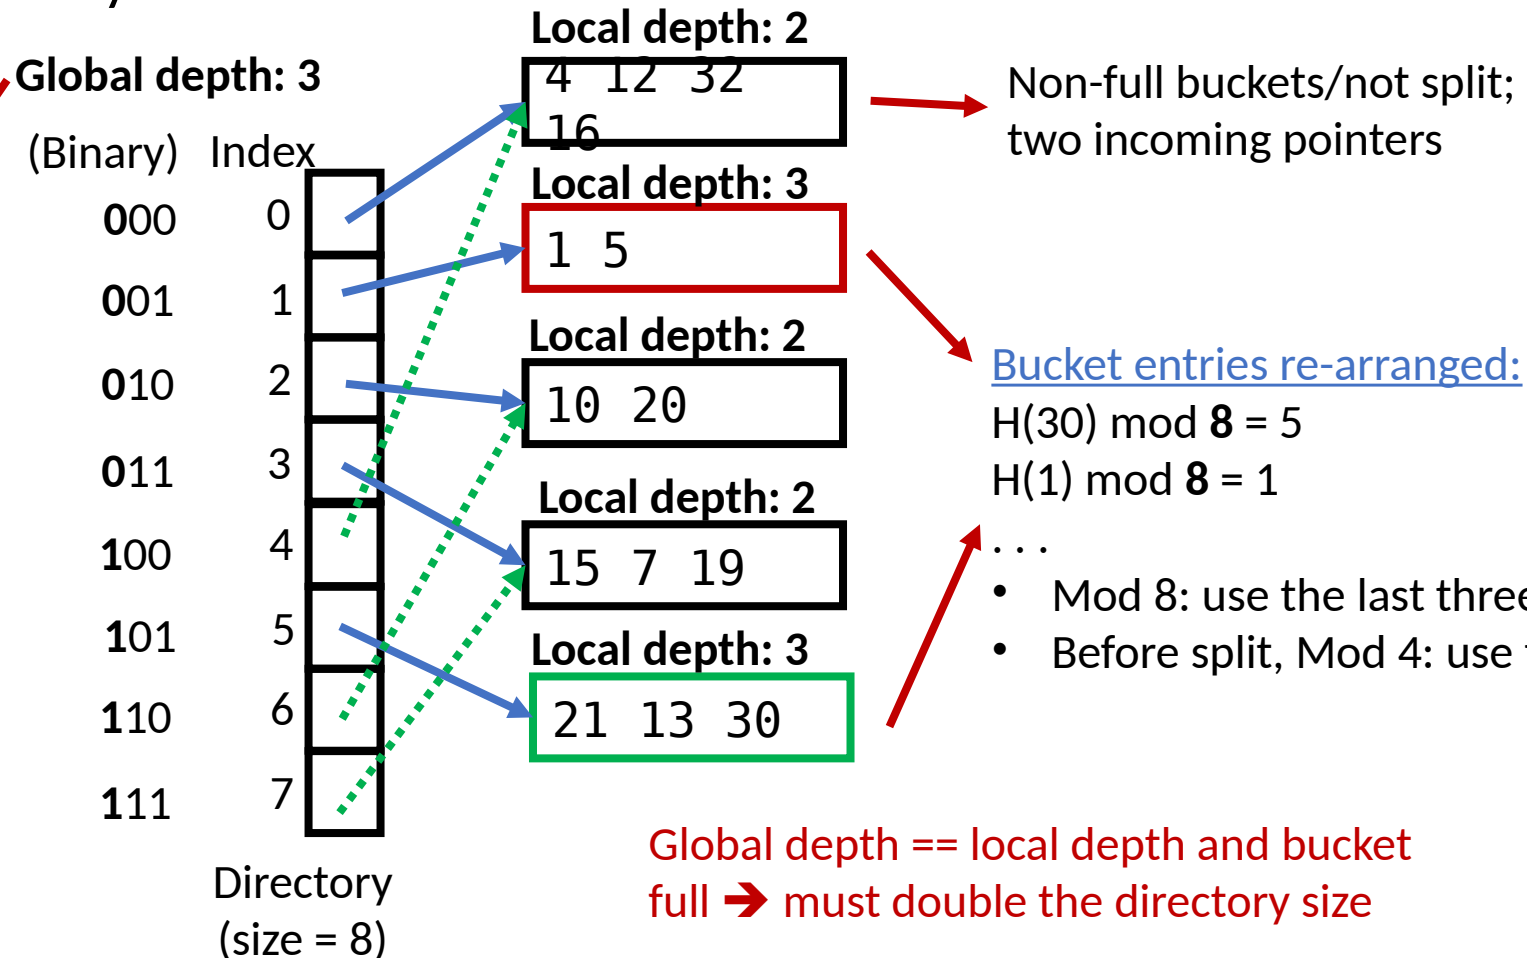


Extensible Hashing

Insert 30: assume $H(30) \bmod 4 = 1$, bucket 1 is already full

- Double the directory size to 8

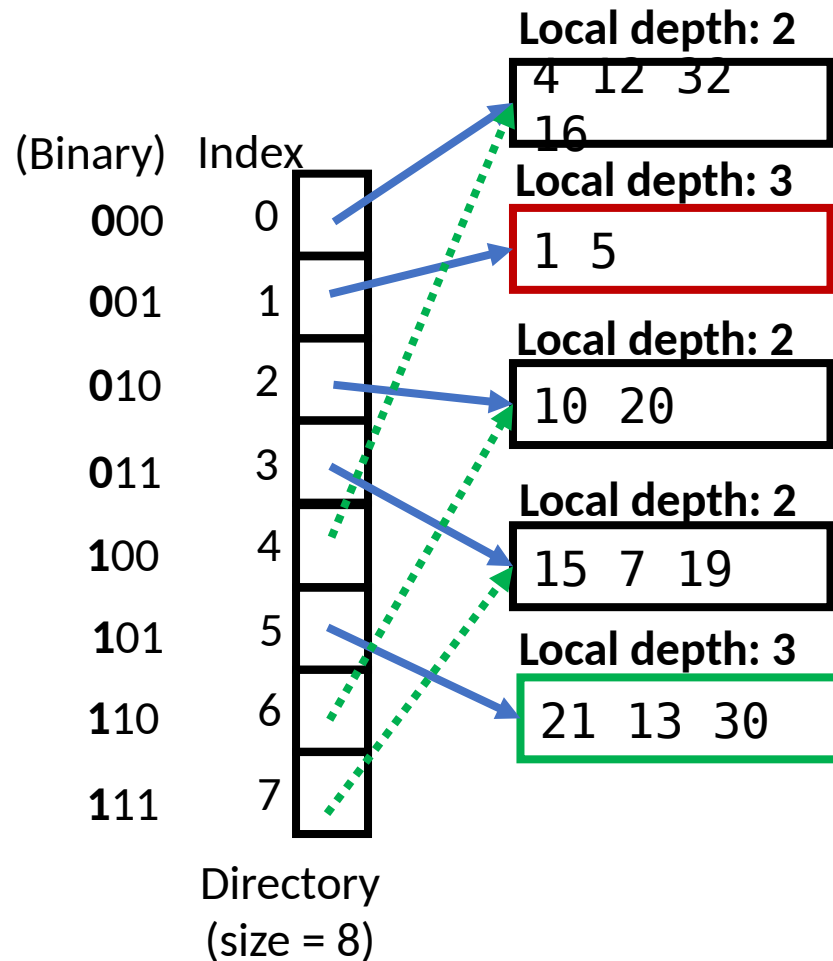
- # of bits to represent all slots in the directory
- +1 when the directory is doubled



Extensible Hashing

- Global depth: the last D bits of $H(\text{key})$ used to decide bucket number
 - $2^{\text{Global depth}}$ = Number of total buckets (i.e., directory size)
 - Example: suppose $H(\text{"abc"}) = 20$,
 - Decimal 20 = binary 10100, global depth = 2
 - ➔ Bucket number = $20 \bmod 4$ (or taking the last 2 bits of 20) = 0
 - Increment global depth by 1 whenever directory is doubled
- Directory doubling not always necessary when a bucket is full
 - Only happens if local depth of the bucket == global depth
 - I.e., no available slot in the directory for the new bucket
 - Bucket pages in file are allocated gradually as they are split
- Each directory entry contains a pointer (page ID)
 - Much smaller than the actual file, more chance to fit in memory
 - Search cost: one I/O if directory already in memory, two if not

Extensible Hashing



Delete:

- Remove entry
- Merge empty bucket
- Decrement local depth
- Halve the directory if every slot points to the same split image

Linear Hashing

Alternative to extendible hashing, without using directory

Basic idea:



- Use a family of hash functions H_0, H_1, H_2, \dots
 - Each hash function's range is twice that of its predecessor
 - H_0 maps to M buckets, H_1 maps to $2M$ buckets
 - E.g., H_0 : 0-7, H_1 : 0-15, H_2 : 0-31, and so on
 - Can be obtained by having a base function H , and define:
 - $H_n(\text{key}) = H(\text{key}) \bmod (2^n * M)$
 - $H_0(\text{key}) = H(\text{key}) \bmod (1 * M)$
 - $H_1(\text{key}) = H(\text{key}) \bmod (2 * M)$
 - $H_2(\text{key}) = H(\text{key}) \bmod (4 * M)$
 - ➔ H_{n+1} doubles the range of H_n
- Use overflow pages and choose bucket to split round-robin
- Main buckets are stored sequentially in file

Linear Hashing

Rounds of splitting


- Split buckets in rounds; at round n , use H_n and H_{n+1}
- The bucket to split is chosen linearly – one bucket after another, eventually doubling the number of buckets
 - Always start from the **first** bucket

Example: initially four buckets ($M = 4$), max two elements per bucket



Linear Hashing

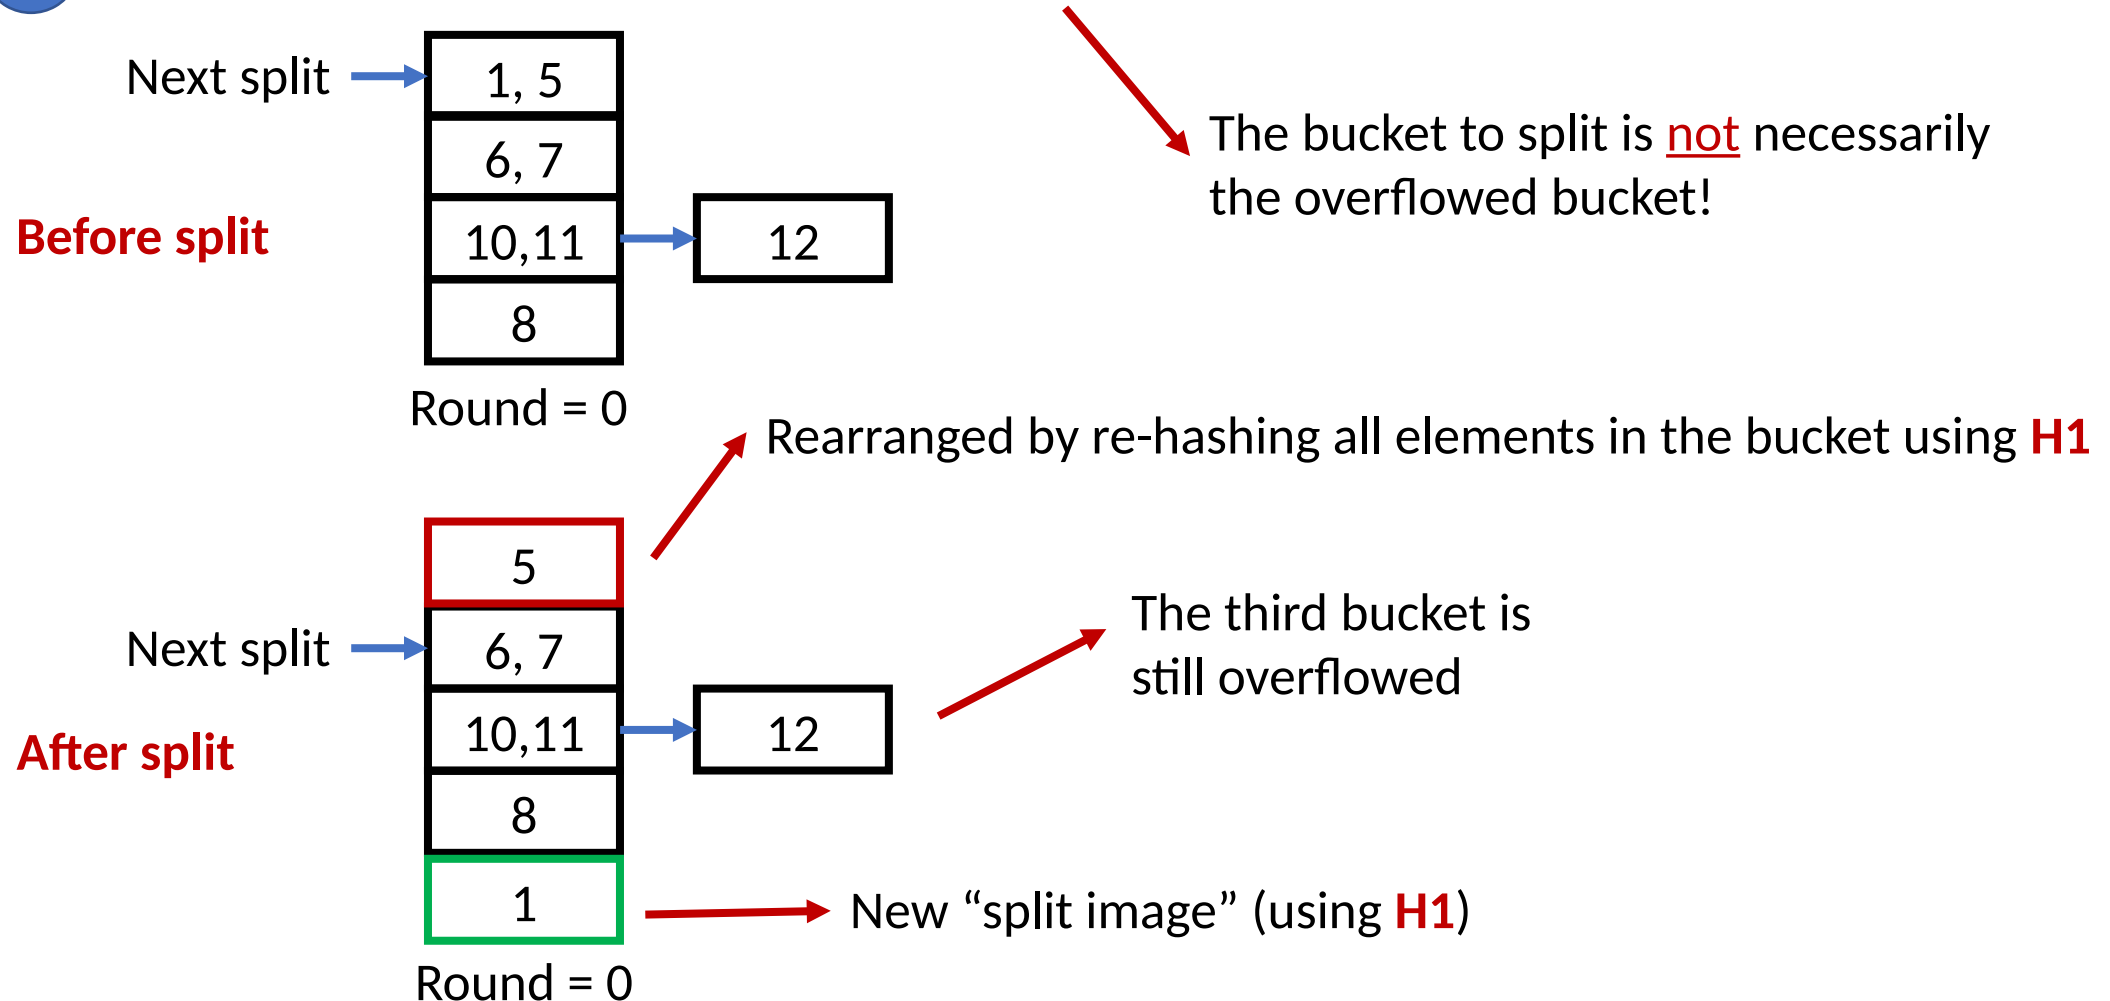
- When should we trigger splits?
- Can use many different criteria
 - E.g., number of buckets with overflows $>$ threshold
 - E.g., longest overflow chain $>$ threshold
 - **E.g., as long as there is any overflow**



Simplistic example
policy in the rest of
the slide deck

Linear Hashing

- 3 Start by splitting the bucket pointed to by [next split] (the first bucket [1, 5])



Linear Hashing

Search for Key K:

- Need two hash function: one for buckets before the split, the other for buckets that are already split


Inputs:

- Key K
- Current split round N

Decide which hash functions to use

Algorithm:

- Do $b = H_n(K)$ first
- If $b \geq \text{next_split}$: done (b is the final bucket number)
- Otherwise $b = H_{n+1}(K)$



Linear Hashing

Insert Key K:

- Similar to search: need two hash function: one for buckets before the split, the other for buckets that are already split


Inputs:

- Key K
- Current split round N

Decide which hash functions to use


Algorithm:

- Let $b = H_n(K)$
- If $b < \text{next_split}$, let $b = H_{n+1}(K)$
- Insert into bucket b, overflow if needed




Linear Hashing

4 Continue to split the second bucket




Before split



After split

Linear Hashing


5 Continue to split the third bucket



Round = 0

Before split

10, 11, and 12 re-hashed using H1 and placed in the corresponding buckets




Round = 0

After split

Linear Hashing

6 After all splits




Round = 0

7

Final steps at the end of a split round:

- Move next_split pointer back to bucket 0
- Increment round number by 1



I.e., double the
number of buckets

Round = 1

H1 and H2 in use

Other Issues

- “Skewed data distribution”
 - Good hash function outputs uniformly distributed hash values
 - Typically ‘skewed data distribution’ == hash values are not uniformly distributed
- Space utilization
 - Load factor: number of stored keys vs. capacity
 - Raw utilization: metadata vs. data (key-value) size
- Concurrency
 - Global latch – low concurrency
 - Per-bucket latch
 - Latch-free – resizing is non-trivial

Summary

- Hash tables
 - Best for point queries, cannot do range search (at least not easily)
 - Hash function maps each key to a bucket
 - Each bucket may store multiple key-value pairs but has limited capacity
- Static hashing: Fixed number of buckets + overflow buckets upon collision
 - Collision: $H(K1) == H(K2)$ but $K1 \neq K2$
 - Long overflow chains possible, impact performance
- Dynamic approaches
 - Extendible hashing – use directory to avoid overflow buckets
 - Linear hashing – avoid directory, split buckets round-robin
 - Adaptations for in-memory use
- Other issues: skewness, space utilization and concurrency