

CSE 541: Database Systems I

Query Processing – Join Operations

Relational Algebra Operators

- Selection (σ)
 - Select a subset of rows from a table (relation)
- Projection (π)
 - Pick a subset of columns
- Set difference
 - $S1 - S2$: return tuples that appear in $S1$ but not in $S2$
- Union
 - $S1 \cup S2$: return tuples in $S1$ and $S2$
- Aggregation
 - Operations such as SUM, MIN and GROUP BY
- Join
 - Combine two relations based on some conditions

Join Operators

Goal: join two relations R and S based on a predicate

```
Sailors(sid: integer, sname: string, rating: integer, age:
real)
Reserves(sid: integer, bid: integer, day: dates, rname:
string)
Reserves ⋈Reserves.sid = Sailors.sid Sailors
```

Algorithms:

- Nested loops join
 - Page nested loops join
 - Block nested loops join
 - Index nested loops join
- Sort-merge join
- Grace hash join

Nested Loop Join

Basic idea: scan outer relation **R**, for each tuple in R, scan the entire inner relation **S**

```
foreach tuple r in R do:
    foreach tuple s in S do:
        if join-condition(ri, sj) then add <r, s> to
result
```

“tuple-at-a-time”

- Desired: low cost (fewer I/O operations)

Cost (number of I/Os): $M + pR * M * N$

- N I/Os to scan S once (S stored in N pages)
- S scanned for $pR * M$ times

M I/Os to scan the entire R (R stored in M pages)

Number of records in R (pR records per page)

Nested Loop Join

Cost (number of I/Os): $M + pR * M * N$

```
Reserves(sid: integer, bid: integer, day: dates, rname: string)
Sailors(sid: integer, sname: string, rating: integer, age: real)
```

Reserves table:

- Tuple size: 40 bytes
- 100 tuples per page
- 1000 pages in total

Sailors table:

- Tuple size: 50 bytes
- 80 tuples per page
- 500 pages in total

- If $R = \text{Reserves}$, $S = \text{Sailors}$, then $M = 1000$, $pR = 100$, $N = 500$
 - Cost = $1000 + 100 * 1000 * 500 = 1000 + 5 * 10^7$ I/Os
 - Assuming 5ms per I/O → about 70 hours to finish!
- What if we change the join order so Sailors becomes the outer table?
 - $M = 500$, $N = 1000$, $pR = 80$, Cost = $500 + 80 * 500 * 1000$ I/Os
 - Smaller than using Reserves as outer relation

→ Should choose the smaller relation as the outer relation

Page Nested Loop Join

Using the smaller table as the outer relation reduces cost, but not significantly

Solution: scan the inner relation for each page (instead of each tuple) of the outer relation (“page-at-a-time”)

```
foreach rpage in R do:
  foreach spage in S do:
    foreach tuple r in rpage do:
      foreach tuple s in spage do:
        if join-condition(ri, sj) then add <r, s> to result
```

Cost (number of I/Os): $M + M * N$

M I/Os to scan the entire R
(R stored in M pages)

Scan S for each page in M

- If R = Sailors, S = Reserves, total cost = $500 + 500 * 1000$ I/Os
 - Much lower than tuple-at-a-time

Block Nested Loop Join

“Block” at a time

- A block may include many pages,
- Read in a block of outer relation each time


Suppose there are B pages in the buffer pool:

```
foreach rblock of B-2 pages in R do:
  foreach spage in S do:
    for all matching tuples in spage and rblock do:
      add <r, s> to result
```

Note: block size = B - 2

- One remaining page for writing out results
- One remaining page for reading in the inner relation

Cost: $M + \lceil M / (B - 2) \rceil * N$

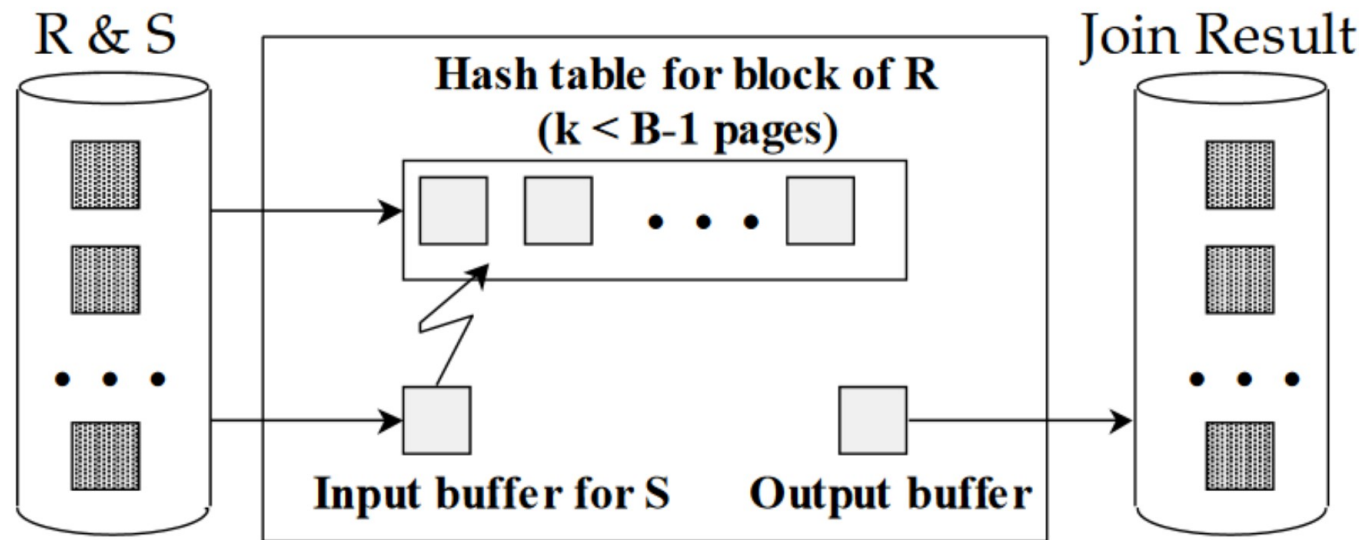


M I/Os to scan R

Number of times needed to scan N

Block Nested Loop Join Refinements

- Build a main-memory hash table for outer relation block
 - Take each S-tuple and query the R-block hash table in memory



- The hash table occupies slightly larger space than the block itself, but reduces complexity in finding matching pairs
- If there is enough memory: build a hash table for the entire inner relation → Cost = M + N

Index Nested Loop Join

If there is an index on the join column of one relation, make it the inner and exploit the index

- E.g., join R and S where R's column i == S's column j
 - There is an index on S's join column S_j
 - Use R's join column as the search key to look for matching tuples in S

foreach tuple r in R do:

 foreach tuple s in S where $r_i == s_j$:

 add $\langle r, s \rangle$ to result

Look up r_i in the
index on S

Cost: $M + pR * M$ * cost to find matching tuples in S (using index)

M I/Os to scan R
(outer relation)

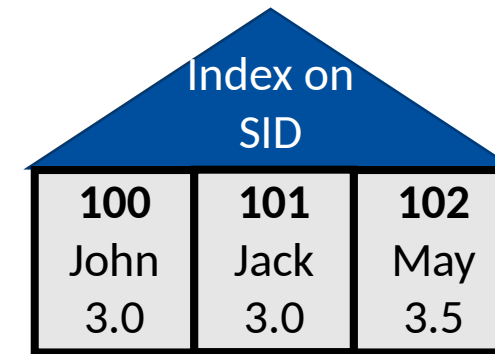
Number of records in R (i.e., number of lookups in S's index)

- Cost of accessing index may vary
 - What's stored in data entries
 - Clustered vs. unclustered

Recap: Index Data Entries

Alternative 1: Actual record

- A special file organization, index == file
- At most one such index per table
 - May need to duplicate data otherwise



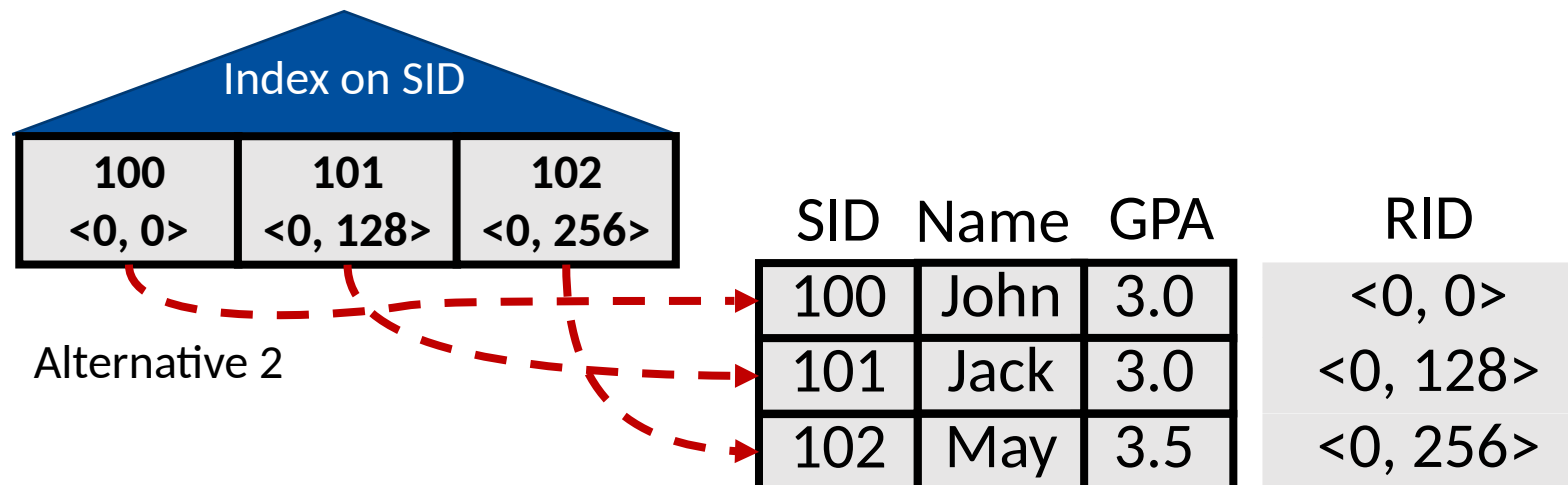
The diagram shows a blue triangle labeled "Index on SID" pointing down to a table. The table has three columns, each with a header (100, 101, 102) and two rows of data (Name and GPA).

| 100 | 101 | 102 |
|------|------|-----|
| John | Jack | May |
| 3.0 | 3.0 | 3.5 |

Alternative 1

Alternative 2: <key, RID>

- Map keys to RIDs
- Independent of the table's organization




Recap: Index Data Entries

Alternative 3: <key, RID list>

- A list of records that match the search key
- Independent of the table's organization

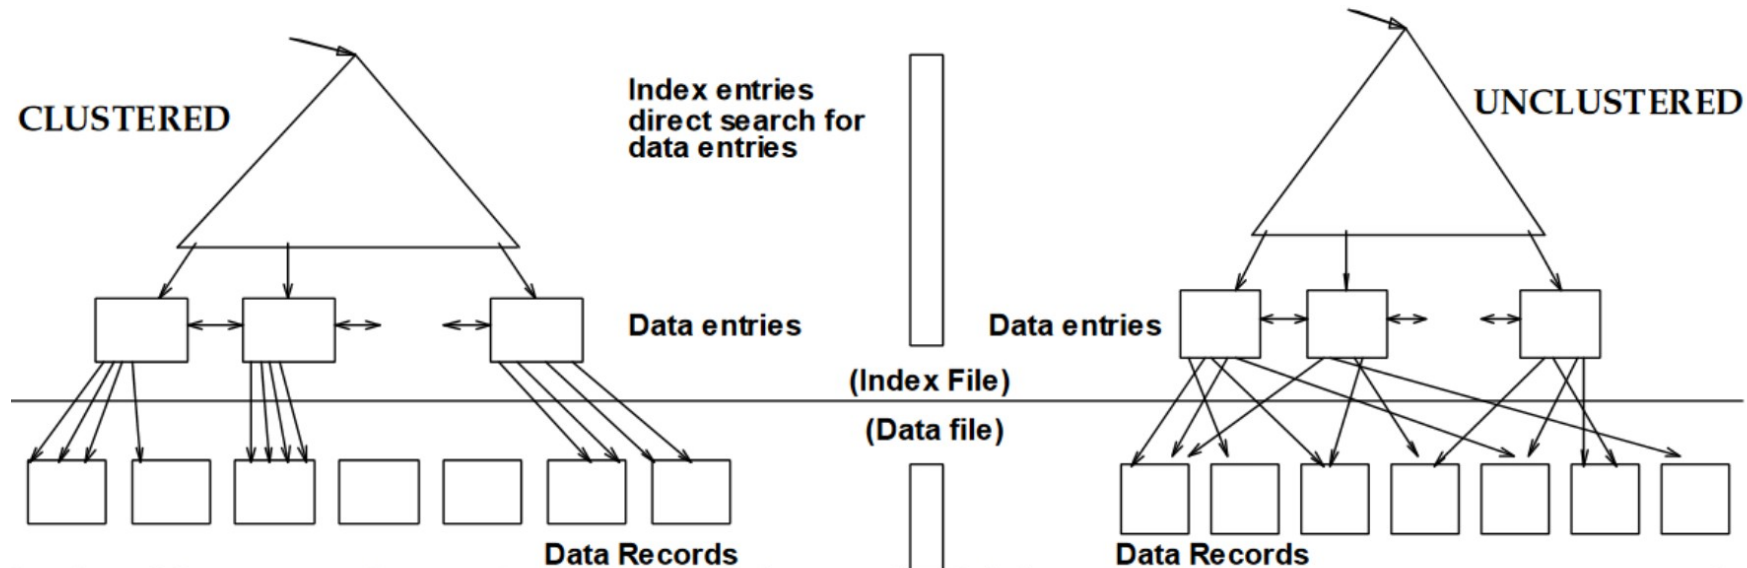
| | Name | GPA | RID |
|-----|------|-----|----------|
| 100 | John | 3.0 | <0, 0> |
| 101 | Jack | 3.0 | <0, 128> |
| 102 | May | 3.5 | <0, 256> |



Alternatives 2 vs. 3:

- Basically the same for unique indexes
- For non-unique indexes, Alternative 3 provides better space utilization (no repeated key storage)

Recap: Clustered vs. Unclustered Index



Order in index == order in data file

- Better performance for scans (faster sequential reads)

Order in index != order in data file

- May be slower: potentially more random reads

Alternative 1 (key → actual record): always clustered

Alternative 2 and 3 (key → RIDs): could be clustered or non-clustered

Index Nested Loop Join

Cost: $M + pR * M$ * cost to find matching tuples in S (using index)

Case 1: index stores records directly as data entries (Alternative 1)

- Cost to traverse from root to leaf in trees or loading hash table pages

Case 2: index stores RIDs or RID lists (Alternatives 2 and 3)

- Cost = cost to look up RIDs + cost to retrieve actual records
- Look up RIDs
 - E.g., traverse B-tree, typically 2-4 I/Os
- Retrieving records from RIDs
 - Unclustered index: up to 1 I/O per matching S tuple
 - Clustered index: 1 I/O per page of matching S tuples
 - Cannot apply the equation: " $M + pR * M$ * cost to find matching tuples in S"

Example: Index Nested Loop Join

```
Reserves(sid: integer, bid: integer, day: dates, rname: string)
Sailors(sid: integer, sname: string, rating: integer, age: real)
```

Reserves table:

- Tuple size: 40 bytes
- 100 tuples per page
- 1000 pages in total

Sailors table:

- Tuple size: 50 bytes
- 80 tuples per page
- 500 pages in total

Query: Join Sailors and Reserves on **sid** column

Cost: $M + pR * M$ * cost to find matching tuples in S (using index)

- Sailors as the inner table, with index on Sailors.sid
 - $M = 1000$ (# of pages in Reserves), $pR = 100 * 1000$
- Cost to find matching tuples using index (assuming a B-tree with 3 levels including internal and leaf levels)
 - With an unclustered B-tree index: 3 I/Os for tree traversal + 1 I/O for accessing heap file
→ $\text{Cost} = M + pR * M * 4 = 1000 + 100 * 1000 * 4$
 - With a clustered B-tree index: 3 I/O for tree traversal, but heap access is based on the number of distinct values of the outer table (1 I/O per distinct value)
→ $\text{Cost} = M + pR * M * 3 + \text{distinct_vals}(\text{Reservers}) = 1000 + 100 * 1000 * 3 + X$ where $X \leq \text{max number of tuples in Reserves } (100 * 1000)$

Sort-Merge Join

Two stages:

- Sort: sort tuples in relations R and S by the join key
- Join: merge-scan the sorted relations, output matching tuples

Example query: Reserves ⋈_{Reserves.sid = Sailors.sid} Sailors

Reserves (sorted)

| sid | bid | day | rname |
|-----|-----|--------------|--------|
| 28 | 103 | Jan 3, 2019 | Guppy |
| 28 | 103 | Jan 4, 2019 | Yuppy |
| 31 | 101 | Mar 3, 2019 | Dustin |
| 31 | 102 | Apr 25, 2019 | Lubber |
| 31 | 101 | May 12, 2019 | Lubber |
| 58 | 103 | May 20, 2019 | Dustin |

Sailors (sorted)

| sid | sname | rating | age |
|-----|--------|--------|-----|
| 22 | Dustin | 7 | 45 |
| 28 | Yuppy | 9 | 35 |
| 31 | Lubber | 8 | 55 |
| 44 | Guppy | 5 | 35 |
| 58 | Rusty | 10 | 35 |

Sort-Merge Join

- Keep two cursors (r, s) pointing to the “current” Reserves tuple and Sailors tuple
 - Advance r until $*r \geq *s$, advance s until $*s \geq *r$
 - Mark the start of the current “block” in S, i.e., let mark = s and start to generate output tuples
- ➔ There might be repeated values in R, the mark allows us to “come back” to join the matching Sailors tuples for each duplicate R tuple

Step 0:

Reserves (sorted)

| | sid | bid | day | rname |
|-----|-----|-----|--------------|--------|
| r → | 28 | 103 | Jan 3, 2019 | Guppy |
| | 28 | 103 | Jan 4, 2019 | Yuppy |
| | 31 | 101 | Mar 3, 2019 | Dustin |
| | 31 | 102 | Apr 25, 2019 | Lubber |
| | 31 | 101 | May 12, 2019 | Lubber |
| | 58 | 103 | May 20, 2019 | Dustin |

Sailors (sorted)

| | sid | sname | rating | age |
|-----|-----|--------|--------|-----|
| s → | 22 | Dustin | 7 | 45 |
| | 28 | Yuppy | 9 | 35 |
| | 31 | Lubber | 8 | 55 |
| | 44 | Guppy | 5 | 35 |
| | 58 | Rusty | 10 | 35 |

|| ← m

Sort-Merge Join

Step 1: Reserves (sorted)

| | sid | bid | day | rname |
|-----|-----|-----|--------------|--------|
| r → | 28 | 103 | Jan 3, 2019 | Guppy |
| | 28 | 103 | Jan 4, 2019 | Yuppy |
| | 31 | 101 | Mar 3, 2019 | Dustin |
| | 31 | 102 | Apr 25, 2019 | Lubber |
| | 31 | 101 | May 12, 2019 | Lubber |
| | 58 | 103 | May 20, 2019 | Dustin |

Sailors (sorted)

| | sid | sname | rating | age | |
|-----|-----|--------|--------|-----|-----|
| | 22 | Dustin | 7 | 45 | |
| s → | 28 | Yuppy | 9 | 35 | ← m |
| | 31 | Lubber | 8 | 55 | |
| | 44 | Guppy | 5 | 35 | |
| | 58 | Rusty | 10 | 35 | |

- Advanced both r and s, mark s
- Current *r matches *s: output result and advance s

Sort-Merge Join

Step 2: Reserves (sorted)

| r → | sid | bid | day | rname |
|------------|-----|-----|--------------|--------|
| | 28 | 103 | Jan 3, 2019 | Guppy |
| | 28 | 103 | Jan 4, 2019 | Yuppy |
| | 31 | 101 | Mar 3, 2019 | Dustin |
| | 31 | 102 | Apr 25, 2019 | Lubber |
| | 31 | 101 | May 12, 2019 | Lubber |
| | 58 | 103 | May 20, 2019 | Dustin |

Sailors (sorted)

| sid | sname | rating | age |
|------------|--------|--------|-----|
| 22 | Dustin | 7 | 45 |
| 28 | Yuppy | 9 | 35 |
| s → | 31 | Lubber | 8 |
| | 44 | Guppy | 5 |
| | 58 | Rusty | 10 |

← **m**

Output
(so far):

| sid | bid | day | rname | sname | rating | age |
|-----|-----|-------------|-------|-------|--------|-----|
| 28 | 103 | Jan 3, 2019 | Guppy | Yuppy | 9 | 35 |

Now *r and *s do not match, next step:

- Reset s to m, advance r, reset mark to null

Sort-Merge Join

Step 3: Reserves (sorted)

| | | | | |
|------------|------------|------------|--------------|--------------|
| | sid | bid | day | rname |
| | 28 | 103 | Jan 3, 2019 | Guppy |
| r → | 28 | 103 | Jan 4, 2019 | Yuppy |
| | 31 | 101 | Mar 3, 2019 | Dustin |
| | 31 | 102 | Apr 25, 2019 | Lubber |
| | 31 | 101 | May 12, 2019 | Lubber |
| | 58 | 103 | May 20, 2019 | Dustin |

Sailors (sorted)

| | | | | |
|------------|------------|--------------|---------------|------------|
| | sid | sname | rating | age |
| | 22 | Dustin | 7 | 45 |
| s → | 28 | Yuppy | 9 | 35 |
| | 31 | Lubber | 8 | 55 |
| | 44 | Guppy | 5 | 35 |
| | 58 | Rusty | 10 | 35 |

|| ← m

Output
(so far):

| | | | | | | |
|------------|------------|-------------|--------------|--------------|---------------|------------|
| sid | bid | day | rname | sname | rating | age |
| 28 | 103 | Jan 3, 2019 | Guppy | Yuppy | 9 | 35 |

Next:

- Position r and s so that for Reserves $*r \geq *s$, for Sailors $*s \geq *r$, and mark s

Sort-Merge Join

Step 4: Reserves (sorted)

| r → | sid | bid | day | rname |
|------------|-----|-----|--------------|--------|
| | 28 | 103 | Jan 3, 2019 | Guppy |
| | 28 | 103 | Jan 4, 2019 | Yuppy |
| | 31 | 101 | Mar 3, 2019 | Dustin |
| | 31 | 102 | Apr 25, 2019 | Lubber |
| | 31 | 101 | May 12, 2019 | Lubber |
| | 58 | 103 | May 20, 2019 | Dustin |

Sailors (sorted)

| s → | sid | sname | rating | age | m ← |
|------------|-----|--------|--------|-----|------------|
| | 22 | Dustin | 7 | 45 | |
| | 28 | Yuppy | 9 | 35 | |
| | 31 | Lubber | 8 | 55 | |
| | 44 | Guppy | 5 | 35 | |
| | 58 | Rusty | 10 | 35 | |

Output
(so far):

| sid | bid | day | rname | sname | rating | age |
|-----|-----|-------------|-------|-------|--------|-----|
| 28 | 103 | Jan 3, 2019 | Guppy | Yuppy | 9 | 35 |

Current *r matches *s, next:

- Output result and advance s

Sort-Merge Join

Step 5: Reserves (sorted)

| | sid | bid | day | rname |
|-----|-----|-----|--------------|--------|
| | 28 | 103 | Jan 3, 2019 | Guppy |
| r → | 28 | 103 | Jan 4, 2019 | Yuppy |
| | 31 | 101 | Mar 3, 2019 | Dustin |
| | 31 | 102 | Apr 25, 2019 | Lubber |
| | 31 | 101 | May 12, 2019 | Lubber |
| | 58 | 103 | May 20, 2019 | Dustin |

Sailors (sorted)

| | sid | sname | rating | age |
|-----|-----|--------|--------|-----|
| | 22 | Dustin | 7 | 45 |
| | 28 | Yuppy | 9 | 35 |
| s → | 31 | Lubber | 8 | 55 |
| | 44 | Guppy | 5 | 35 |
| | 58 | Rusty | 10 | 35 |

← m

Output
(so far):

| sid | bid | day | rname | sname | rating | age |
|-----|-----|-------------|-------|-------|--------|-----|
| 28 | 103 | Jan 3, 2019 | Guppy | Yuppy | 9 | 35 |
| 28 | 103 | Jan 4, 2019 | Yuppy | Yuppy | 9 | 35 |

Now *r and *s do not match, next step:

- Reset s to m, advance r, reset mark to null

Sort-Merge Join

Step 6: Reserves (sorted)

| sid | bid | day | rname |
|---------------|-----|--------------|--------|
| 28 | 103 | Jan 3, 2019 | Guppy |
| 28 | 103 | Jan 4, 2019 | Yuppy |
| r → 31 | 101 | Mar 3, 2019 | Dustin |
| 31 | 102 | Apr 25, 2019 | Lubber |
| 31 | 101 | May 12, 2019 | Lubber |
| 58 | 103 | May 20, 2019 | Dustin |

Sailors (sorted)

| sid | sname | rating | age |
|---------------|--------|--------|-----|
| 22 | Dustin | 7 | 45 |
| s → 28 | Yuppy | 9 | 35 |
| 31 | Lubber | 8 | 55 |
| 44 | Guppy | 5 | 35 |
| 58 | Rusty | 10 | 35 |

|| ← m

Output
(so far):

| sid | bid | day | rname | sname | rating | age |
|-----|-----|-------------|-------|-------|--------|-----|
| 28 | 103 | Jan 3, 2019 | Guppy | Yuppy | 9 | 35 |
| 28 | 103 | Jan 4, 2019 | Yuppy | Yuppy | 9 | 35 |

Next:

- Position r and s so that for Reserves $*r \geq *s$, for Sailors $*s \geq *r$, and mark s

Sort-Merge Join

Step 7: Reserves (sorted)

| | sid | bid | day | rname |
|-----|-----|-----|--------------|--------|
| | 28 | 103 | Jan 3, 2019 | Guppy |
| | 28 | 103 | Jan 4, 2019 | Yuppy |
| r → | 31 | 101 | Mar 3, 2019 | Dustin |
| | 31 | 102 | Apr 25, 2019 | Lubber |
| | 31 | 101 | May 12, 2019 | Lubber |
| | 58 | 103 | May 20, 2019 | Dustin |

Sailors (sorted)

| | sid | sname | rating | age |
|-----|-----|--------|--------|-----|
| | 22 | Dustin | 7 | 45 |
| | 28 | Yuppy | 9 | 35 |
| s → | 31 | Lubber | 8 | 55 |
| | 44 | Guppy | 5 | 35 |
| | 58 | Rusty | 10 | 35 |

← m

Output
(so far):

| sid | bid | day | rname | sname | rating | age |
|-----|-----|-------------|-------|-------|--------|-----|
| 28 | 103 | Jan 3, 2019 | Guppy | Yuppy | 9 | 35 |
| 28 | 103 | Jan 4, 2019 | Yuppy | Yuppy | 9 | 35 |

Repeat the previous process until r is EOF

Sort-Merge Join Algorithm

Advance r until $*r \geq *s$, advance s until $*s \geq *r$

```
while (*r != EOF) {  
    if (m == null) {  
        while (*r < *s) { advance r; }  
        while (*s < *r) { advance s; }  
        m = s  
    }  
    if (*r == *s) {  
        output <*r, *s>  
        advance s  
    } else {  
        s = m;  
        advance r  
        m = null  
    }  
}
```

Mark the start of the potential block of matching records in S

*r and *s match:

- Output the join result
- Advance to the next S tuple to find more matches for the same R tuple (*r)

*r and *s do not match

- Advance to the next R tuple and reset to the marked S tuple (*m)
- Need to reset m to null to allow advancing r and s if needed in the next round

Sort-Merge Join (R ⋈ S)

- R is scanned once
- S could be scanned multiple times
 - Each 'block' is scanned per matching R tuple
 - But usually the needed pages are already in buffer pool

Cost: sorting cost for R and S + scanning cost

- Scanning cost
 - Best scenario: $M + N$, if no S block is scanned multiple times
 - Worst scenario: $M * N$, if the repeatedly scanned S block cannot be buffered
 - E.g., the page is no longer in buffer pool when it is requested the second time
 - Query optimizer tries to avoid this when deciding on plans

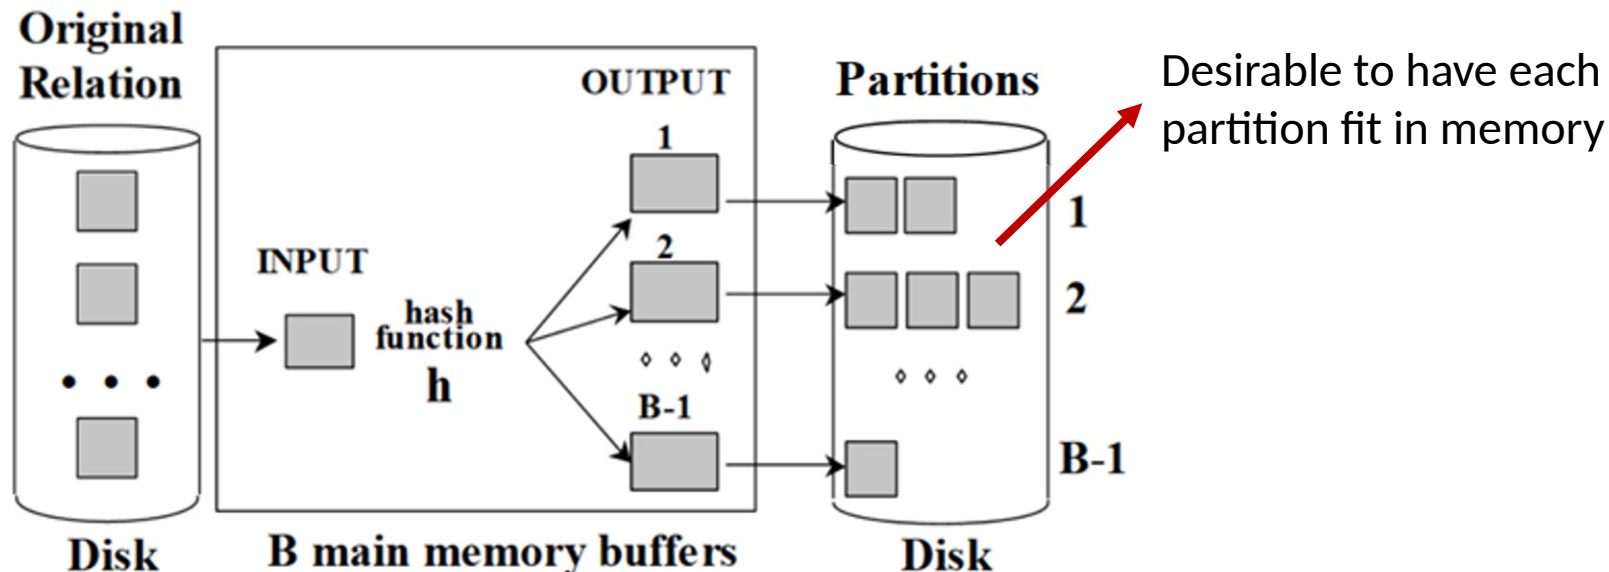
(Grace) Hash Join ($R \bowtie S$)

Leverage hashing to partition input relations, then join by partitions

- Two phases: building and probing

Building (aka partitioning) phase:

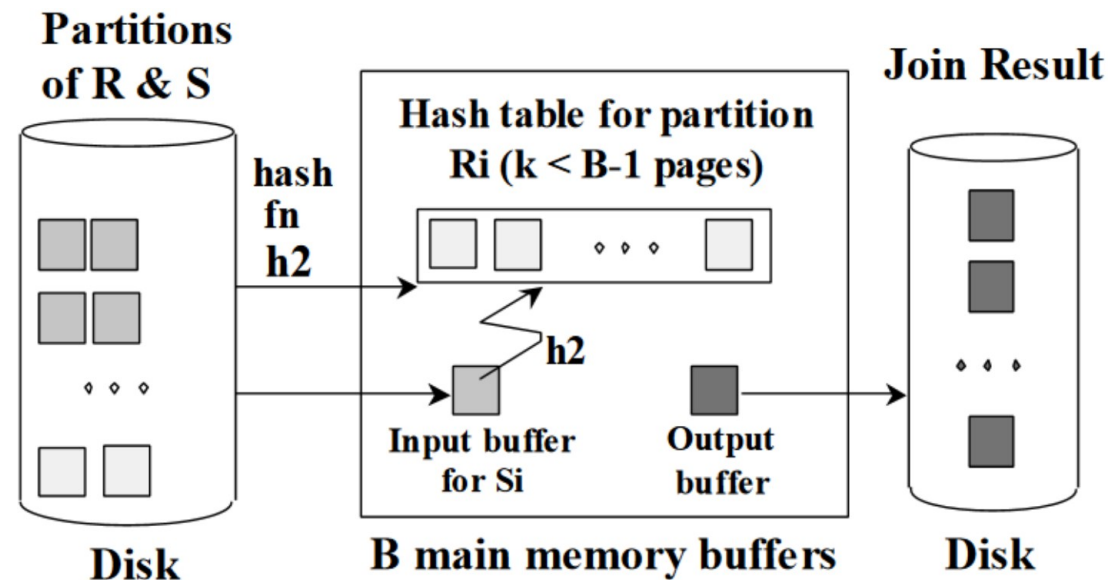
- Hash both relations on the join attribute using the same hash function h into k partitions
- Partitions could be on different disks or even machines



(Grace) Hash Join ($R \bowtie S$)

Probing (aka matching) phase:

- After the building phase, R tuples in partition i can join only with S tuples in its partition i
- Read each R partition and scan just the corresponding S partition for matches
- Reduce CPU cost on matching:
 - Build a hash table for the R partition with a different hash function **h2**
 - Scan S partition and probe the hash table using h2 to find matches



Cost of Grace Hash Join ($R \bowtie S$)

Suppose R has M pages and S has N pages

Number of I/O operations:

- Partitioning phase
 - Read and write both relations: $2 * (M + N)$ I/Os
- Matching phase
 - Read both relations once: $(M + N)$ I/Os
- Total: $3 * (M + N)$ I/Os

Sort-Merge Join vs. Hash Join

Sort-merge join:

- Good if the input is already sorted or the output needs to be sorted
- Less sensitive to data skew or bad hash functions

Hash join:

- Good if the input is already hashed or the output needs to be hashed
- May be vulnerable to data skew/bad hash functions
- Number of passes depends on the size of the smaller relation

General Join Conditions

Equalities over multiple attributes:

- E.g., $R.sid = S.sid$ AND $R.rname = S.sname$
- Index nested loops join
 - Build or use index on $\langle sid, sname \rangle$ (S being the inner relation)
- Sort-merge/hash join
 - Sort/partition on the combined attributes

Inequality conditions:

- E.g., $R.age < S.age$
- Cannot use hash join or sort-merge join
- Index nested loops join: need index that supports range scans
 - Range probes the inner relation
 - Number of matches likely much higher than equality join

Impact of Buffering

- If several operations are executing concurrently, estimating the available and allocating buffer pages is guesswork
- **Repeated access pattern interacts with buffer replacement policy**
- Example 1: the inner relation is often scanned repeatedly in tuple-at-a-time nested loops join
 - LRU could be the worst algorithm to use (sequential flooding)
 - MRU would be much better
- Example 2: in block nested loops join, $B - 2$ buffers are used for the outer relation, 1 for the result
 - The remaining one buffer for reading the inner relation
 - Replacement policy does not matter here

Summary

- Join algorithms
 - Nested loops (tuple/page-at-a-time, block, index)
 - Sort-merge join and hash join
 - Costs, pros and cons
- Common techniques
 - Sorting, hashing, partitioning
 - Index can be helpful and preferred in general to reduce cost
 - Sometimes also need to scan the entire relation
- Buffering impact
 - Relations are sometimes repeatedly scanned
 - Need to watch out for problems like sequential flooding in LRU