

# CSE 541: Database Systems I

## Transactions & Concurrency Control

# Transactions

**Definition:** A transactions is a user program that runs in a DBMS to access (read/write) database records

- A transaction may
  - Access (read/write) one or multiple database records
  - Conduct computation with the records

```
BEGIN TRANSACTION
UPDATE Cities
SET Population = 1000000
WHERE City = Vancouver
COMMIT
```

- Desirable to run multiple transactions concurrently
  - Utilize the abundant parallelism provided by multicore CPUs
  - Keep CPU busy while I/O is in-progress

# **Concurrency vs. Parallelism?**

# Concurrency in DBMSs

- Multiple users can submit transactions to the DBMS
  - The DBMS will interleave read and write actions of different transactions
- A transaction must leave the database in a consistent state
  - DBMS enforces integrity constraints specified in table schema
  - DBMS doesn't know data semantics

```
CREATE TABLE Cities (  
    Id INTEGER,  
    City CHAR(50),  
    Population INTEGER,  
    PRIMARY KEY(Id),  
    CHECK (Population > 0)  
)
```



Integrity  
constraint

# ACID Properties

**Ensured despite concurrent accesses and system failures**

- **A**tomicity
  - Actions in a transaction are either all applied, or not at all
  - Commit – apply all read/write actions
  - Abort – rollback all changes so far, as if nothing happened
- **C**onsistency
  - Must leave the database in a consistent state
  - No data corruption, anomalies, etc. (more on this later)
- **I**solation
  - As if the transaction were the only one in the system
- **D**urability
  - Successful changes must be correctly persisted in storage

# Example

Initial account  
balances:

Alice	Bob
100	100

Consider two transactions submitted **at the same time**

**T1** }

```
BEGIN
Alice = Alice + 100
Bob = Bob - 100
COMMIT
```

Transfer 100 from  
Bob to Alice

**T2** }



```
BEGIN
Alice = Alice * 1.05
Bob = Bob * 1.05
COMMIT
```

Credit a 5% interest  
on both accounts

- Actions are interleaved for high performance
- Many ways to interleave (schedules), but not all are valid

# Transaction Schedules

Schedule: a list of actions from a set of transactions

- Actions: read, write, abort, commit
- Example:
  - T1: Read A, Write B, Commit → denoted  $R_1(A)$ ,  $W_1(B)$ ,  $C_1$
  - T2: Write C, Read B, Commit → denoted  $R_2(C)$ ,  $W_2(B)$ ,  $C_2$
  - A possible schedule:  $R_1(A)$ ,  $R_2(C)$ ,  $W_1(B)$ ,  $C_1$ ,  $W_2(B)$ ,  $C_2$
- Two actions of a transaction T that appear in the schedule, must be in the same order as they are in T
  - Consider the previous example
  - $R_1(A)$ ,  $R_2(C)$ ,  $W_1(B)$ ,  $C_1$ ,  $W_2(B)$ ,  $C_2$  
  - $W_1(B)$ ,  $R_2(C)$ ,  $R_1(A)$ ,  $C_1$ ,  $W_2(B)$ ,  $C_2$  

# Transaction Schedules

Serial schedule: Schedule that does not interleave actions of different transactions

- Transactions executed one after another:

- $R_1(A), W_1(B), C_1, R_2(C), W_2(B), C_2 \rightarrow T_1, T_2$

- $R_2(C), W_2(B), C_2, R_1(A), W_1(B), C_1 \rightarrow T_2, T_1$

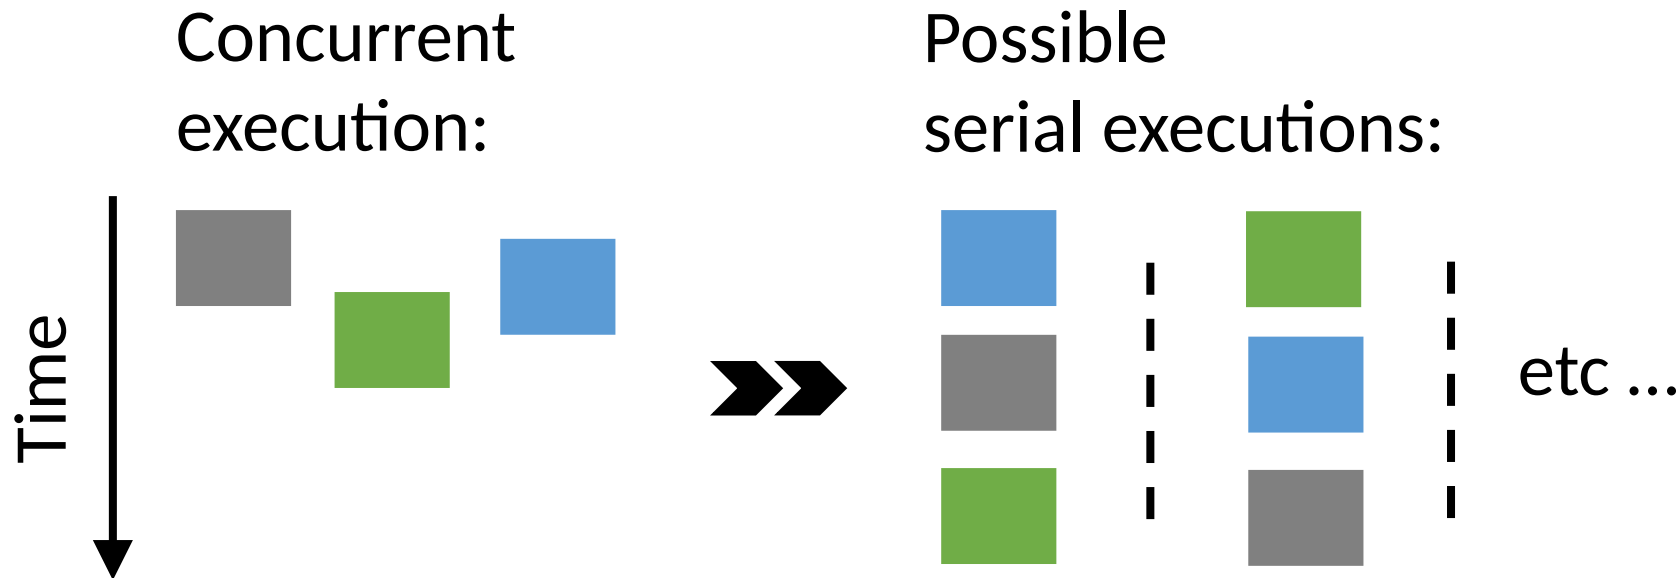
Equivalent schedules: two schedules are equivalent if their final states are the same



# Serializability

Serializable schedule: A schedule that is equivalent to some serial execution's schedule (**doesn't matter which one**)

➔ A DBMS provides serializability if it produces serializable schedules




# Example


Initial account  
balances:

Alice	Bob
100	100

Consider two transactions submitted **at the same time**

**T1**   
BEGIN  
Alice = Alice + 100  
Bob = Bob - 100  
COMMIT

Transfer 100 from  
Bob to Alice

**T2**   
BEGIN  
Alice = Alice \* 1.05  
Bob = Bob \* 1.05  
COMMIT

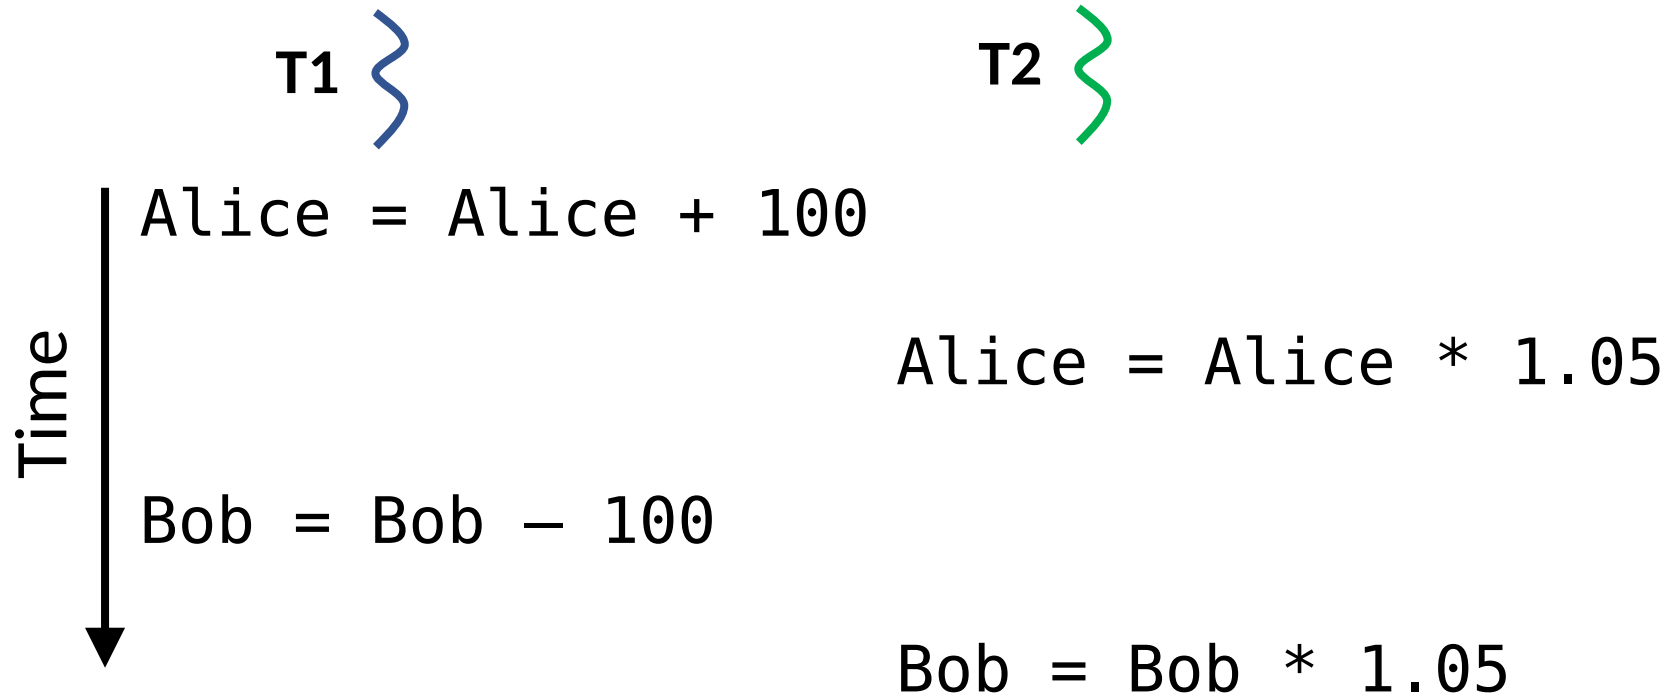
Credit a 5% interest  
on both accounts

Serializable: final effect must be equivalent to that of executing **[T1 then T2]**  
or **[T2 then T1]**

# Schedule 1

Initial account  
balances:

Alice	Bob
100	100



Final results:

Alice	Bob
210	0

**Correct:** as if we executed T1, then T2

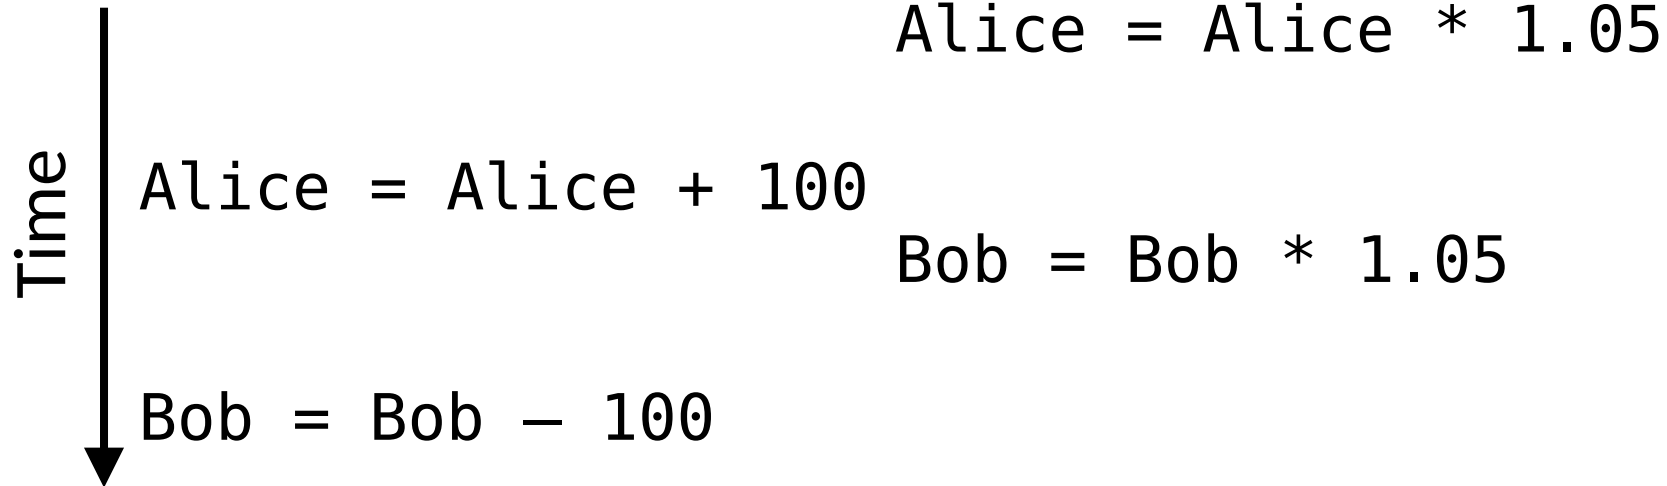
# Schedule 2

Initial account  
balances:

Alice	Bob
100	100

T1 }

T2 }



Final results:

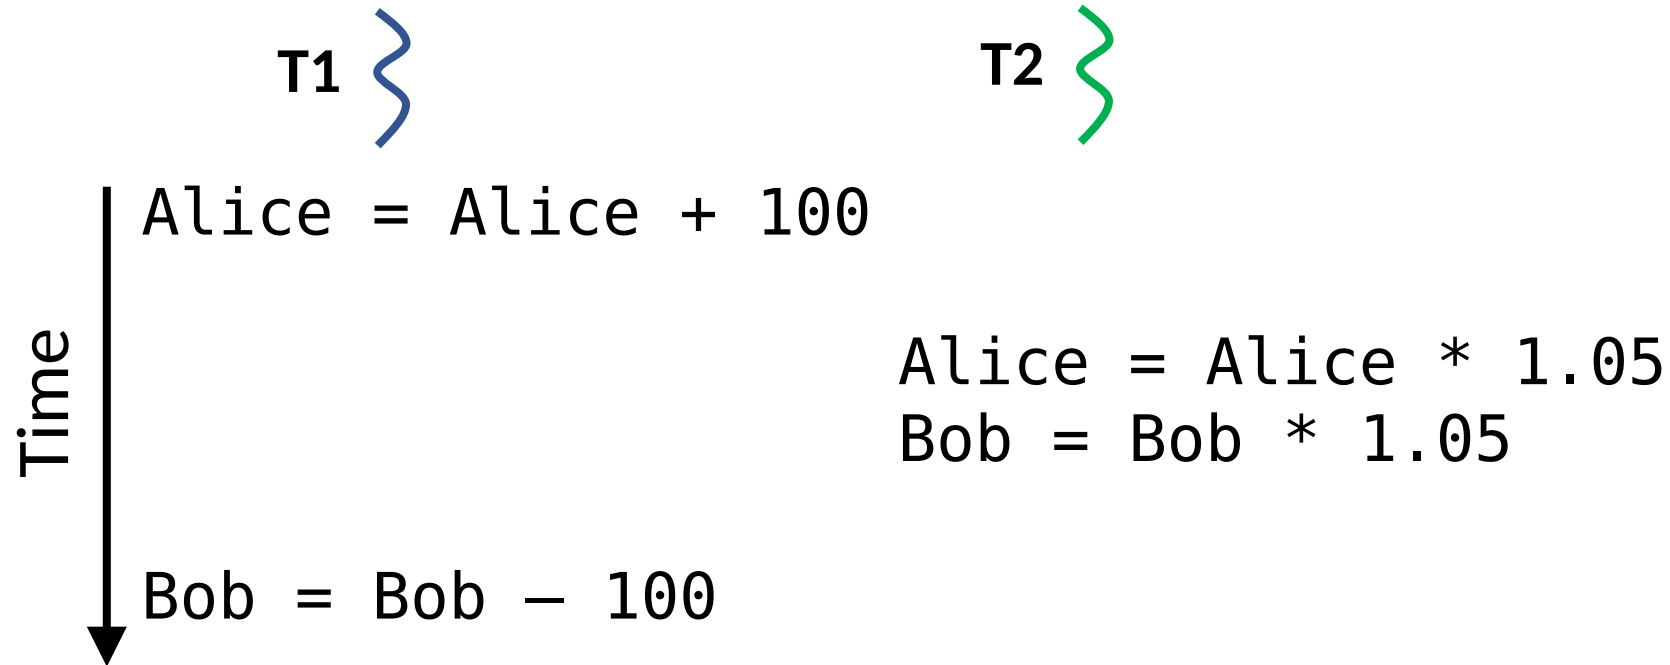
Alice	Bob
205	5

**Correct:** as if we executed T2, then T1

# Schedule 3

Initial account  
balances:

Alice	Bob
100	100



Final results:

Alice	Bob
210	5

**Non-serializable schedule:** no serial execution produces this result. A serializable concurrency control protocol is needed to prevent this.

# Note on Serializability

- Guaranteeing serializability will ensure correct results
- But, **not guaranteeing serializability != wrong results**
  - i.e., lower isolation levels may be “good enough” to give correct results
  - Depends on the workload. Possible already no anomalies arise under a weaker isolation level, e.g., read committed.
  - Guaranteeing serializability incurs overhead, so it might be advantageous to use lower isolation levels
- Example: the classic TPC-C benchmark exhibits no anomalies (i.e., wrong results) under snapshot isolation, which is not serializable
  - More on this later

# Conflict Serializability

Conflicts: In a schedule, two actions conflict if they operation on the same data record and at least one of them is a write

- Write-write conflict: two transactions both write R
- Read-write conflict: one reads R, the other writes R

Non-conflicting operations: order change does not affect final outcome

- Operations from different transactions and accessing different records, or reading the same record
- Consider schedule: R<sub>1</sub>(A), R<sub>2</sub>(C), W<sub>1</sub>(B), C<sub>1</sub>, W<sub>2</sub>(B), C<sub>2</sub>
  - Switching R<sub>1</sub>(A) and R<sub>2</sub>(C) will not change the final outcome

Conflicting operations: changing order leads to different final outcome, e.g., two writes

# Conflict Serializability

Conflict equivalent schedules: Two schedules S1 and S2 are conflict equivalent iff S1 can be transformed to S2 by swapping adjacent non-conflicting operations

- Consider H1 and H2, transform H2 to H1

- S1: R<sub>1</sub>(A)   R<sub>2</sub>(C)   W<sub>1</sub>(B)   C<sub>1</sub>   W<sub>2</sub>(B)   C<sub>2</sub>
- S2: R<sub>1</sub>(A)   W<sub>1</sub>(B)   R<sub>2</sub>(C)   W<sub>2</sub>(B)   C<sub>1</sub>   C<sub>2</sub>



Conflict serializable: A schedule S that is conflict equivalent to a serial schedule

- S can be transformed by swapping non-conflicting operations to a serial schedule



# Conflict Serializability

## Necessary but not sufficient condition for serializability

- Every conflict serializable schedule is serializable, but not vice versa

- Example


- H1: W<sup>1</sup>(A)    W<sup>2</sup>(B)    W<sup>1</sup>(B)    W<sup>2</sup>(A)    W<sup>3</sup>(B)

- H2: W<sup>1</sup>(A)    W<sup>1</sup>(B)    W<sup>2</sup>(B)    W<sup>2</sup>(A)    W<sup>3</sup>(B)

T1

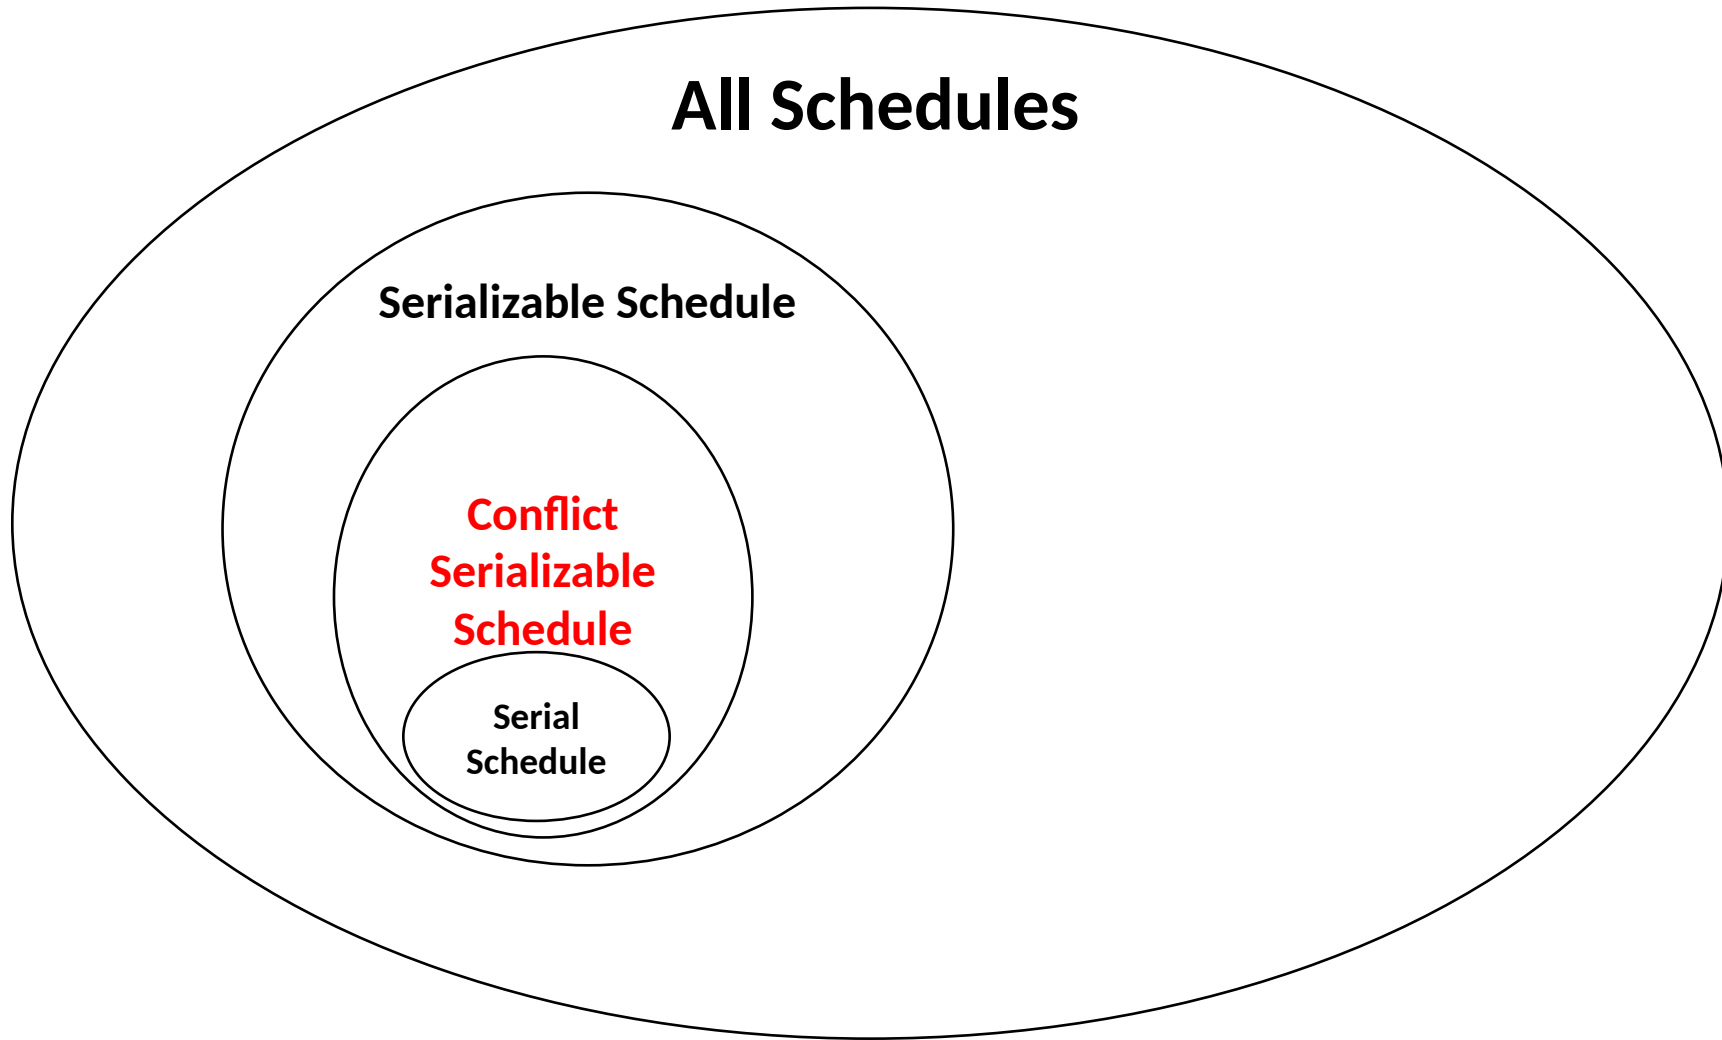
T2

T3

Serial  
schedule  


- The final result of H1 and H2 are the same
  - Definition of serializability: final outcome same as some serial schedule's
  - ➔ H1 is serializable!
- But there are no way to transform H1 to H2 by swapping **adjacent, non-conflicting pairs**, so H1 is not conflict-serializable

# Conflict Serializable Schedule




# Dependency Graph

An effective way to capture potential conflicts between transactions in a schedule (aka “precedence graph”)

- Node: committed transaction
- Directed Edge: from T1 to T2 if T1 precedes and conflicts with one of the actions of T2

**Schedule is conflict-serializable iff the dep. graph is acyclic**



# Anomalies

- Reading uncommitted data (“dirty reads”)

T1: R(A), W(A), R(B), W(B), Abort  
T2: R(A), W(A), C

- Unrepeatable Reads

T1: R(A), R(A), W(A), C  
T2: R(A), W(A), C

- Overwriting committed data (“lost updates”)

T1: W(A), W(B), C  
T2: W(A), W(B), C

- Phantoms

- The same range scan executed twice in the same transaction return different results

# ANSI Isolation Levels

- Serializability gives the highest correctness guarantee
- There are also lower-level guarantees (may corrupt data or generate wrong results, depending on the workload)
- **Isolation levels defined by ANSI**
  - Read Uncommitted
  - Read Committed: no dirty reads
  - Repeatable Read: Read Committed + no non-repeatable reads
  - Serializable: Repeatable Reads + no phantom
- ANSI definitions are lock-based and not comprehensive
  - E.g., Snapshot Isolation isn't truly serializable, yet does not produce any of the anomalies in ANSI SQL-92